

Reinforcement Learning

Russell and Norvig: Chapter 21 (21.1-21.4)

CSE 240: Winter 2023

Lecture 17 (or 18?)

Announcements

- Assignment 4 is due on Friday at 5pm
- Assignment 5 posted and it is due on Wednesday: 03/22 of finals week
- Last Quiz due *next* Friday: 03/17 at 5pm
 - We will drop the lowest quiz
- Please do the course evaluations
 - We will award 1 point of extra credit if over $> 80\%$ of the course do the evaluations.
 - The evaluations matter.
- ~~Guest speakers~~


Agenda and Topics

- Reinforcement Learning
 - Motivation
 - Model-base Learning
 - Model-free Learning
 - Passive RL
 - Direct Evaluation
 - Temporal Difference Learning
 - Q-Learning (if time)

Reinforcement Learning


- Reinforcement learning:
 - Still assume an MDP:
 - A set of states $s \in S$
 - A set of actions (per state) A
 - A model $T(s,a,s')$
 - A reward function $R(s,a,s')$
 - Still looking for a policy $\pi(s)$
 - New twist: don't know T or R
 - i.e. don't know which states are good or what the actions do
 - Must actually try actions and states out to learn

Reinforcement Learning




- Basic idea:
 - Receive feedback in the form of rewards
 - Agent's utility is defined by the reward function
 - Must (learn to) act so as to maximize expected rewards
 - All learning is based on observed samples of outcomes!

Offline (MDPs) vs. Online RL



Offline Solution




Online Learning

Model-Based Learning

Model-Based Learning

- Model-Based Idea:
 - Learn the model empirically through experience
 - Solve for values as if the learned model were correct
- Step 1: Learn empirical model learning
 - Count outcomes for each s, a
 - Normalize to give estimate of $\hat{T}(s, a, s')$
 - Discover $\hat{R}(s, a, s')$ when we experience (s, a, s')
- Solving the MDP with the learned model
 - Iterative policy evaluation, for example

$$V_{i+1}^{\pi}(s) \leftarrow \sum_{s'} \hat{T}(s, \pi(s), s') [\hat{R}(s, \pi(s), s') + \gamma V_i^{\pi}(s')]$$



Analogy: Expected Age

Goal: Compute expected age of students


Known $P(A)$

$$E[A] = \sum_a P(a) \cdot a = 0.35 \times 20 + \dots$$

Model-Free Learning

Passive Learning

- Simplified task
 - You don't know the transitions $T(s,a,s')$
 - You don't know the rewards $R(s,a,s')$
 - You are given a policy $\pi(s)$
 - **Goal: learn the state values**
 - ... what policy evaluation did
- In this case:
 - Learner “along for the ride”
 - No choice about what actions to take
 - Just execute the policy and learn from experience
 - We'll get to the active case soon
 - This is NOT offline planning! You actually take actions in the world and see what happens...



Direct Evaluation

- Goal: Compute values for each state under π .
- Idea: Average together observed sample values:
 - Act according to π
 - Every time you visit a state, write down what the sum of discounted rewards turned out to be.
 - Average those samples
- This is called direct evaluation

Example: Direct Evaluation

Input Policy π

	<div>A</div>	
<div>B</div>	<div>C</div>	<div>D</div>
	<div>E</div>	

Assume: $\gamma = 1$

Observed Episodes (Training)

Episode 1

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 2

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 3

E, north, C, -1
C, east, D, -1
D, exit, x, +10

Episode 4

E, north, C, -1
C, east, A, -1
A, exit, x, -10


Output Values

	<div>A</div>	
<div>B</div>	<div>C</div>	<div>D</div>
	<div>E</div>	

CE 18: Calculate Output Values

For A, C, D

Input Policy π



Assume: $\gamma = 1$

Observed Episodes (Training)

Episode 1

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 2

B, east, C, -1
C, east, D, -1
D, exit, x, +10


Episode 3

E, north, C, -1
C, east, D, -1
D, exit, x, +10

Episode 4

E, north, C, -1
C, east, A, -1
A, exit, x, -10

Output Values



Example: Direct Evaluation

Input Policy π

	A	
B	C	D
	E	

Assume: $\gamma = 1$

Observed Episodes (Training)

Episode 1

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 2

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 3

E, north, C, -1
C, east, D, -1
D, exit, x, +10

Episode 4

E, north, C, -1
C, east, A, -1
A, exit, x, -10

Output Values

	A	
B	C	D
	E	

Problems with Direct Evaluation

- What's good about direct evaluation?
 - It's easy to understand
 - It doesn't require any knowledge of T, R
 - It eventually computes the correct average values using just sample transitions
- What is bad about it?
 - It wastes information about state connections
 - Each state must be learned separately
 - So, it takes a long time to learn


Output Values

	-10 A	
+8 B	+4 C	+10 D
	-2 E	

Why Not Use Policy Evaluation?

- Simplified Bellman updates to calculate V for a fixed policy:
 - New V is expected one-step-look-ahead using current V
 - Unfortunately, need T and R

$$V_0^\pi(s) = 0$$



$$V_{i+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_i^\pi(s')]$$

- Key question: how can we do this update to V without knowing T and R ?
 - In other words, how to take a weighted average without knowing the weights?

Sample-Based Policy Evaluation?

- We want to improve our estimate of V by computing these averages:

$$V_{k+1}^{\pi}(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^{\pi}(s')]$$

- Idea: Take samples of outcomes s' (by doing the action!) and average

$$sample_1 = R(s, \pi(s), s'_1) + \gamma V_k^{\pi}(s'_1)$$

$$sample_2 = R(s, \pi(s), s'_2) + \gamma V_k^{\pi}(s'_2)$$

...


$$sample_n = R(s, \pi(s), s'_n) + \gamma V_k^{\pi}(s'_n)$$

$$V_{k+1}^{\pi}(s) \leftarrow \frac{1}{n} \sum_i sample_i$$

Temporal Difference Learning

Temporal-Difference Learning

- Big idea: learn from every experience!
 - Update $V(s)$ each time we experience (s,a,s',r)
 - Likely s' will contribute updates more often
- Temporal difference learning
 - Policy still fixed!
 - Move values toward value of whatever successor occurs: running average!



Sample of $V(s)$: $sample = R(s, \pi(s), s') + \gamma V^\pi(s')$

Update to $V(s)$: $V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + (\alpha)sample$

Same update: $V^\pi(s) \leftarrow V^\pi(s) + \alpha(sample - V^\pi(s))$

Example: Temporal Difference Learning

States

	A	
B	C	D
	E	

Assume: $\gamma = 1$, $\alpha = 1/2$

Observed Transitions

B, east, C, -2

	0	
0	0	8
	0	

C, east, D, -2


	0	
-1	0	8
	0	

	0	
-1	3	8
	0	

$$V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + \alpha [R(s, \pi(s), s') + \gamma V^\pi(s')]$$

Problems with TD Value Learning

- TD value learning is a model-free way to do policy evaluation
- However, if we want to turn values into a (new) policy, we're stuck:



$$\pi(s) = \arg \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Idea: learn Q-values directly
- Makes action selection model-free too!

Q-Learning

- Q-Learning: sample-based Q-value iteration
- Learn $Q^*(s,a)$ values
 - Receive a sample (s,a,s',r)
 - Consider your old estimate: $Q(s, a)$
 - Consider your new sample estimate:

$$Q^*(s, a) = \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma \max_{a'} Q^*(s', a') \right]$$

$$sample = R(s, a, s') + \gamma \max_{a'} Q(s', a')$$

- Incorporate the new estimate into a running average:

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + (\alpha) [sample]$$

Summary

- Today: Reinforcement Learning
 - Motivation
 - Model-base Learning
 - Model-free Learning
 - Passive RL
 - Direct Evaluation
 - Temporal Difference Learning
 - Q-Learning (if time)
- Next week
 - Q-Learning
 - Deep RL
 - Neural networks