

Agents that Plan Ahead: Search

Russell and Norvig: Chapter 3.1-3.4, 3.5-3.6

CSE 240: Winter 2023

Lecture 2

Agenda

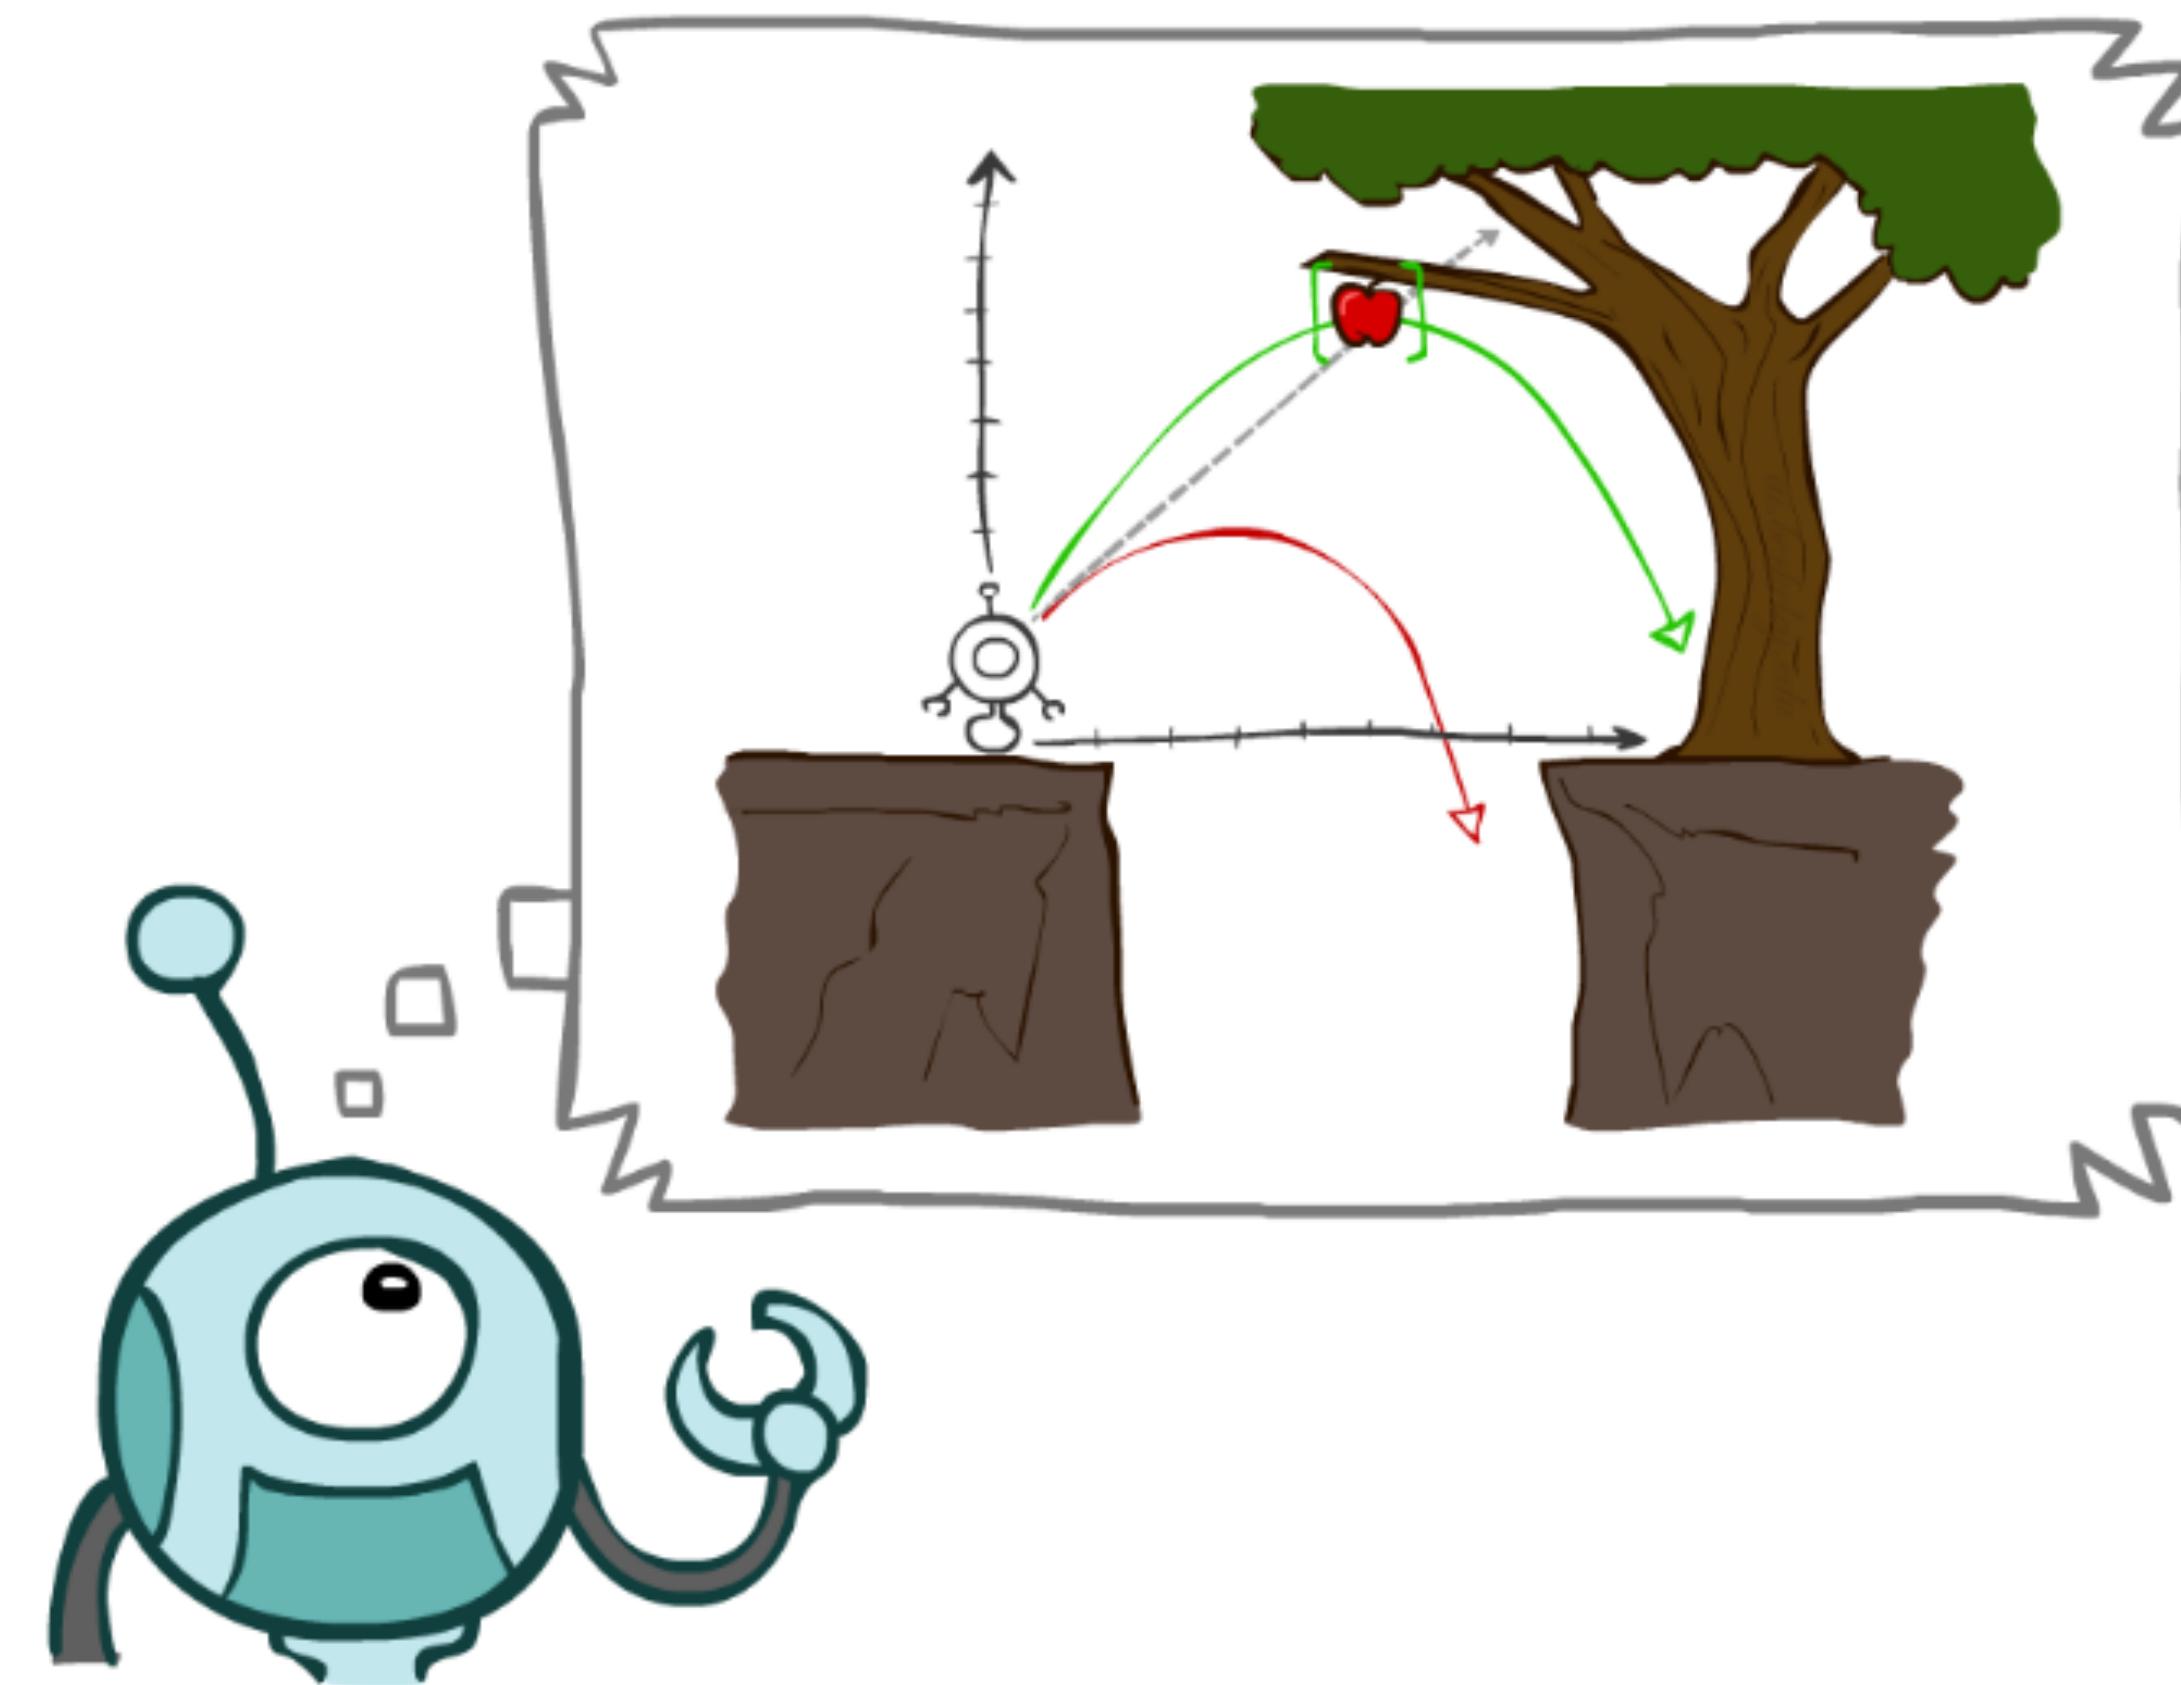
Agents that Plan Ahead

Search Problems

Uninformed Search Methods

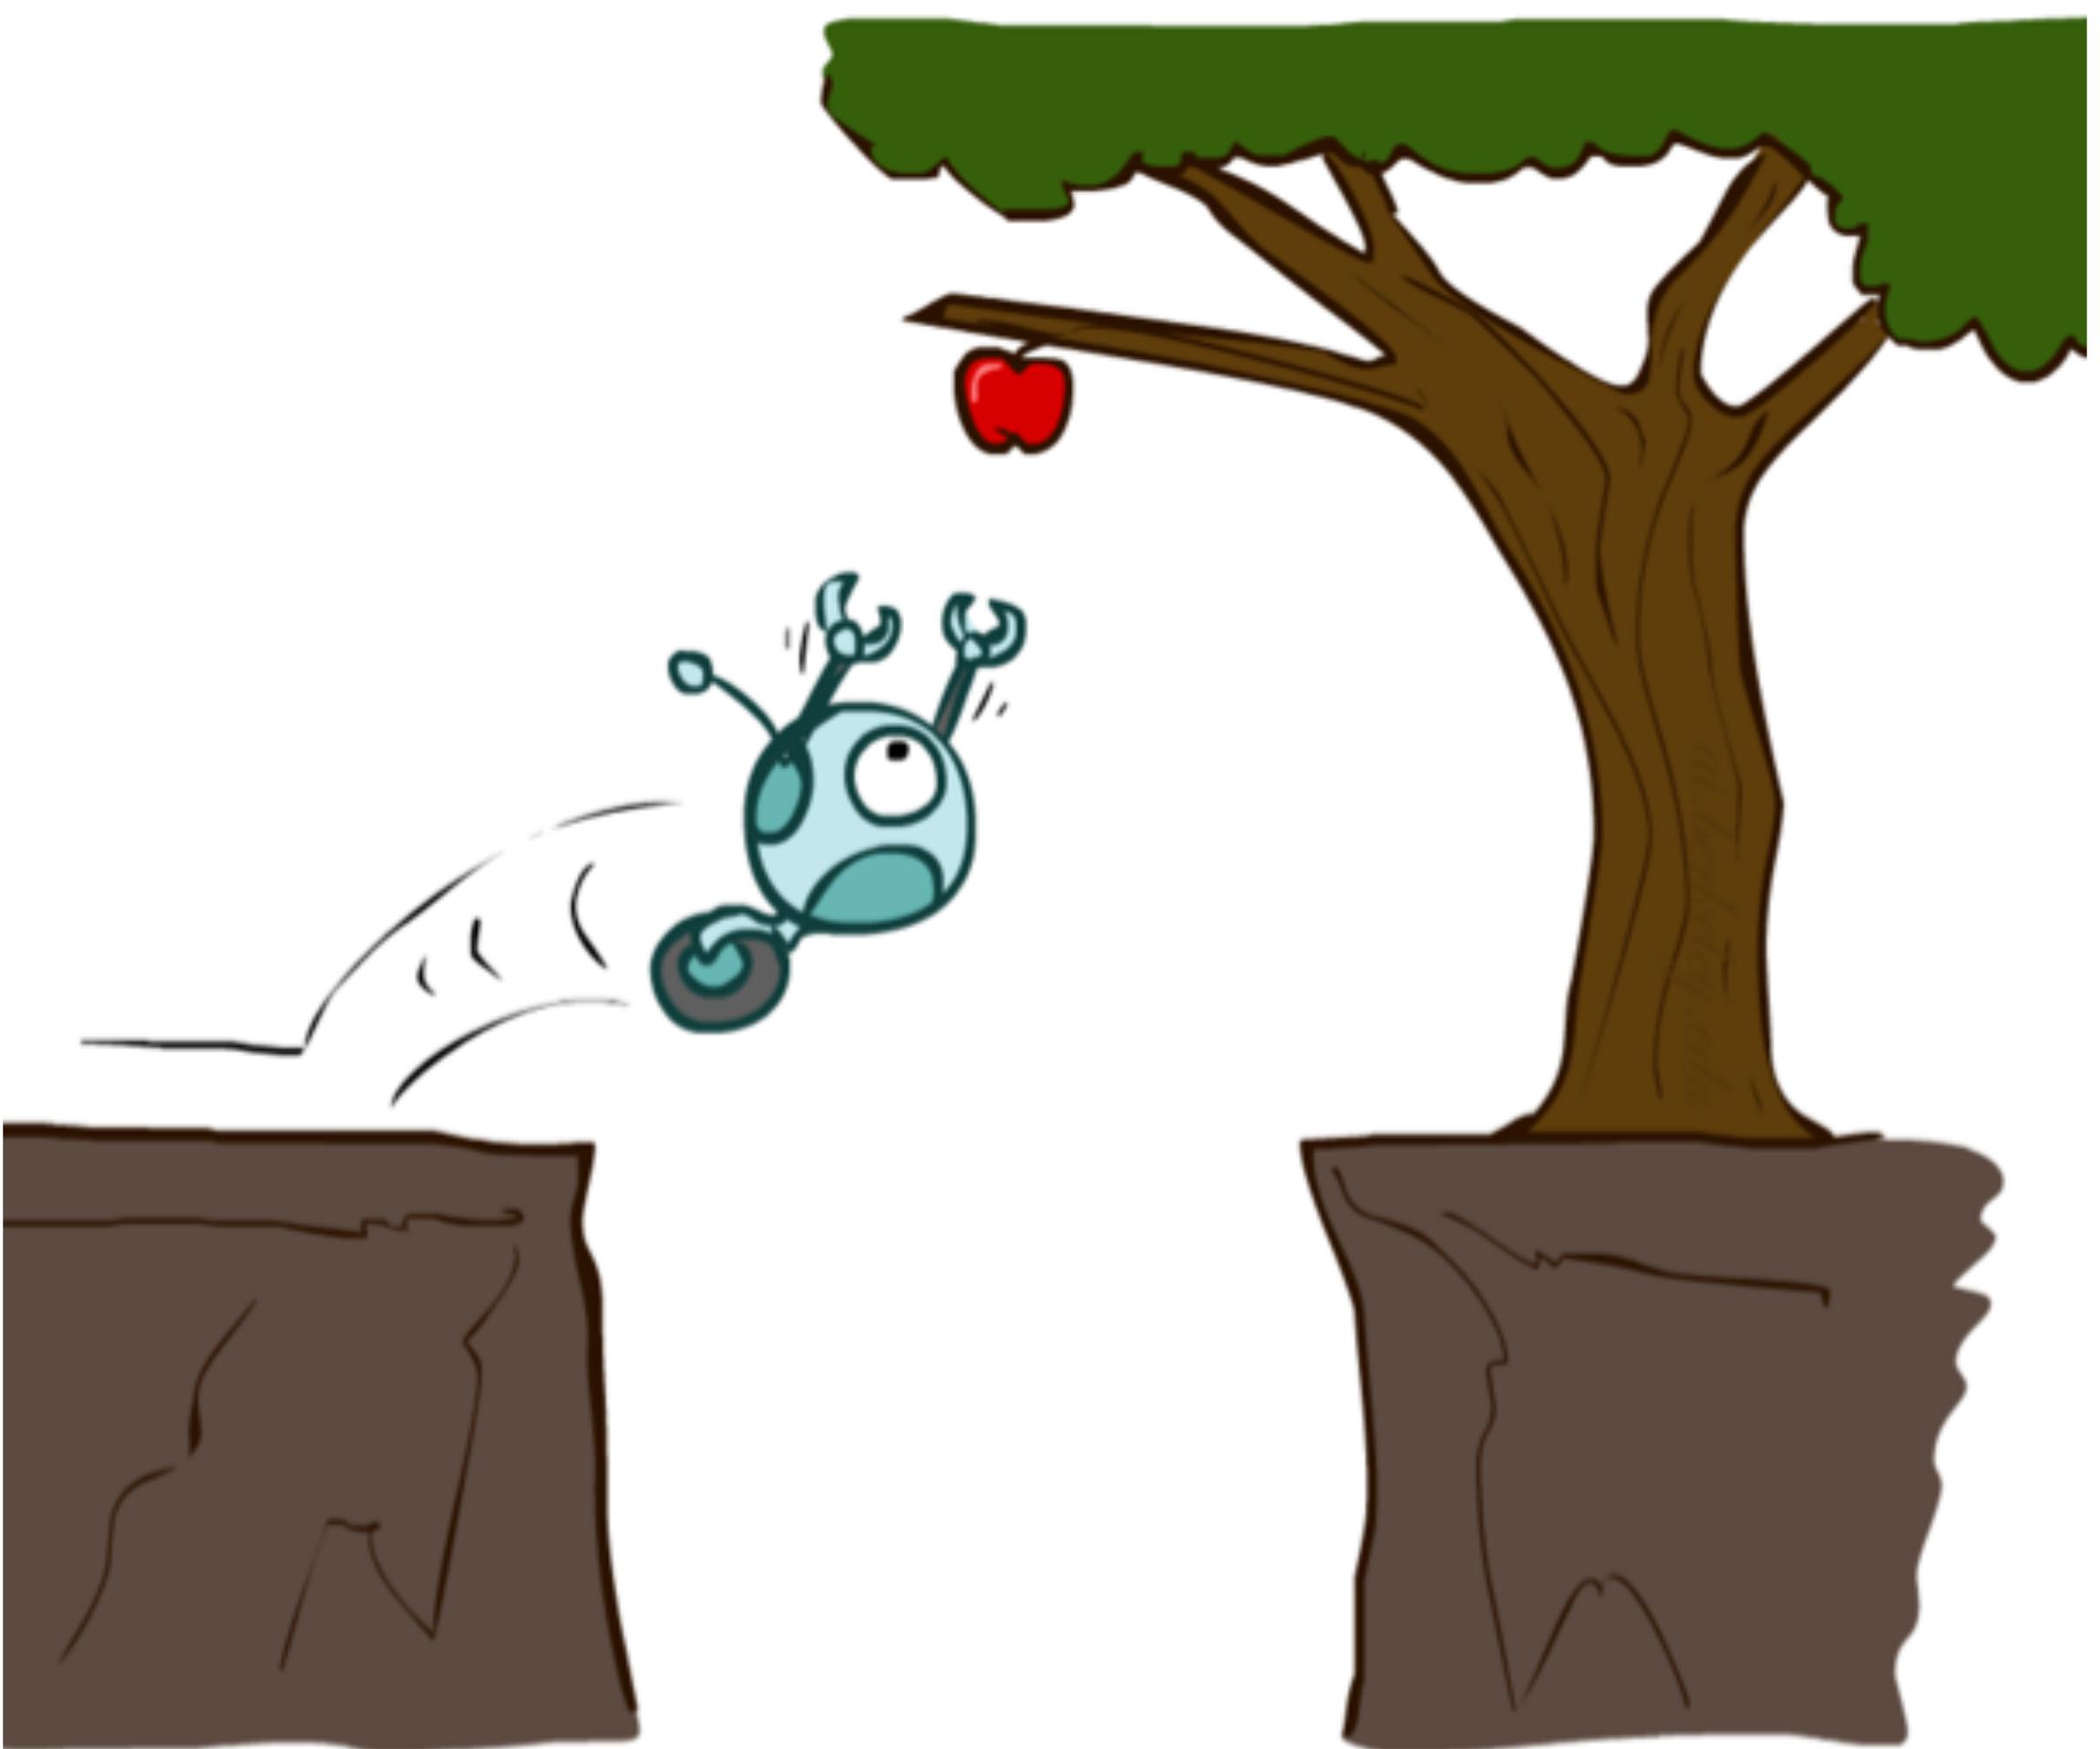
- Backtracking
- Depth-First Search
- Breadth-First Search
- Iterative deepening DFS
- Uniform-Cost Search

Agents that Plan Ahead



Reflex Agents

- Reflex agents:
 - Choose action based on current percept (and maybe memory)
 - May have memory or a model of the world's current state.
 - Do not consider the future consequences of their actions
 - Consider how the world “is”
 - *Can a reflex agent be rational?*



Reflex agent

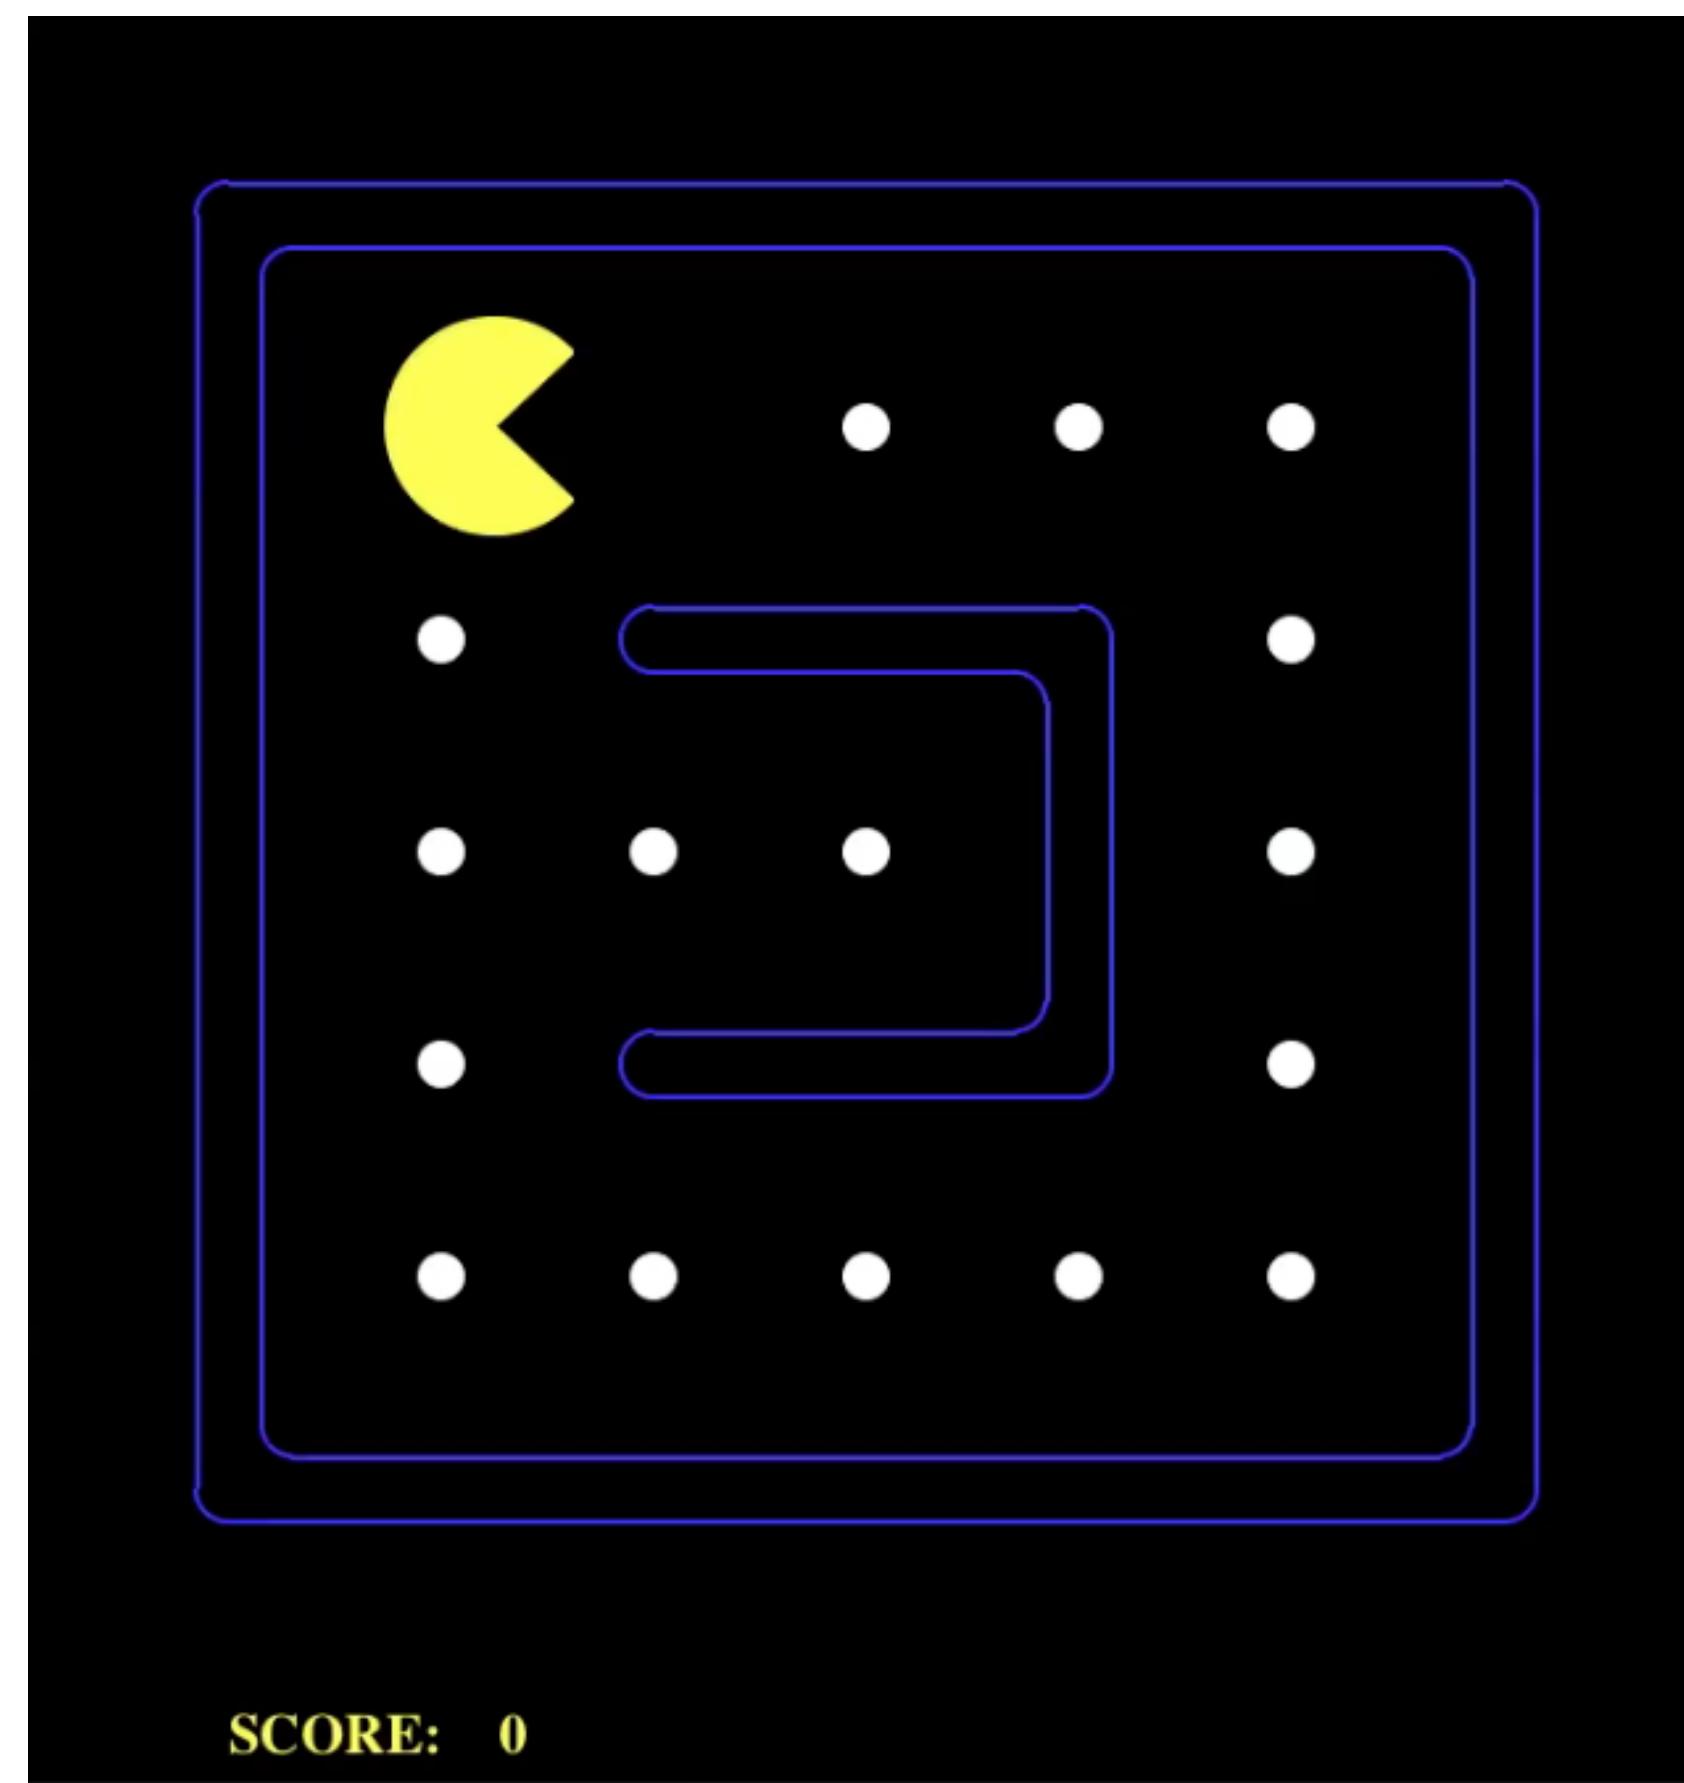
Strategy: Collect closest food item

How good is our agent ? Let's check it out...

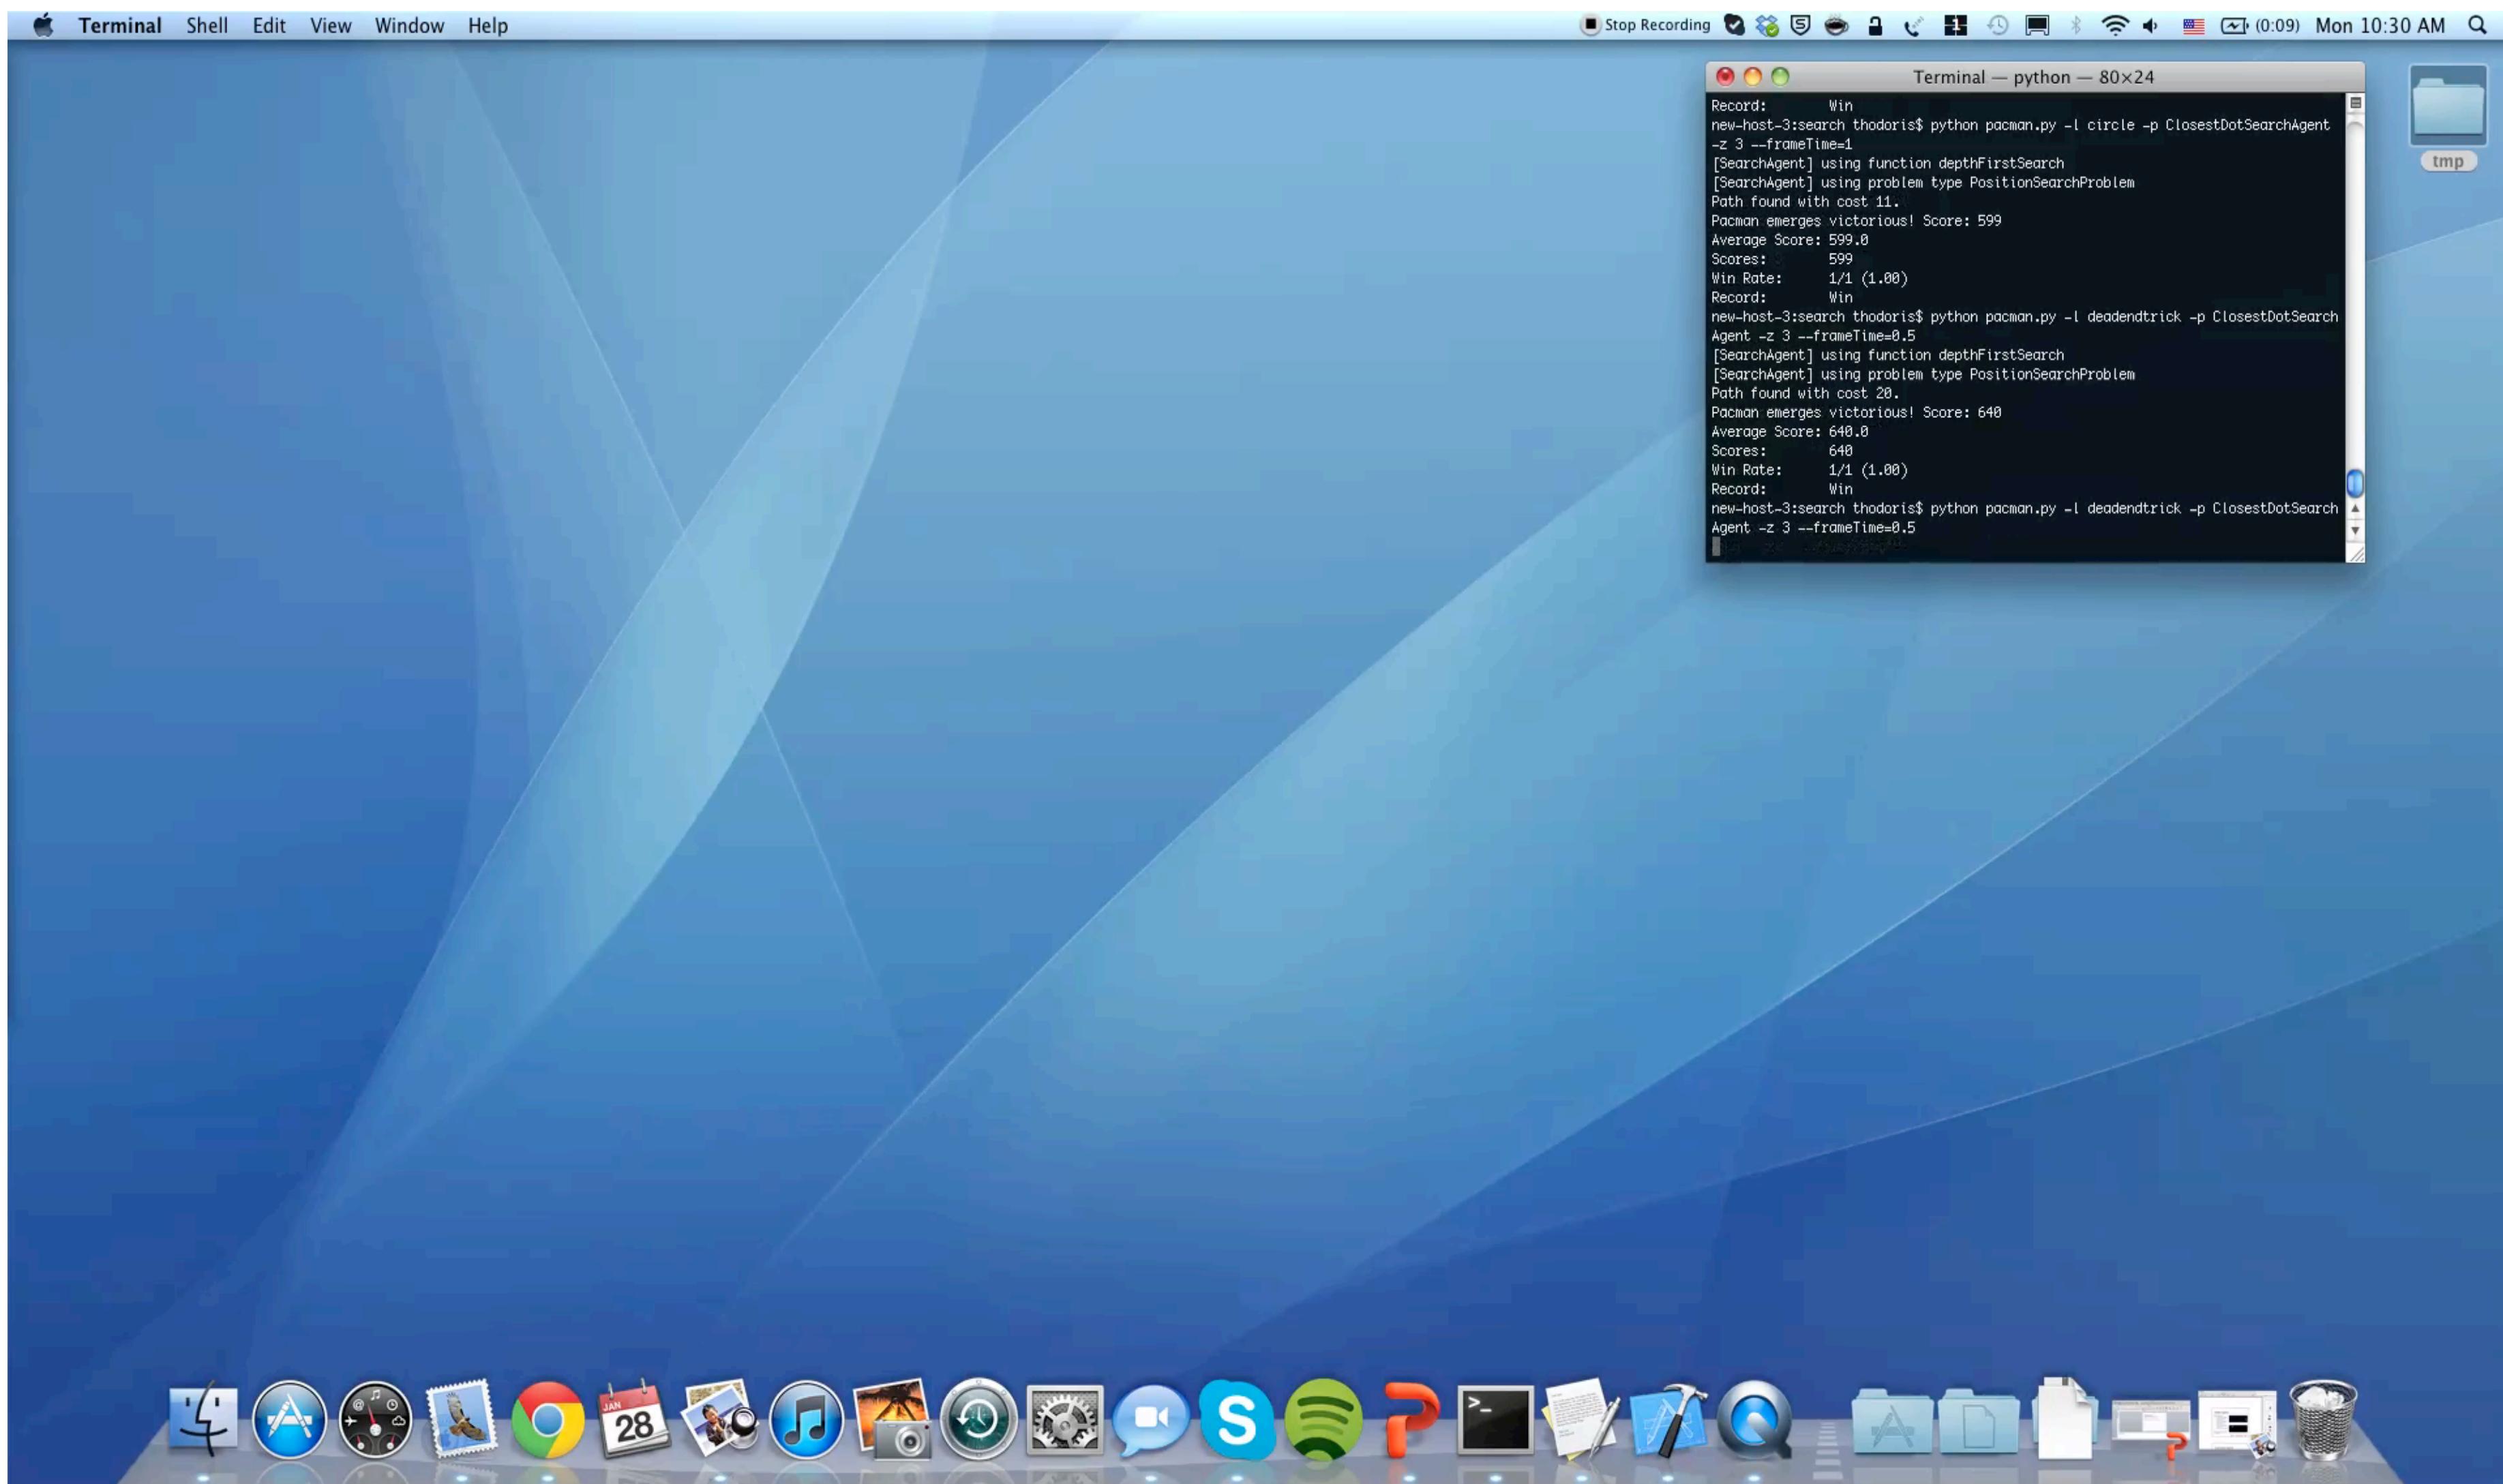


Reflex agent

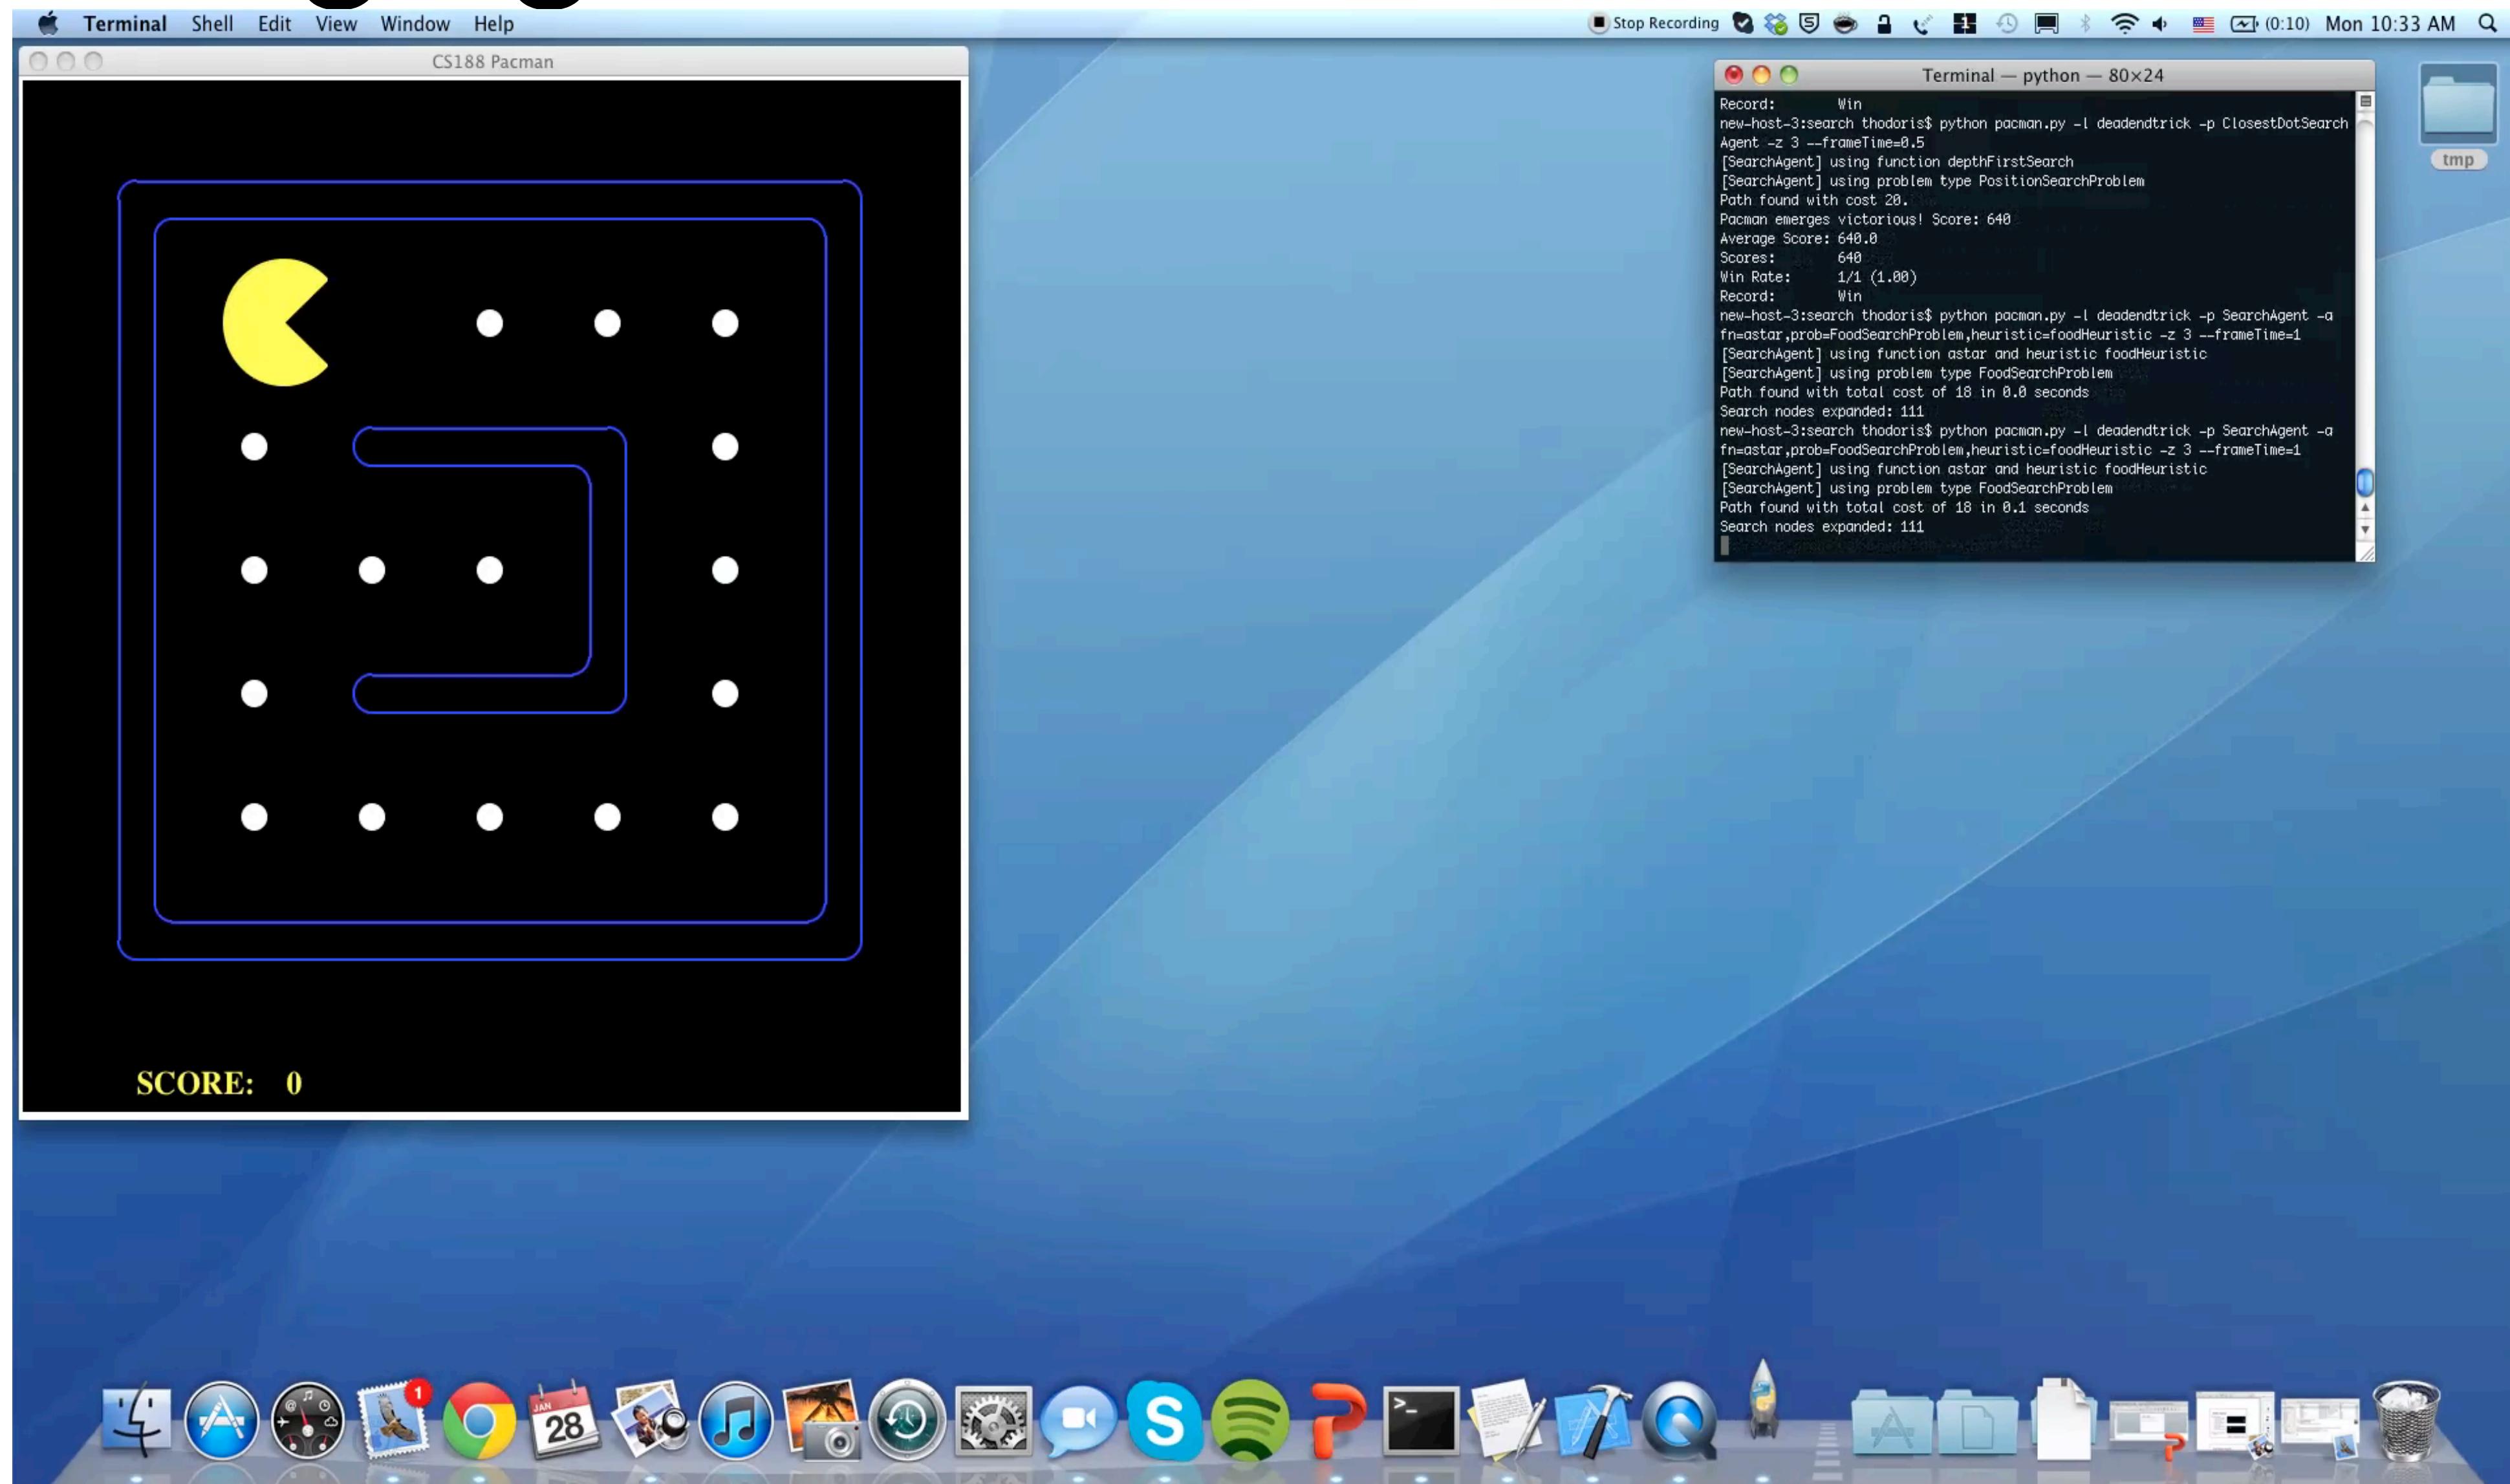
That was easy... What about harder instances?



Reflex agent

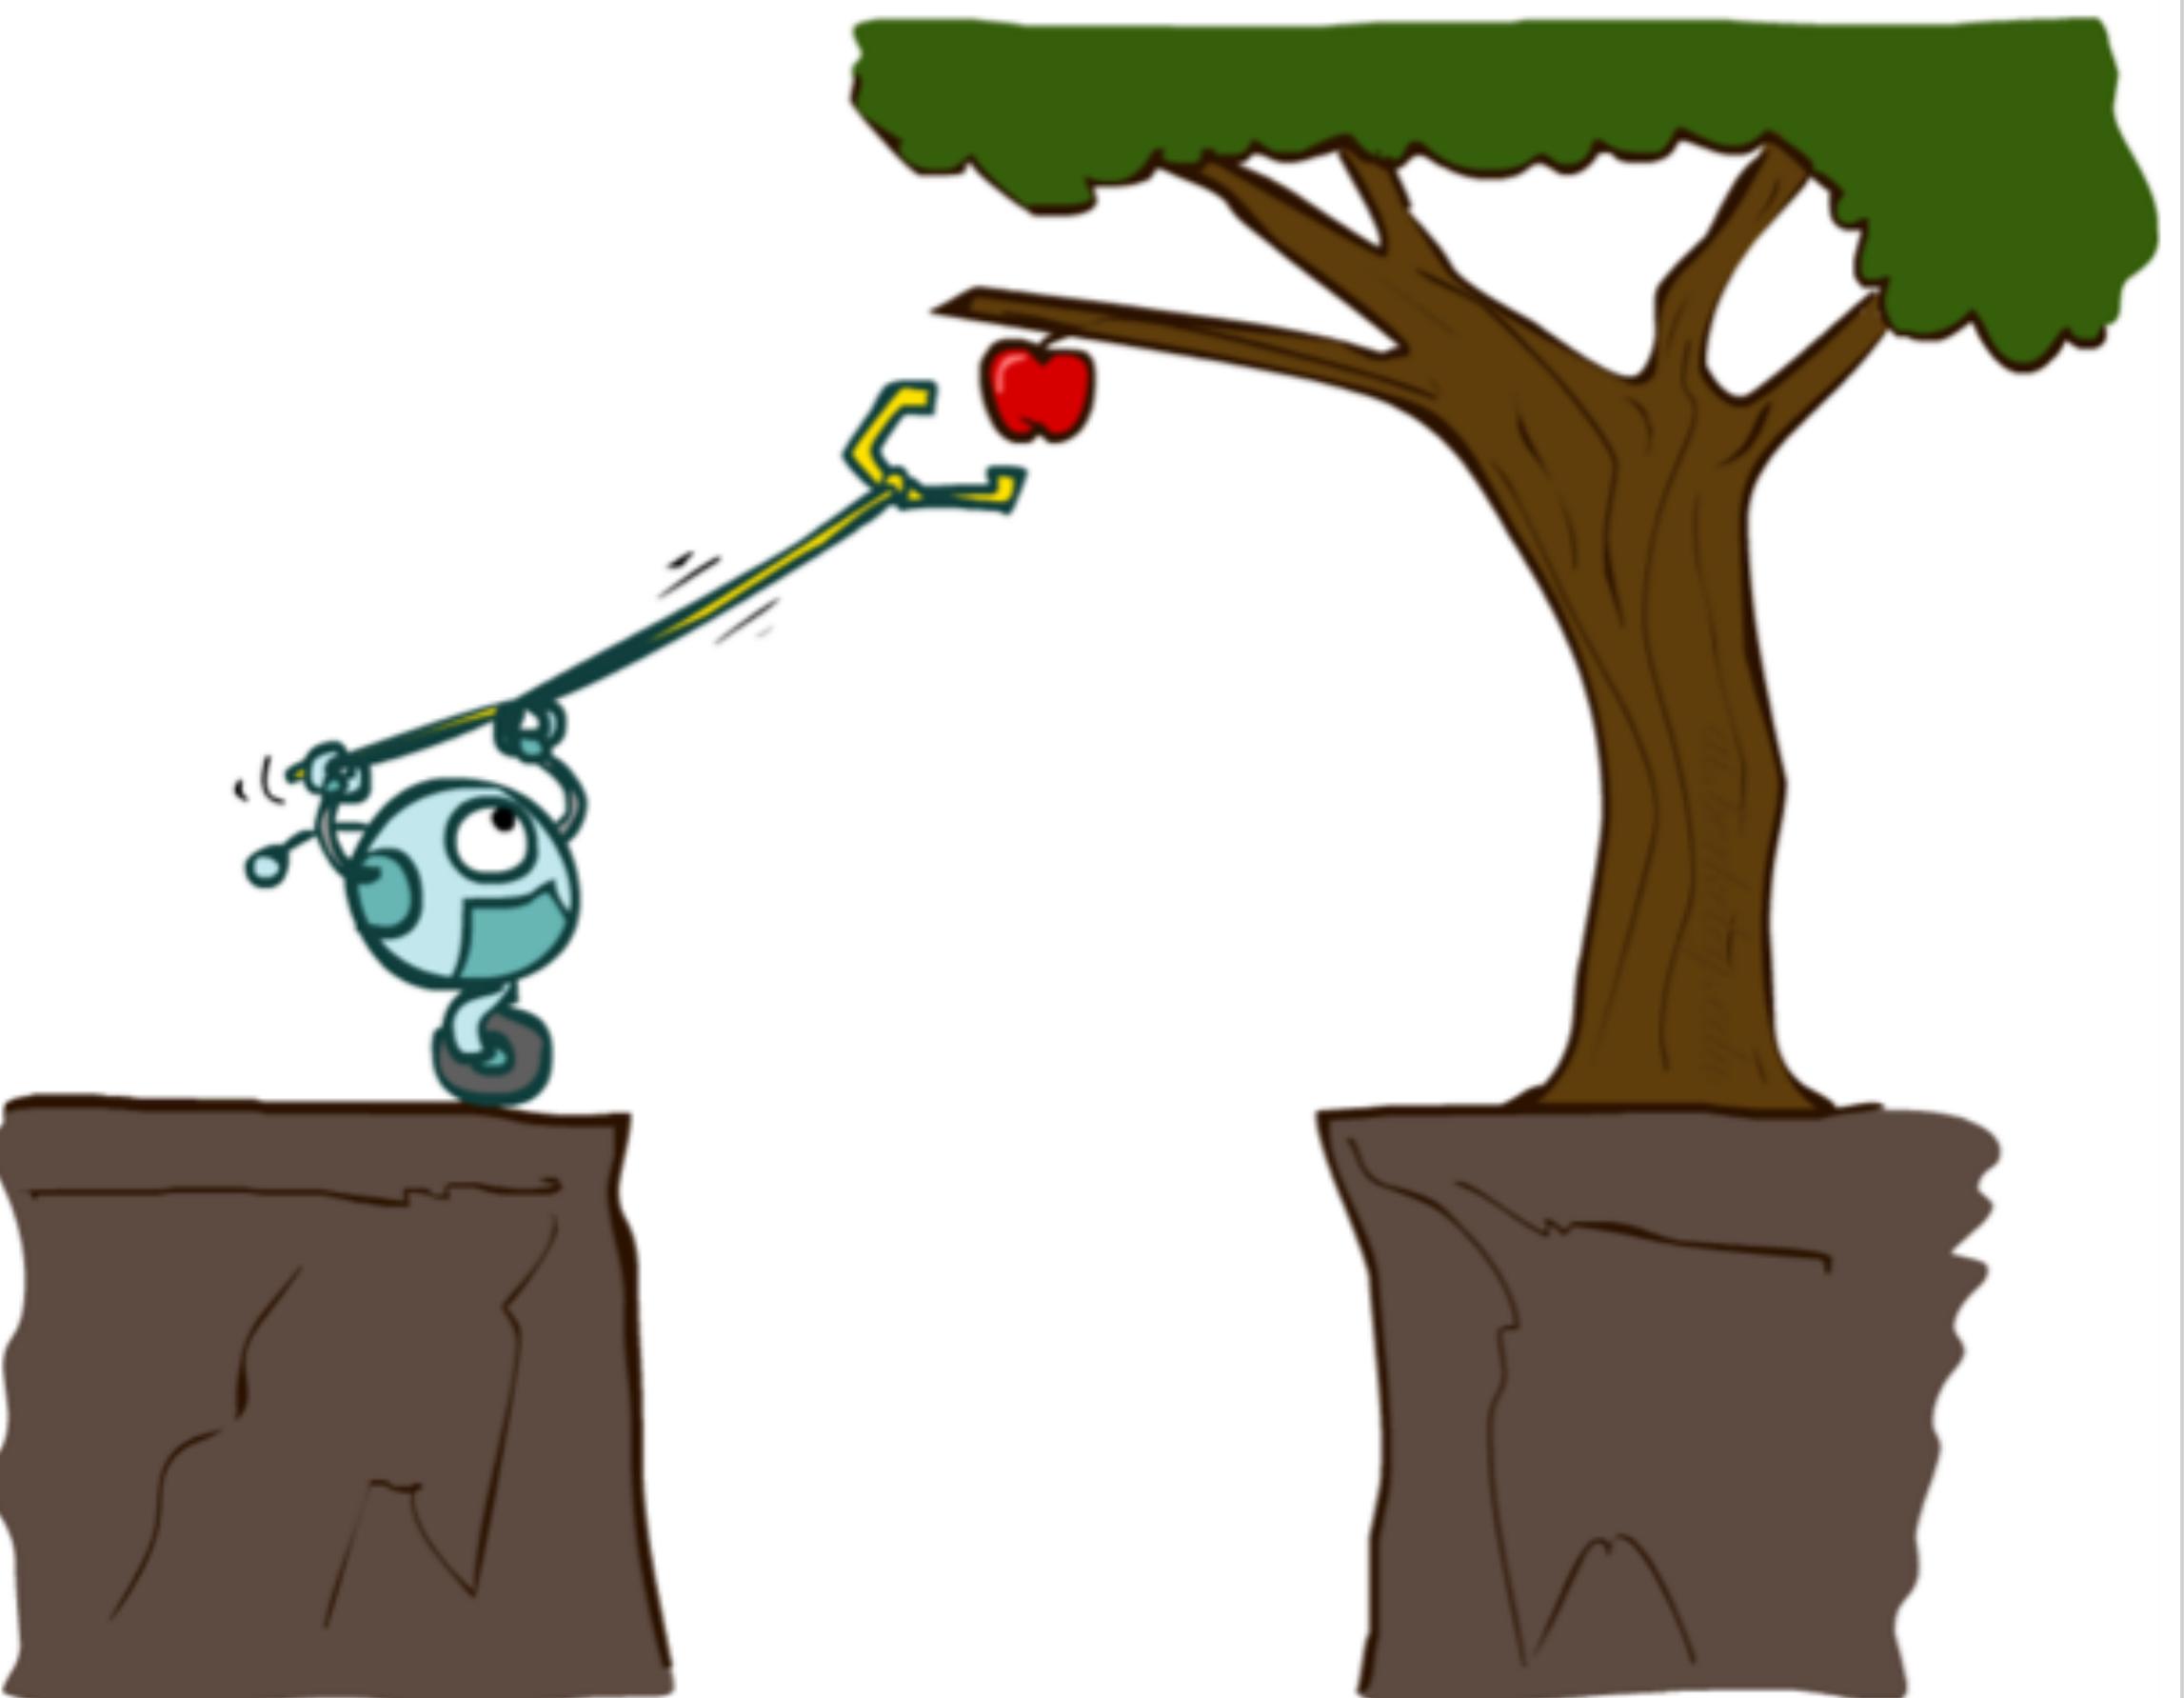


Planning agent



Planning Agents

- Planning agents:
 - Ask “what if”
 - Decisions based on (hypothesized) consequences of actions
 - Must have a model of how the world evolves in response to actions
 - Must formulate a goal (test)
 - Consider how the world WOULD BE
- Optimal vs. complete planning
- Planning vs. re-planning



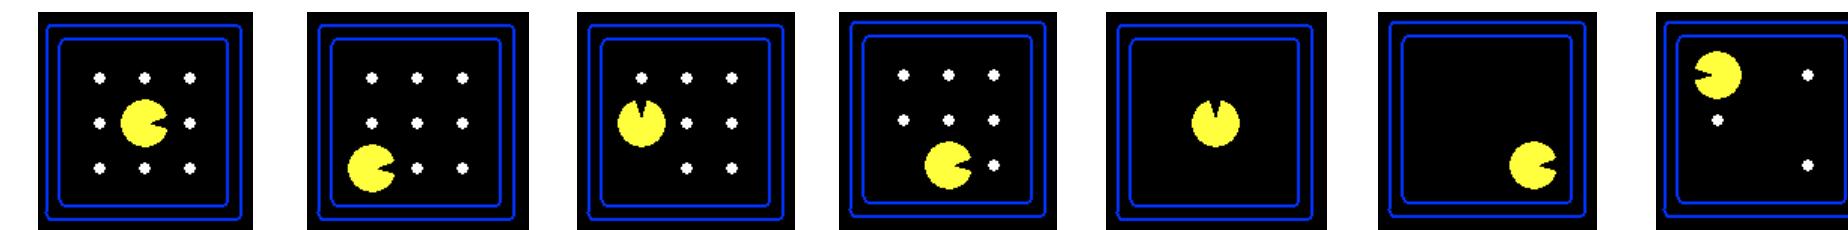
Search Problems



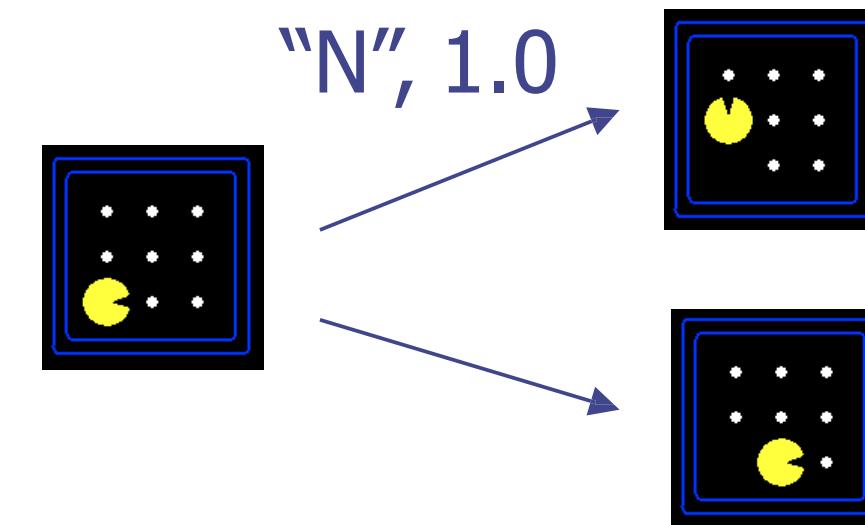
Search Problems

A **search problem** consists of:

- (1) A state space



- (2) A successor function



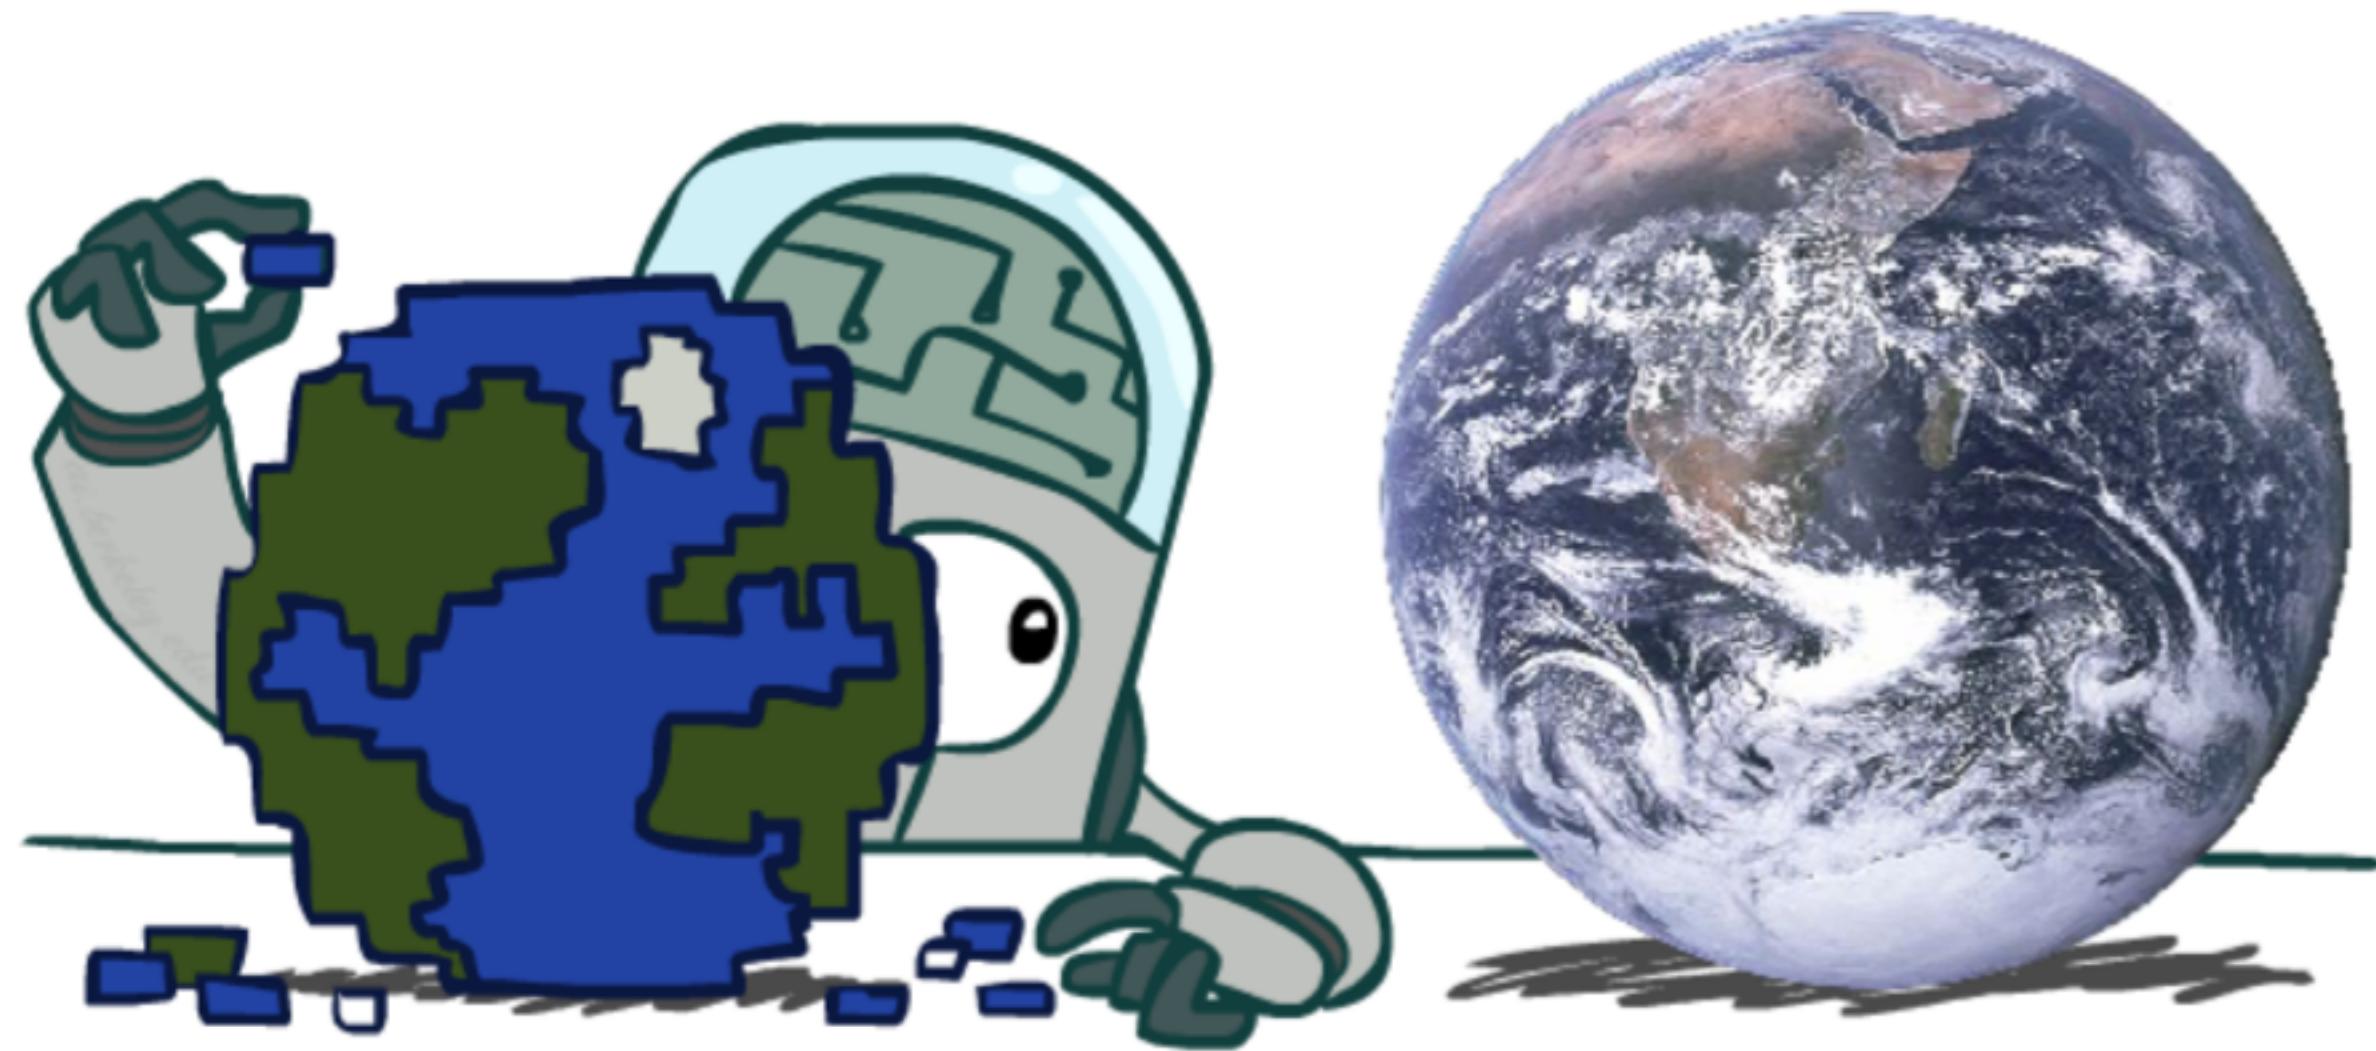
- (3) A start state and (4) a goal test

A solution is a sequence of actions (a plan) which transforms the start state to a goal state

Formal Definition of Search Problems

- s_{start} : starting state
- Actions(s): possible actions
- Cost(s, a): action cost
- Succ(s, a): successor state
- IsGoal(s): reached the goal state?

Search Problems Are Models



Example: Traveling in Romania

State space:

- Cities

Successor function:

- Go to adj city with cost = dist

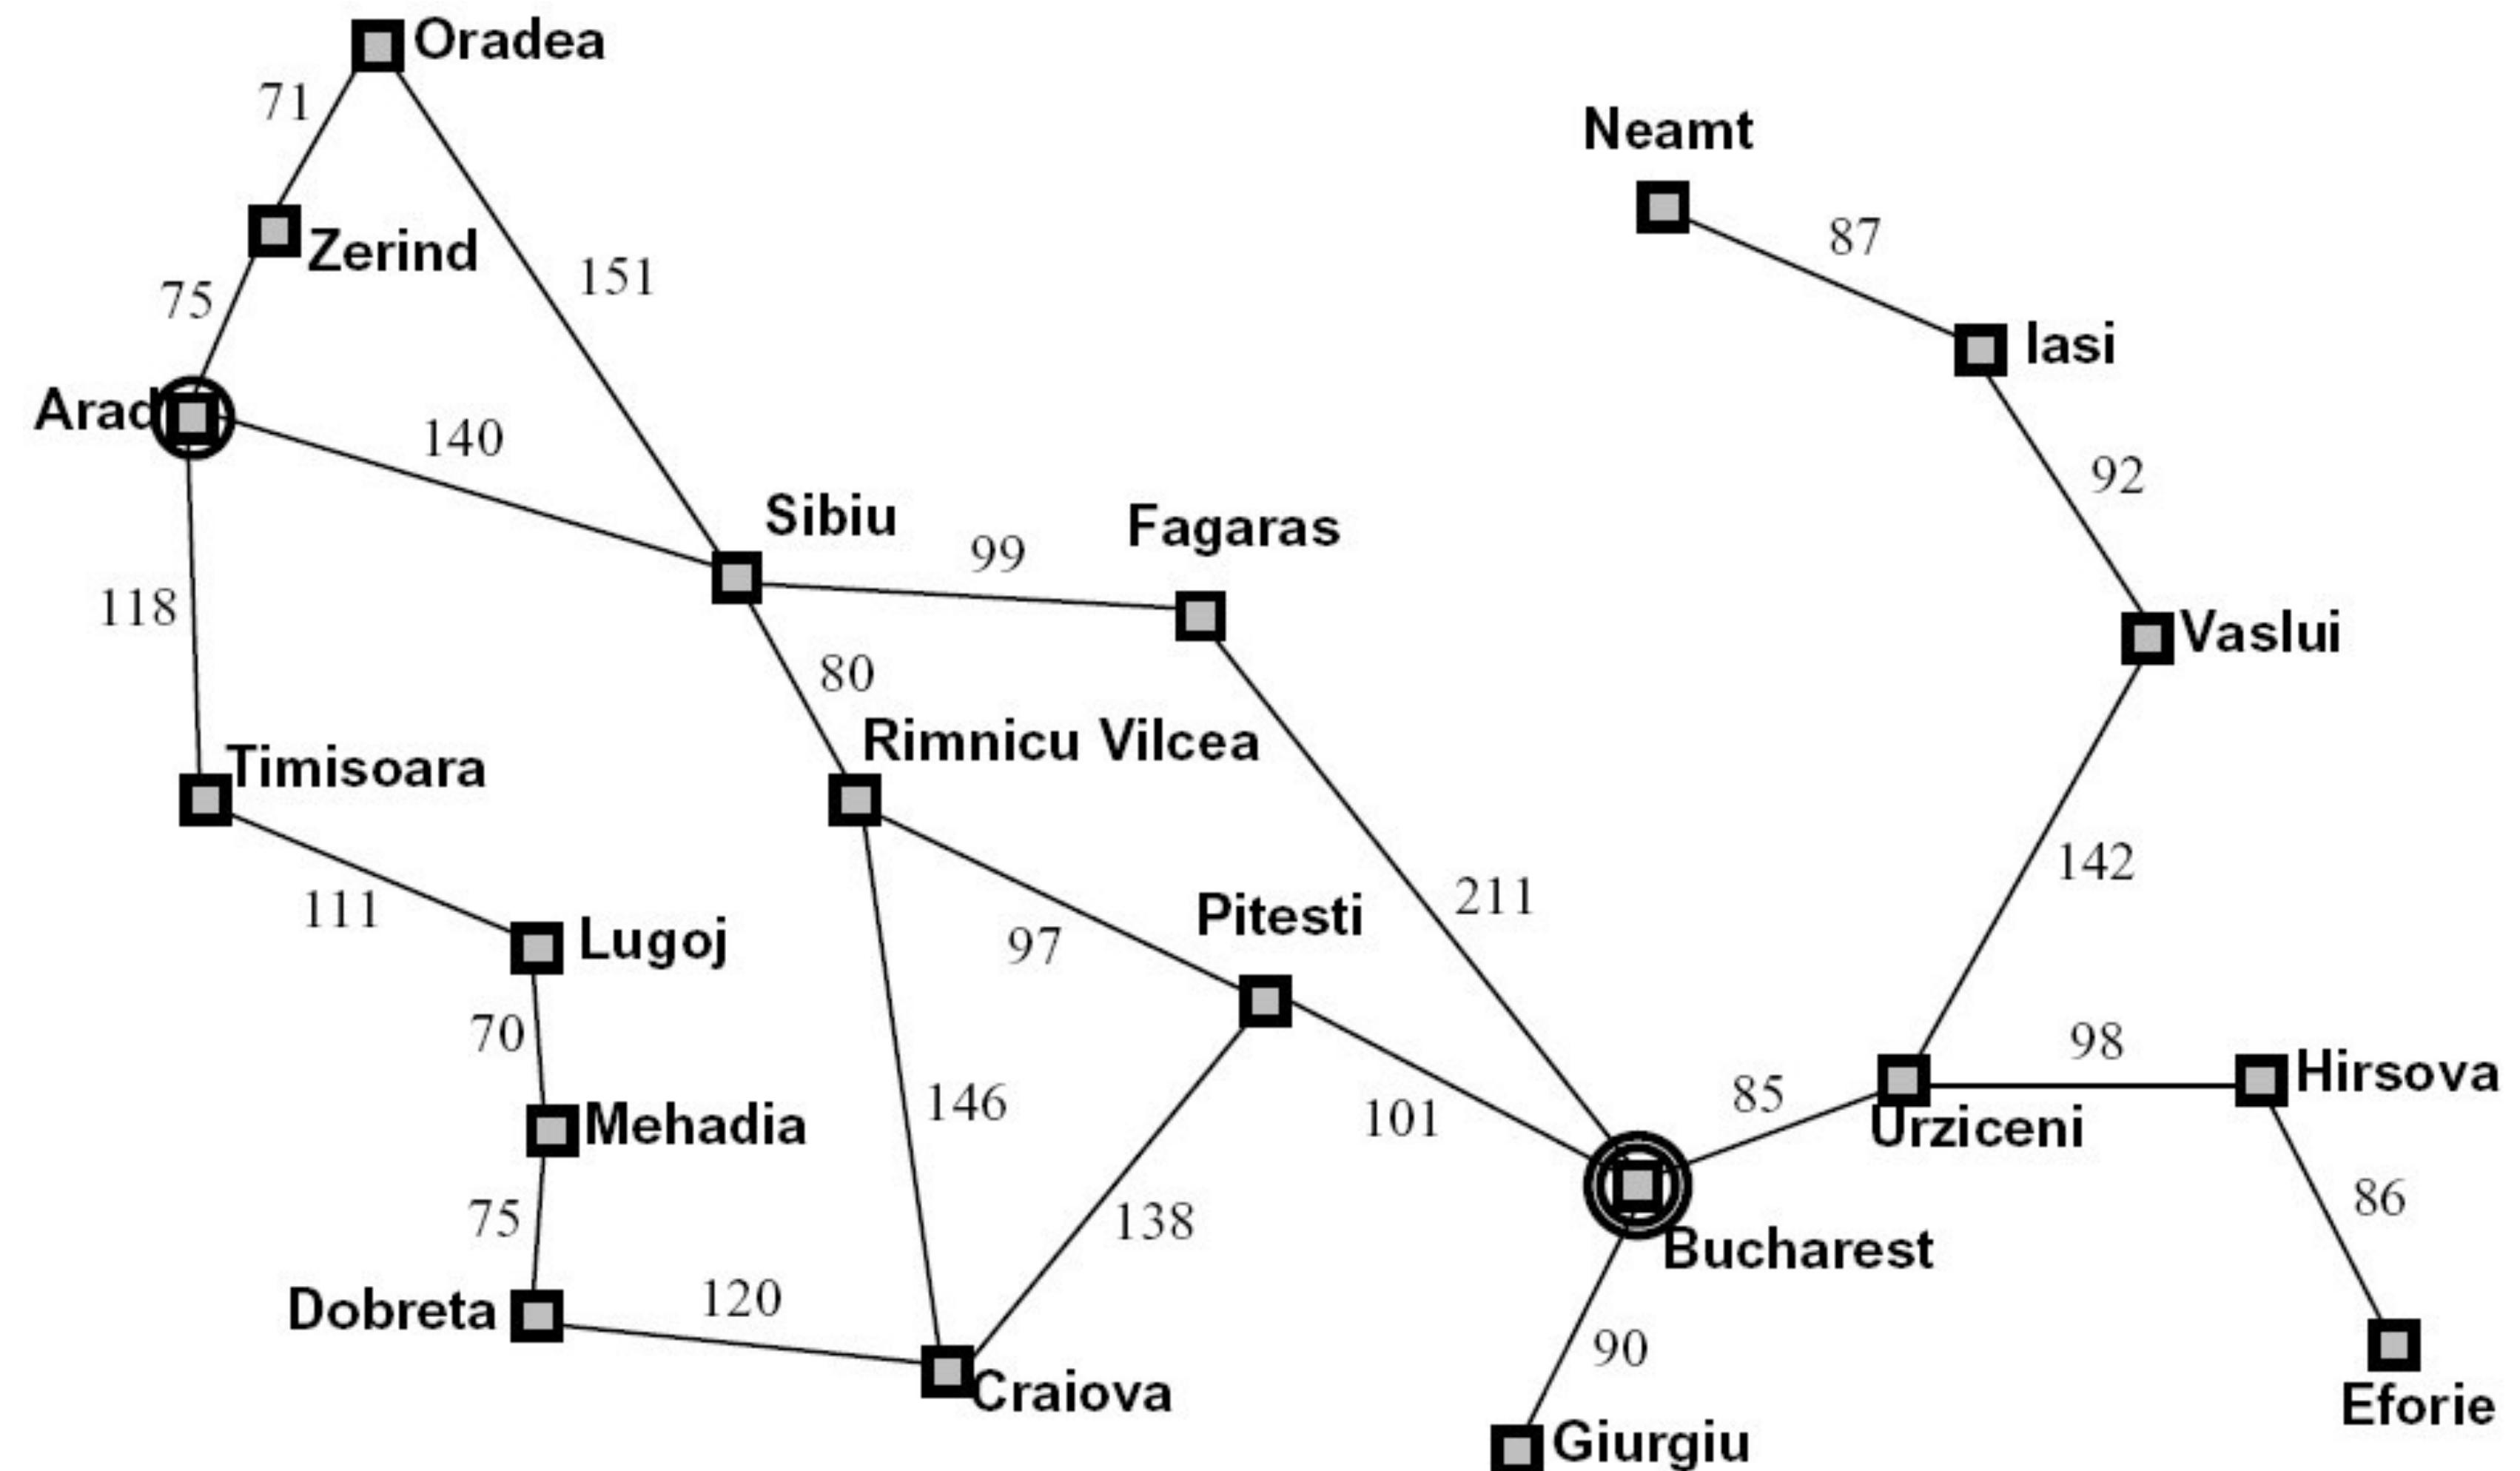
Start state:

- Arad

Goal test:

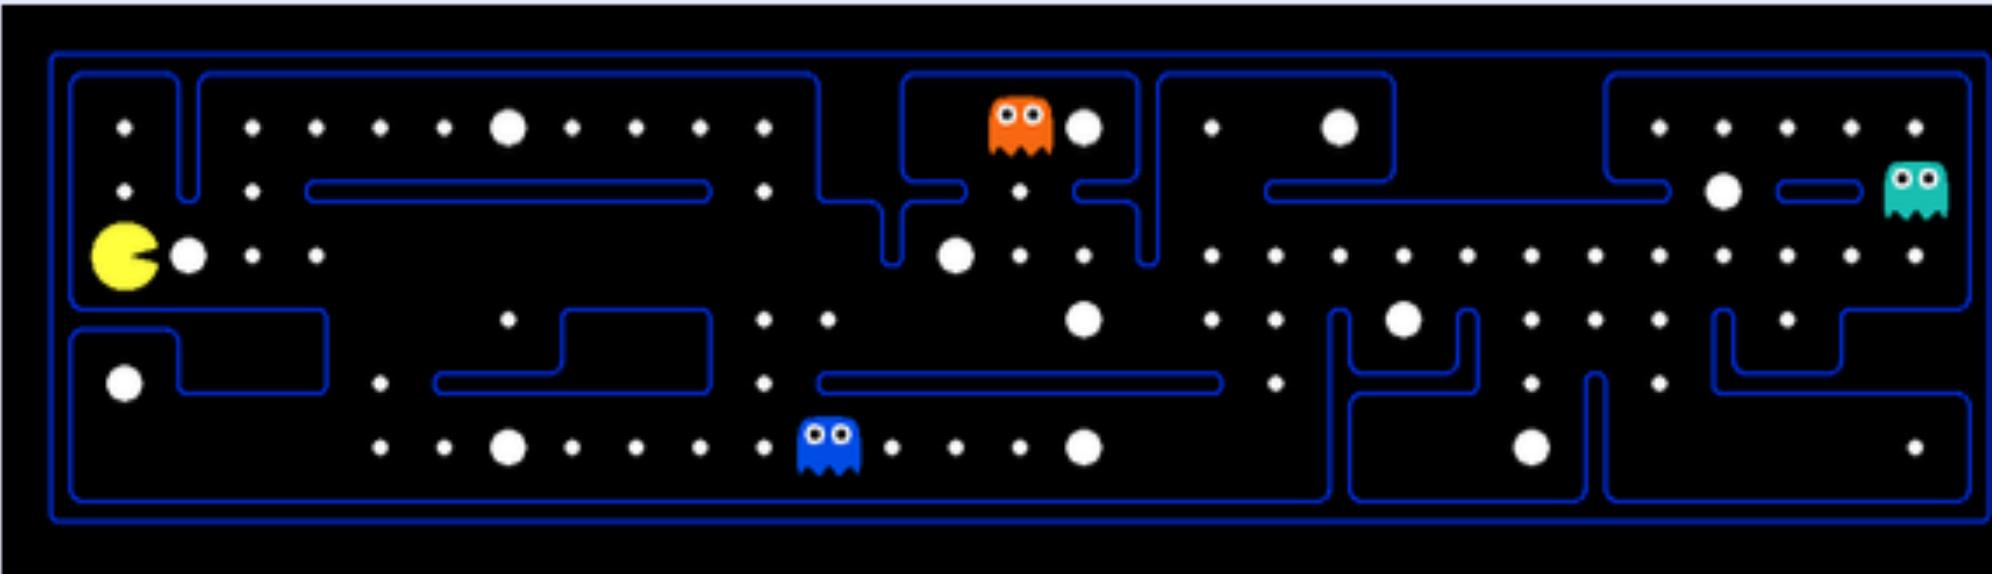
- Is state == Bucharest?

Solution?



What's in a State Space?

The **world state** includes every last detail of the environment

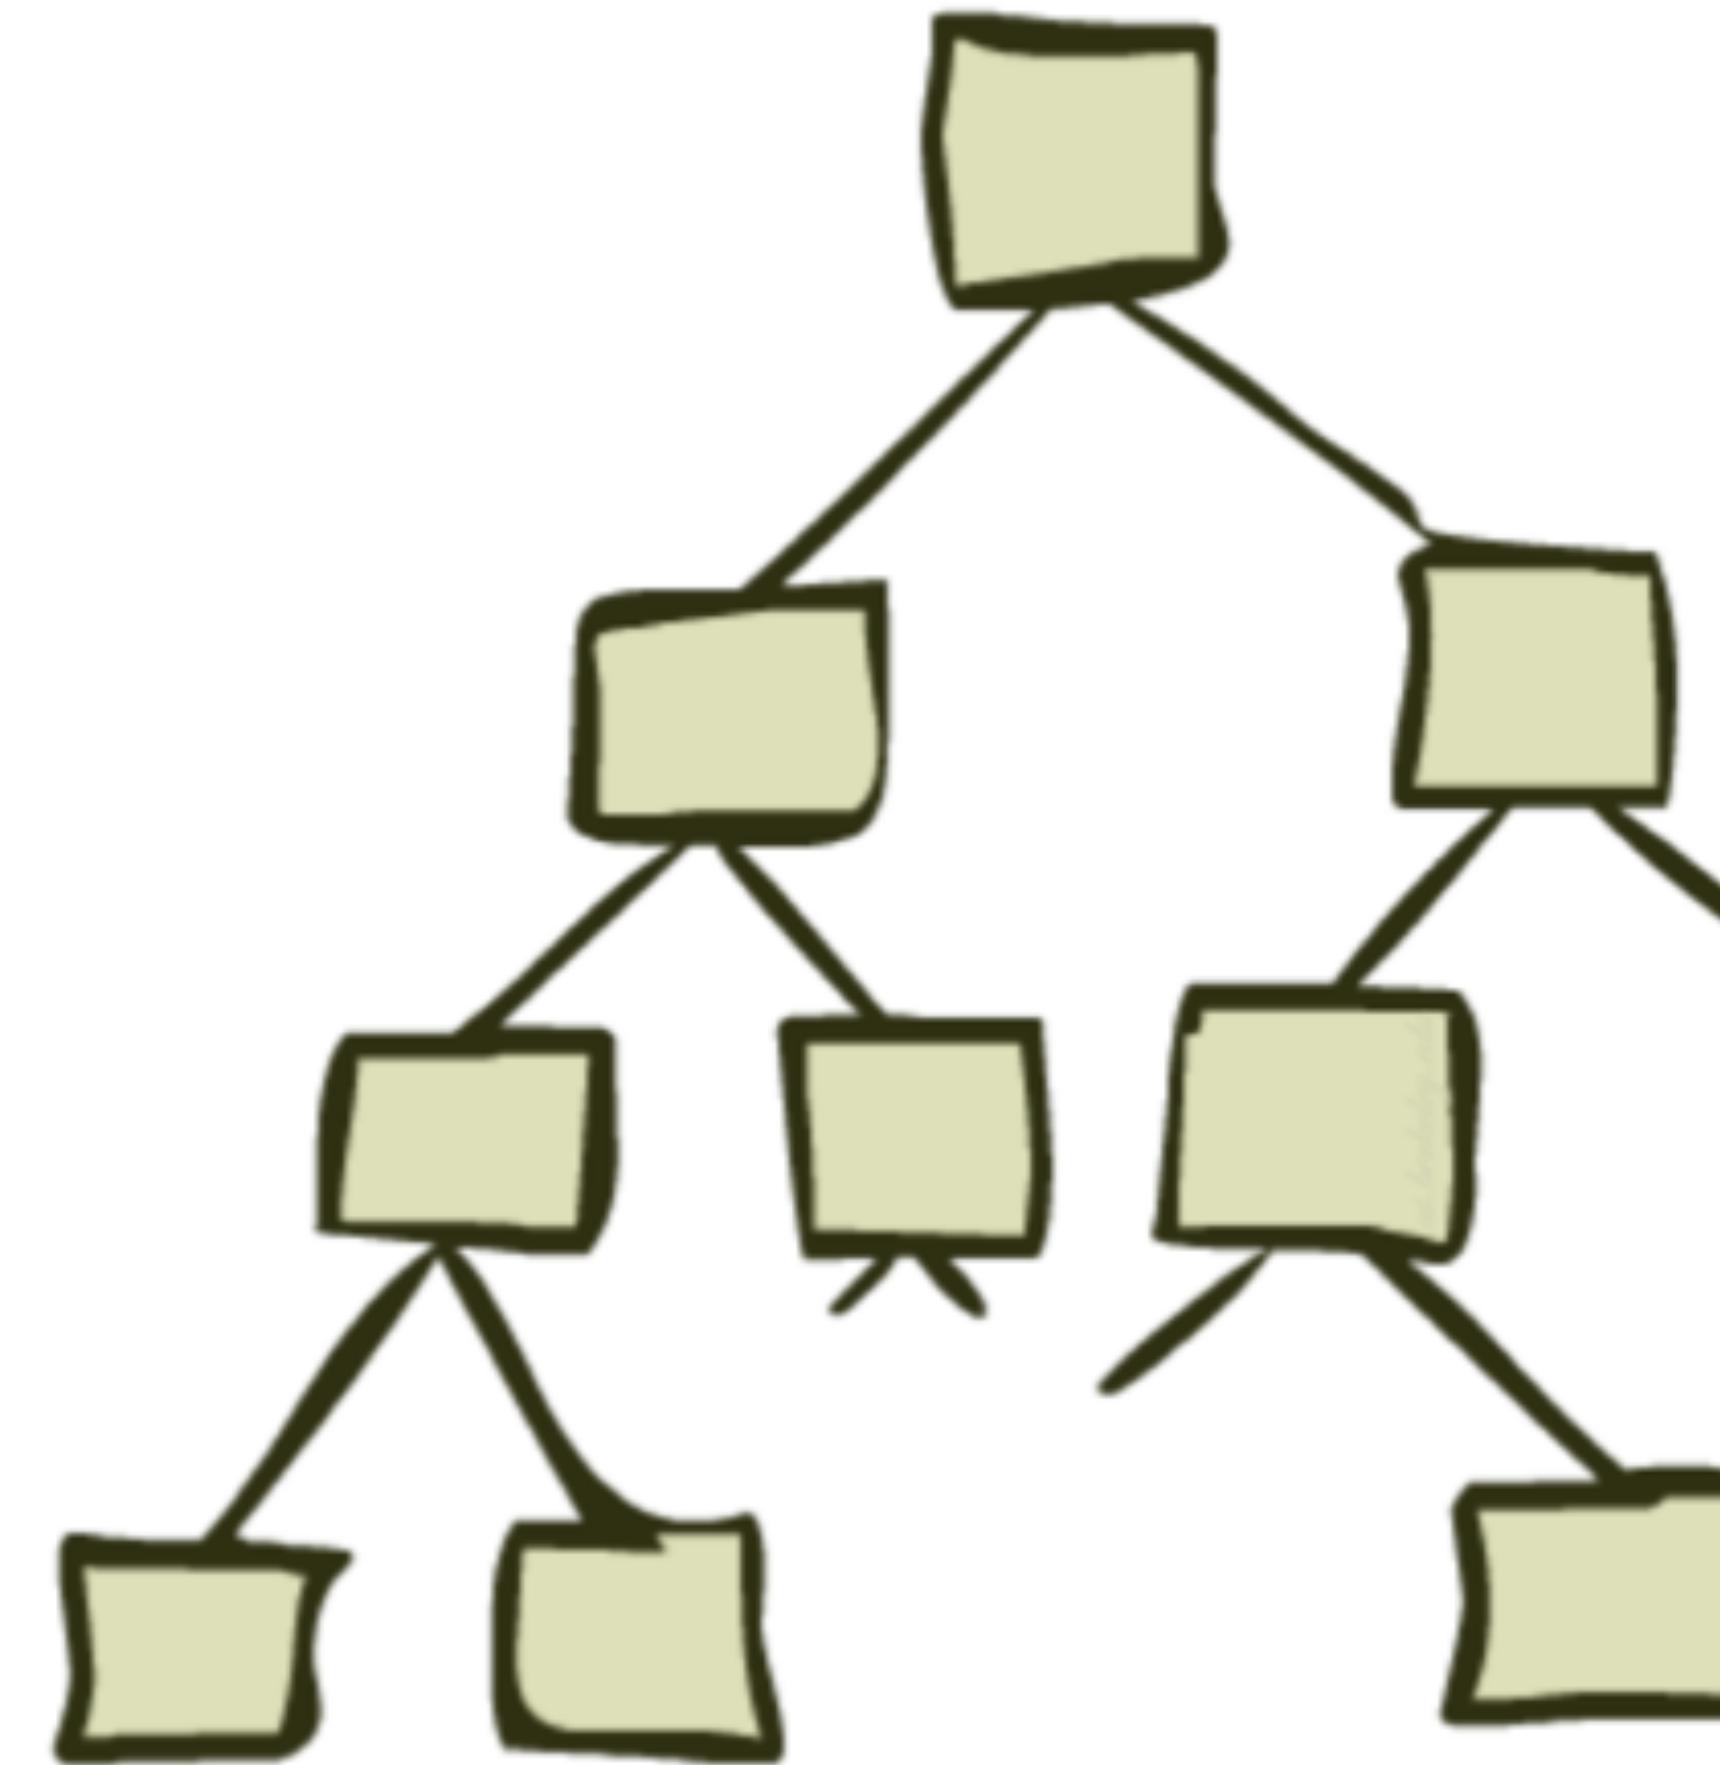


A **search state** keeps only the details needed for planning (abstraction)

- Problem: Finding a path
 - States: (x,y) location
 - Actions: N, S, E, W
 - Successor: update location only
 - Goal test: is (x,y)=END

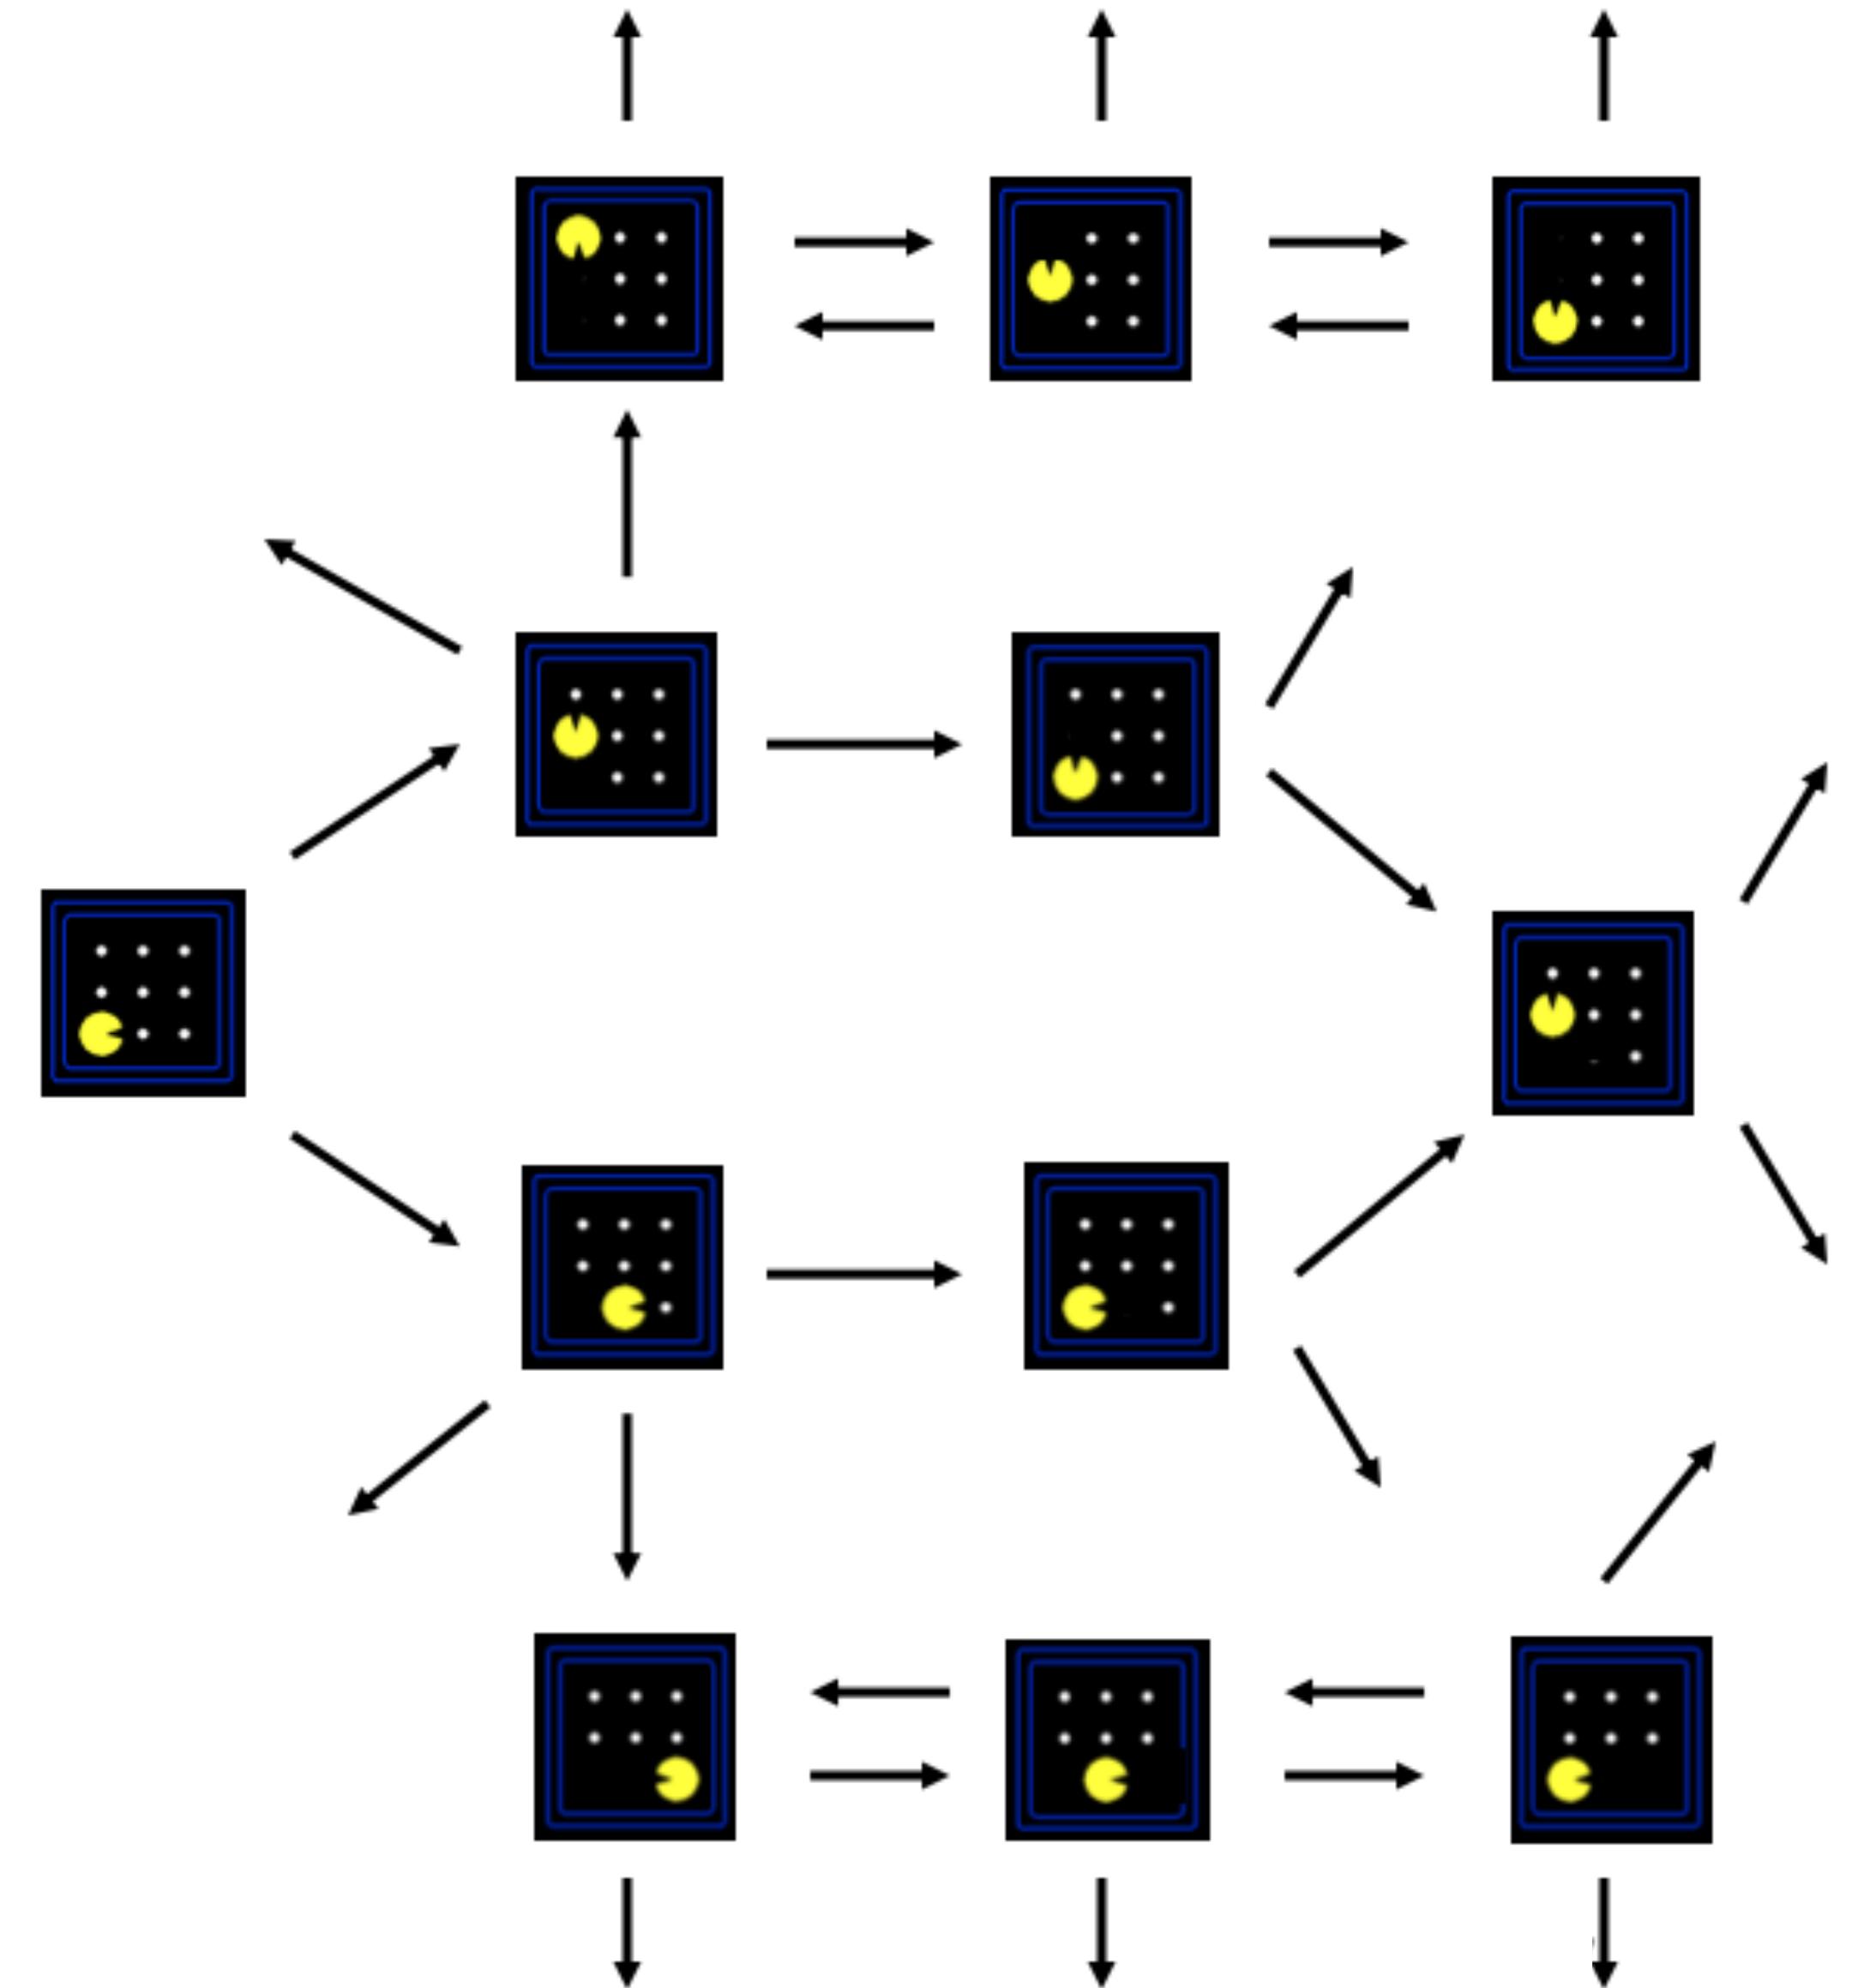
- Problem: Eat-All-Dots
 - States: {(x,y), dot Booleans}
 - Actions: N, S, E, W
 - Successor: update location and possibly a dot Boolean
 - Goal test: dots all false

State Space Graphs and Search Trees

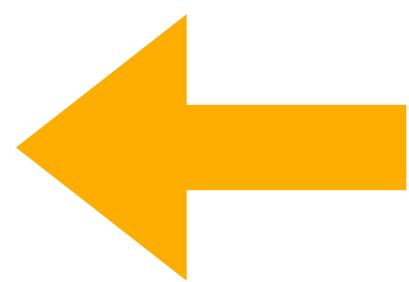
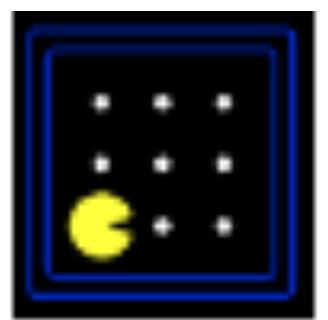


State Space Graphs

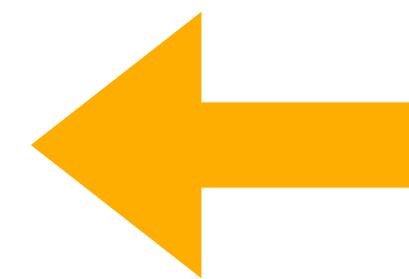
- State space graph: A mathematical representation of a search problem
 - Nodes are (abstracted) world configurations
 - Arcs represent successors (action results)
 - The goal test is a set of goal node(s)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



Search Trees



This is now /start

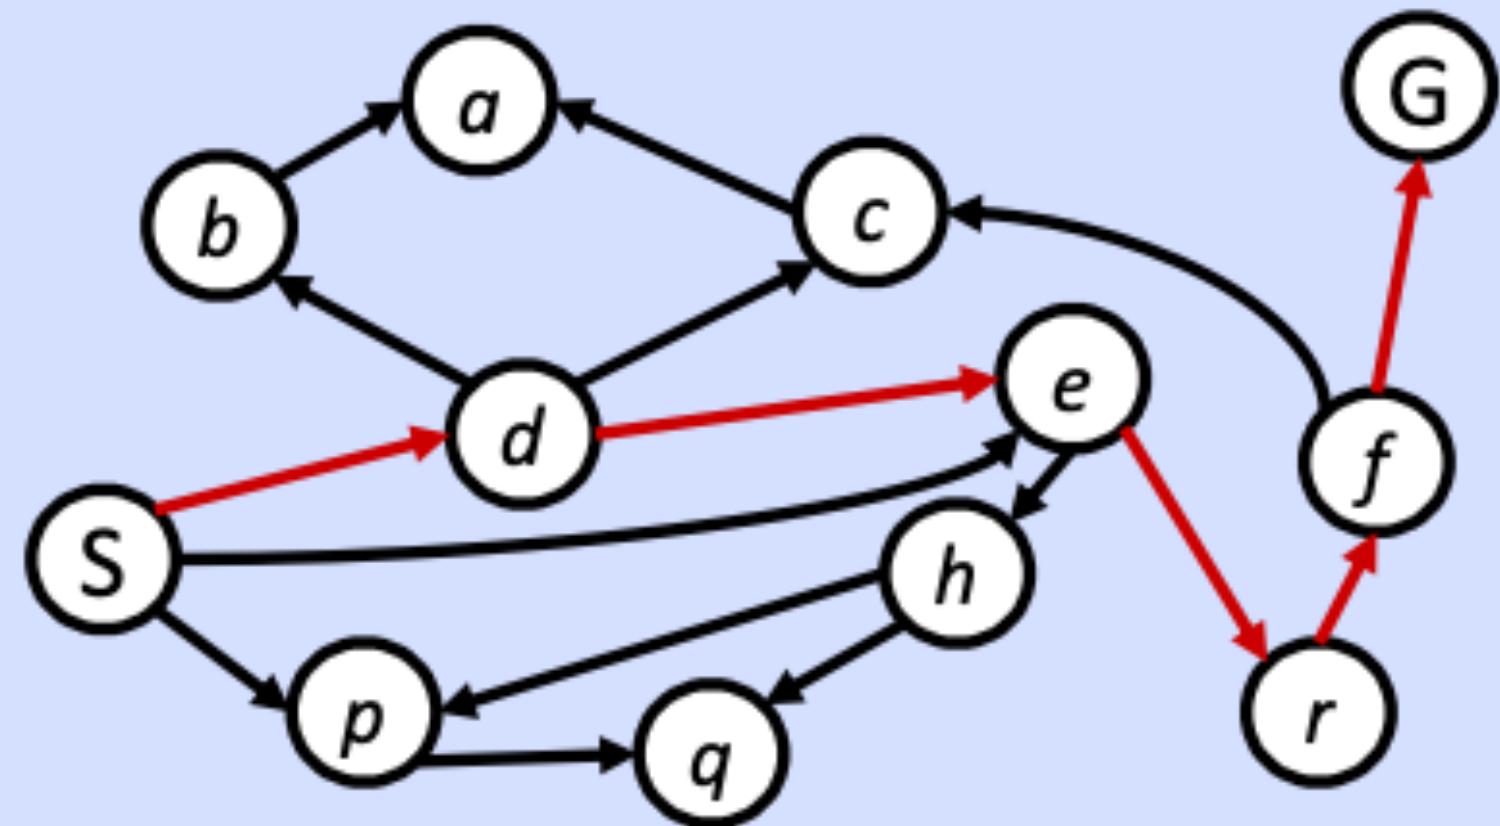


Possible futures

- A search tree:
 - This is a “what if” tree of plans and outcomes
 - Start state at the root node
 - Children correspond to successors
 - Nodes contain states, paths correspond to PLANS to those states
 - For most problems, we can never actually build the whole tree

State Space Graph vs. Search Trees

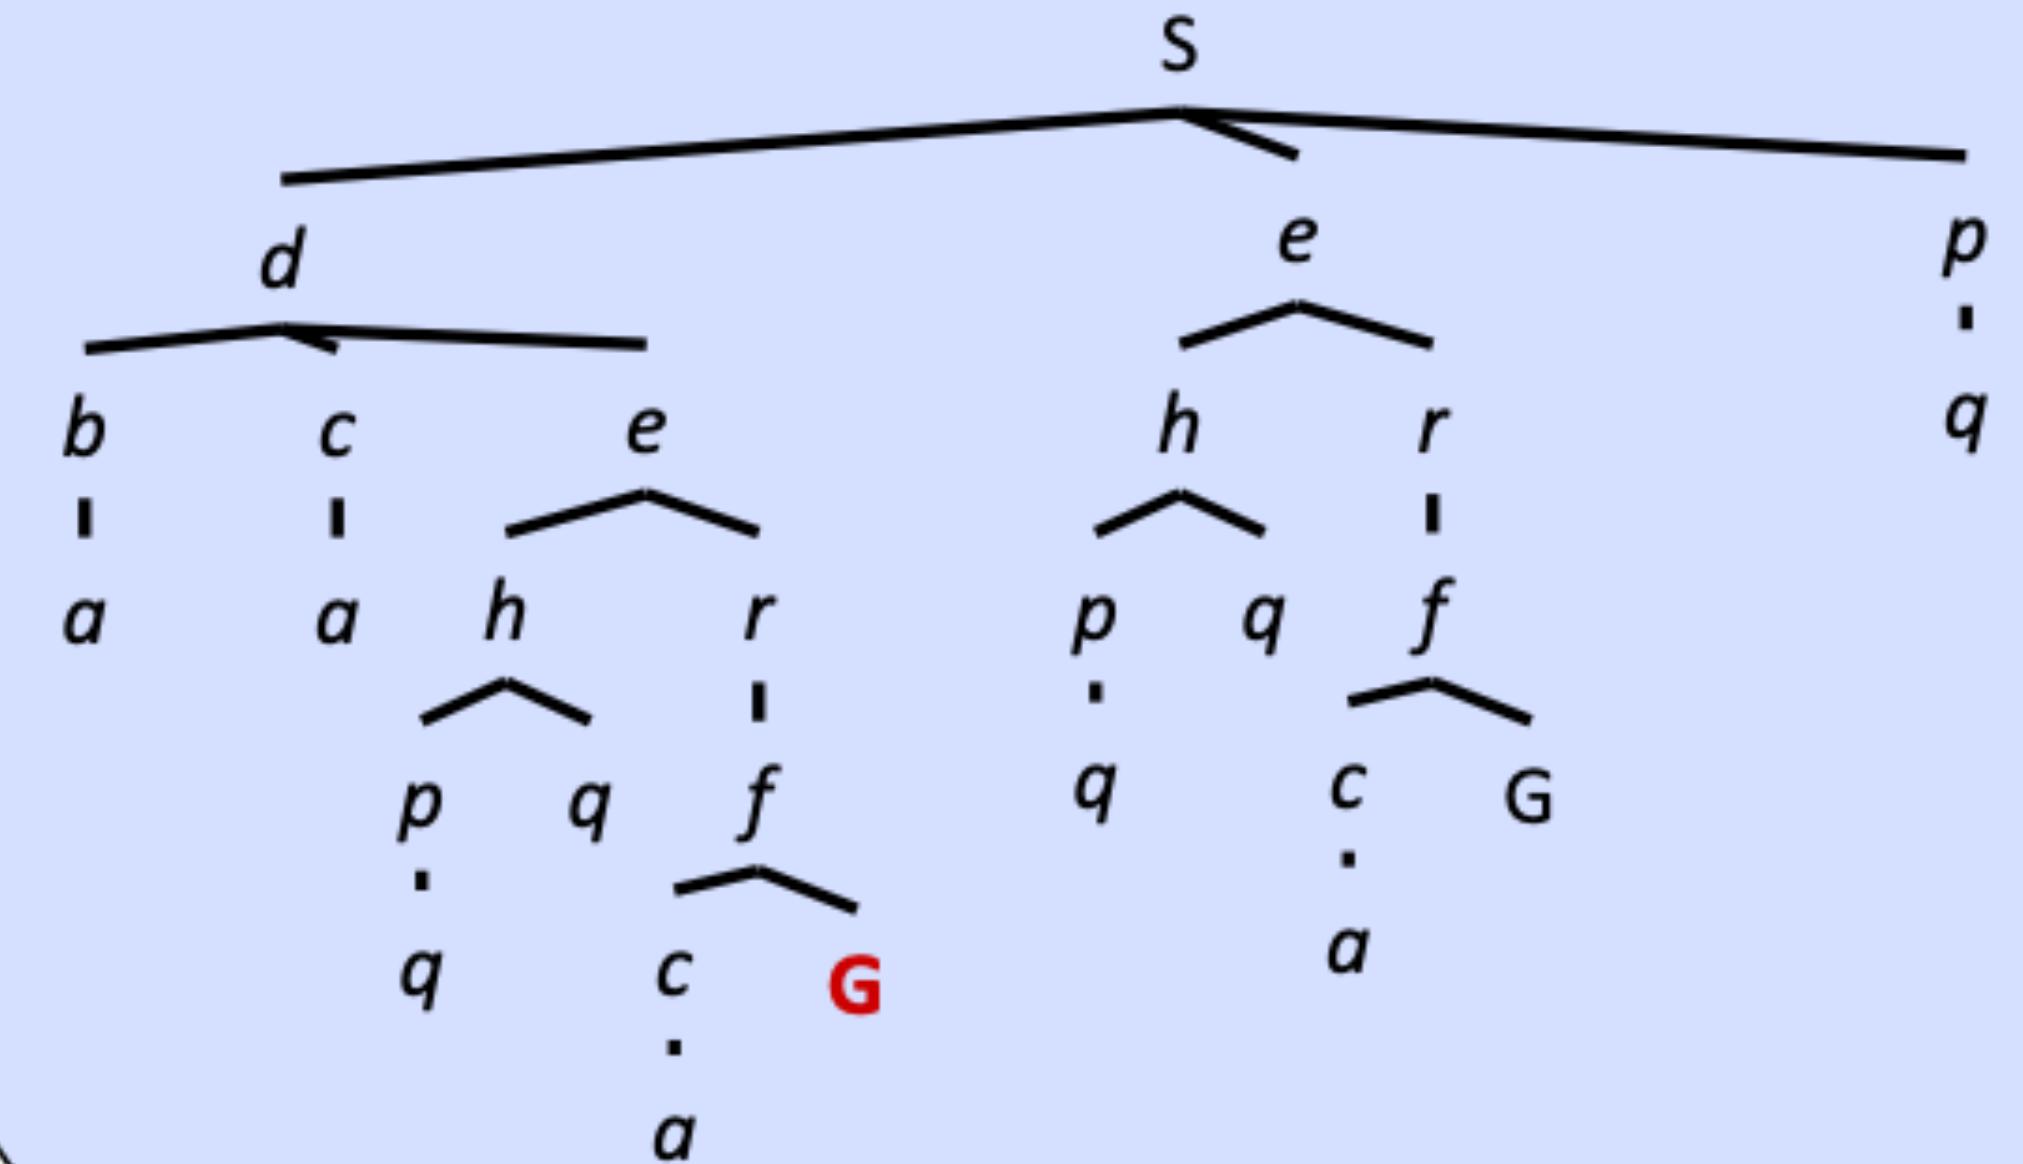
State Space Graph



Each NODE in the search tree is an entire PATH in the state space graph

We construct both on demand – and we construct as little as possible

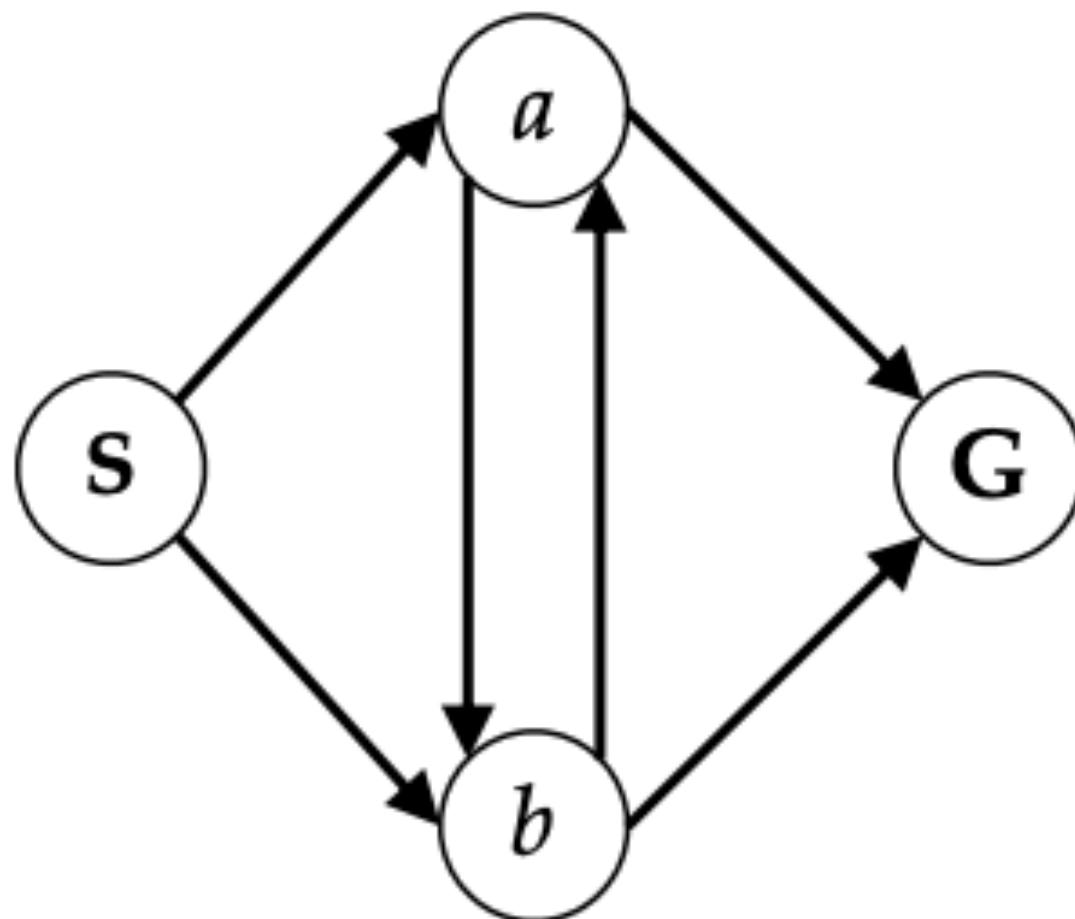
Search Tree



State Space Graph vs. Search Trees

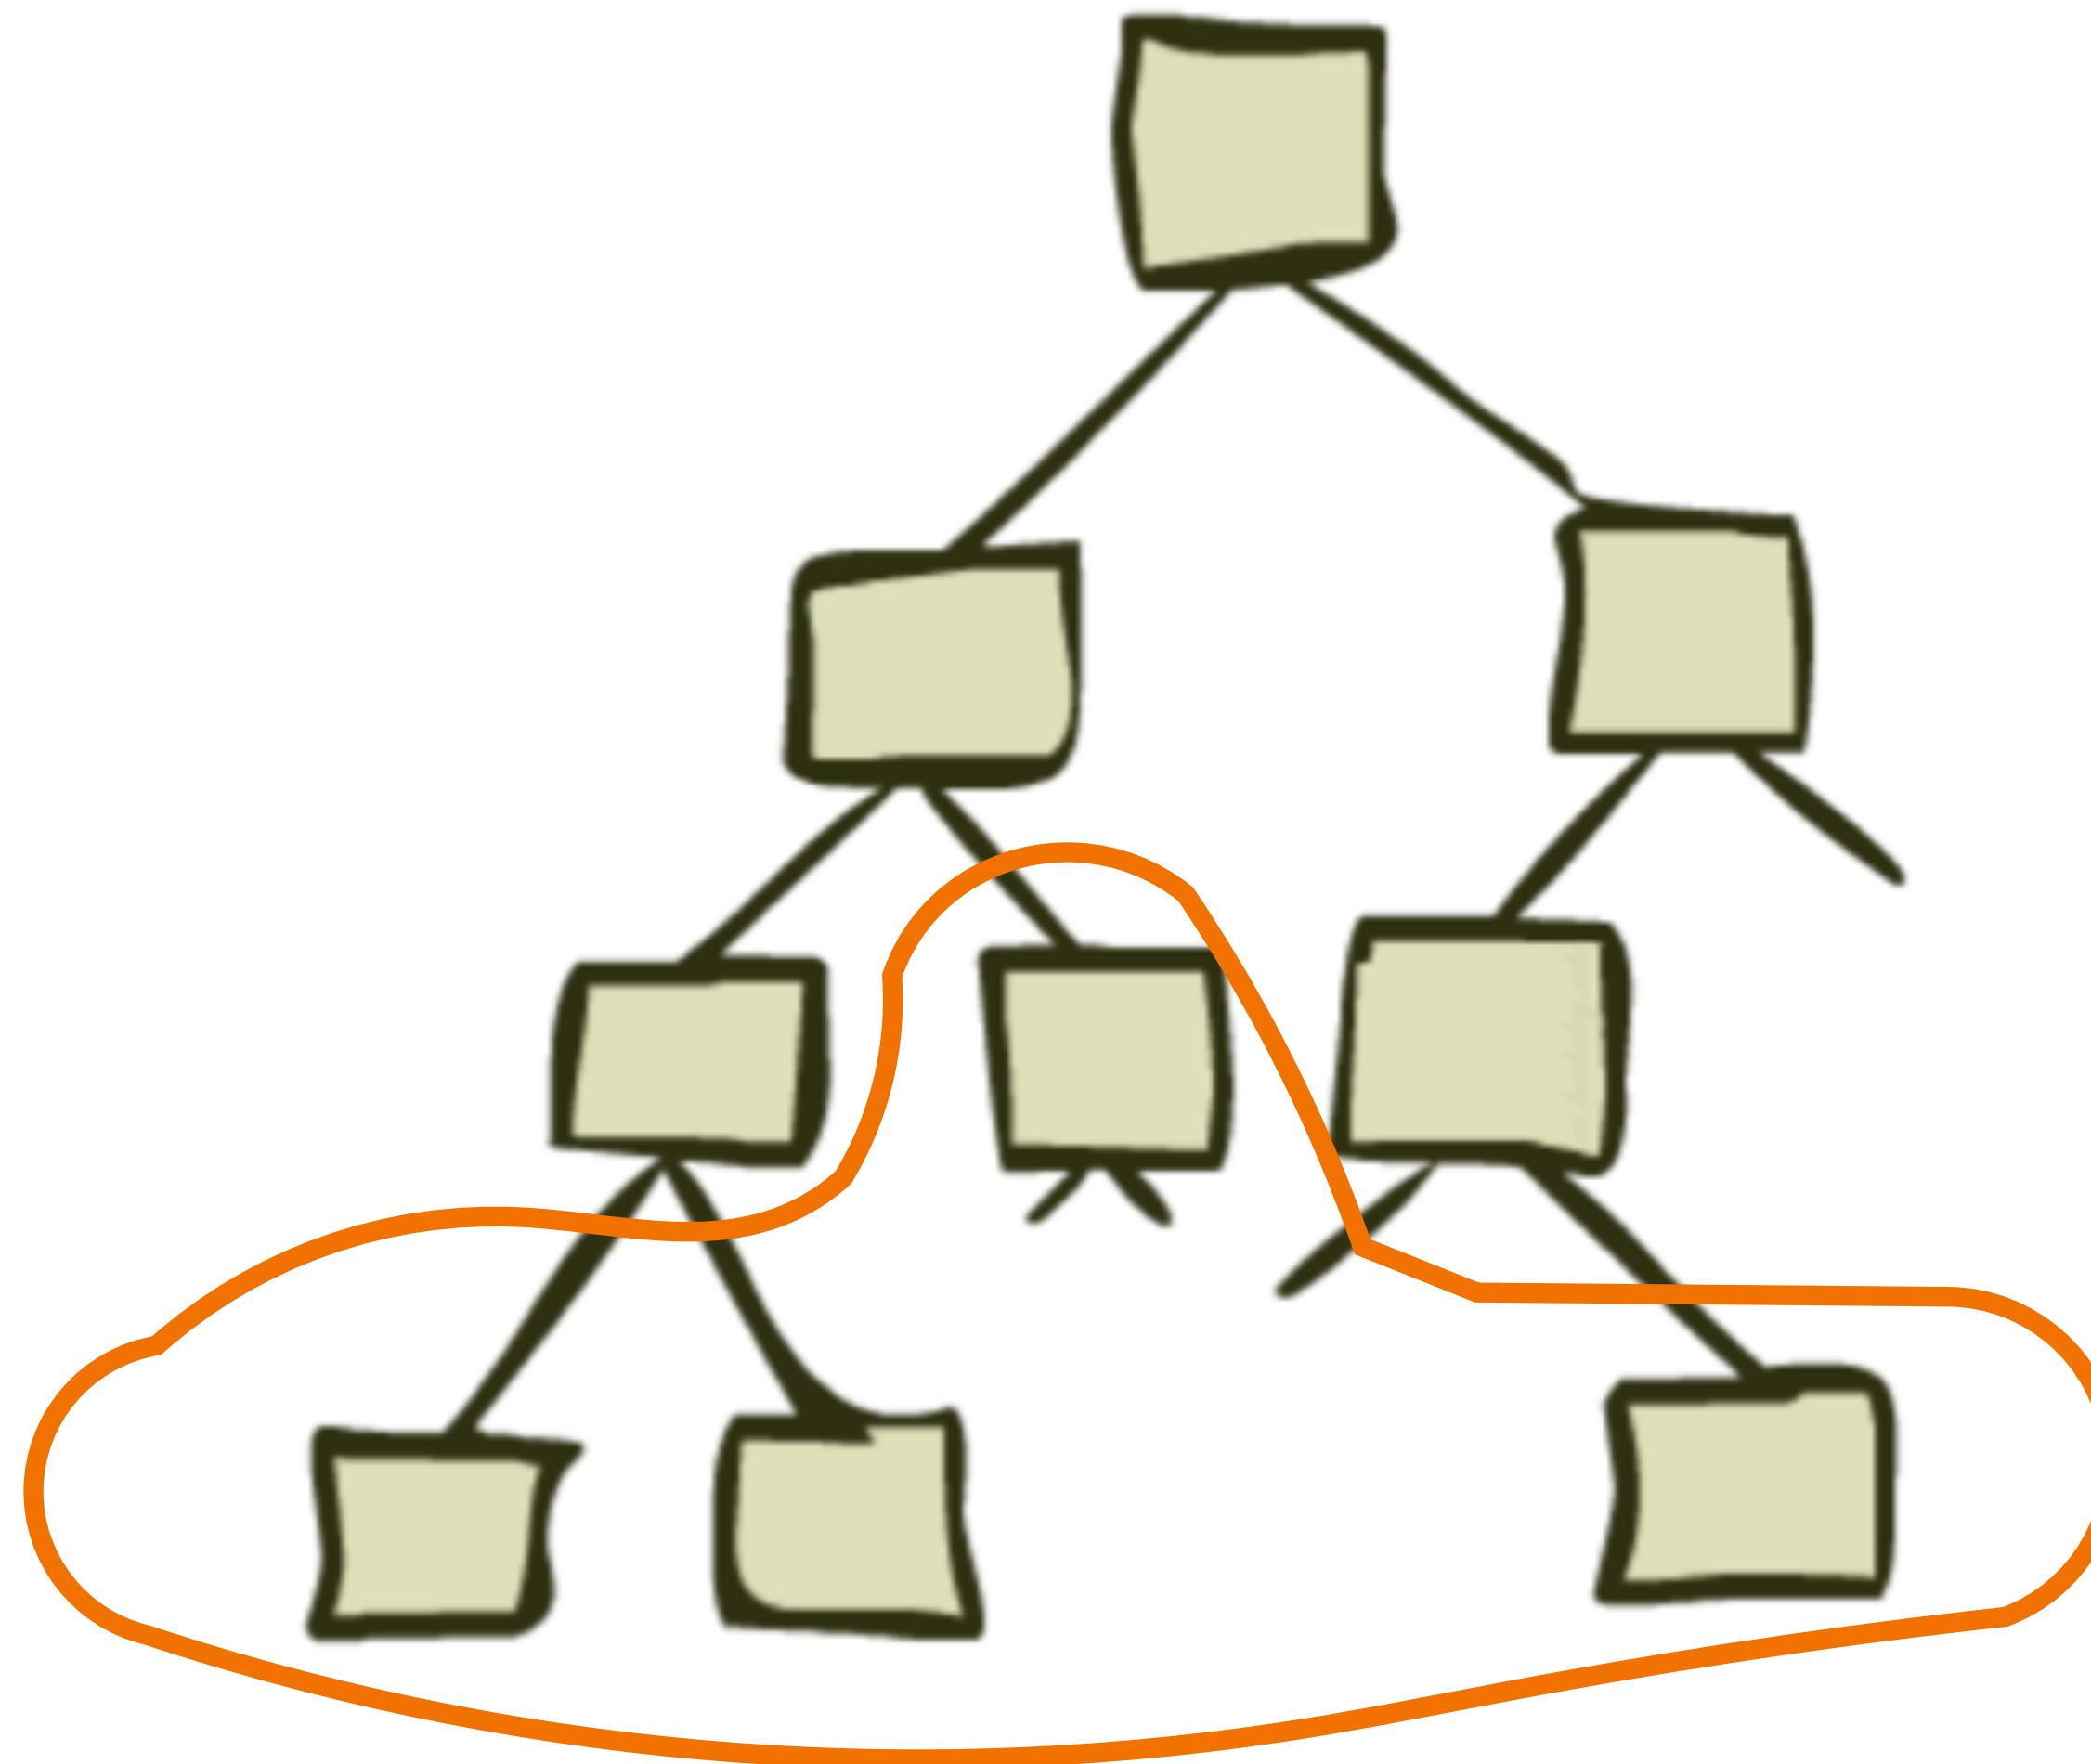
Consider this 4-state graph:

How big is its search tree (from S)?

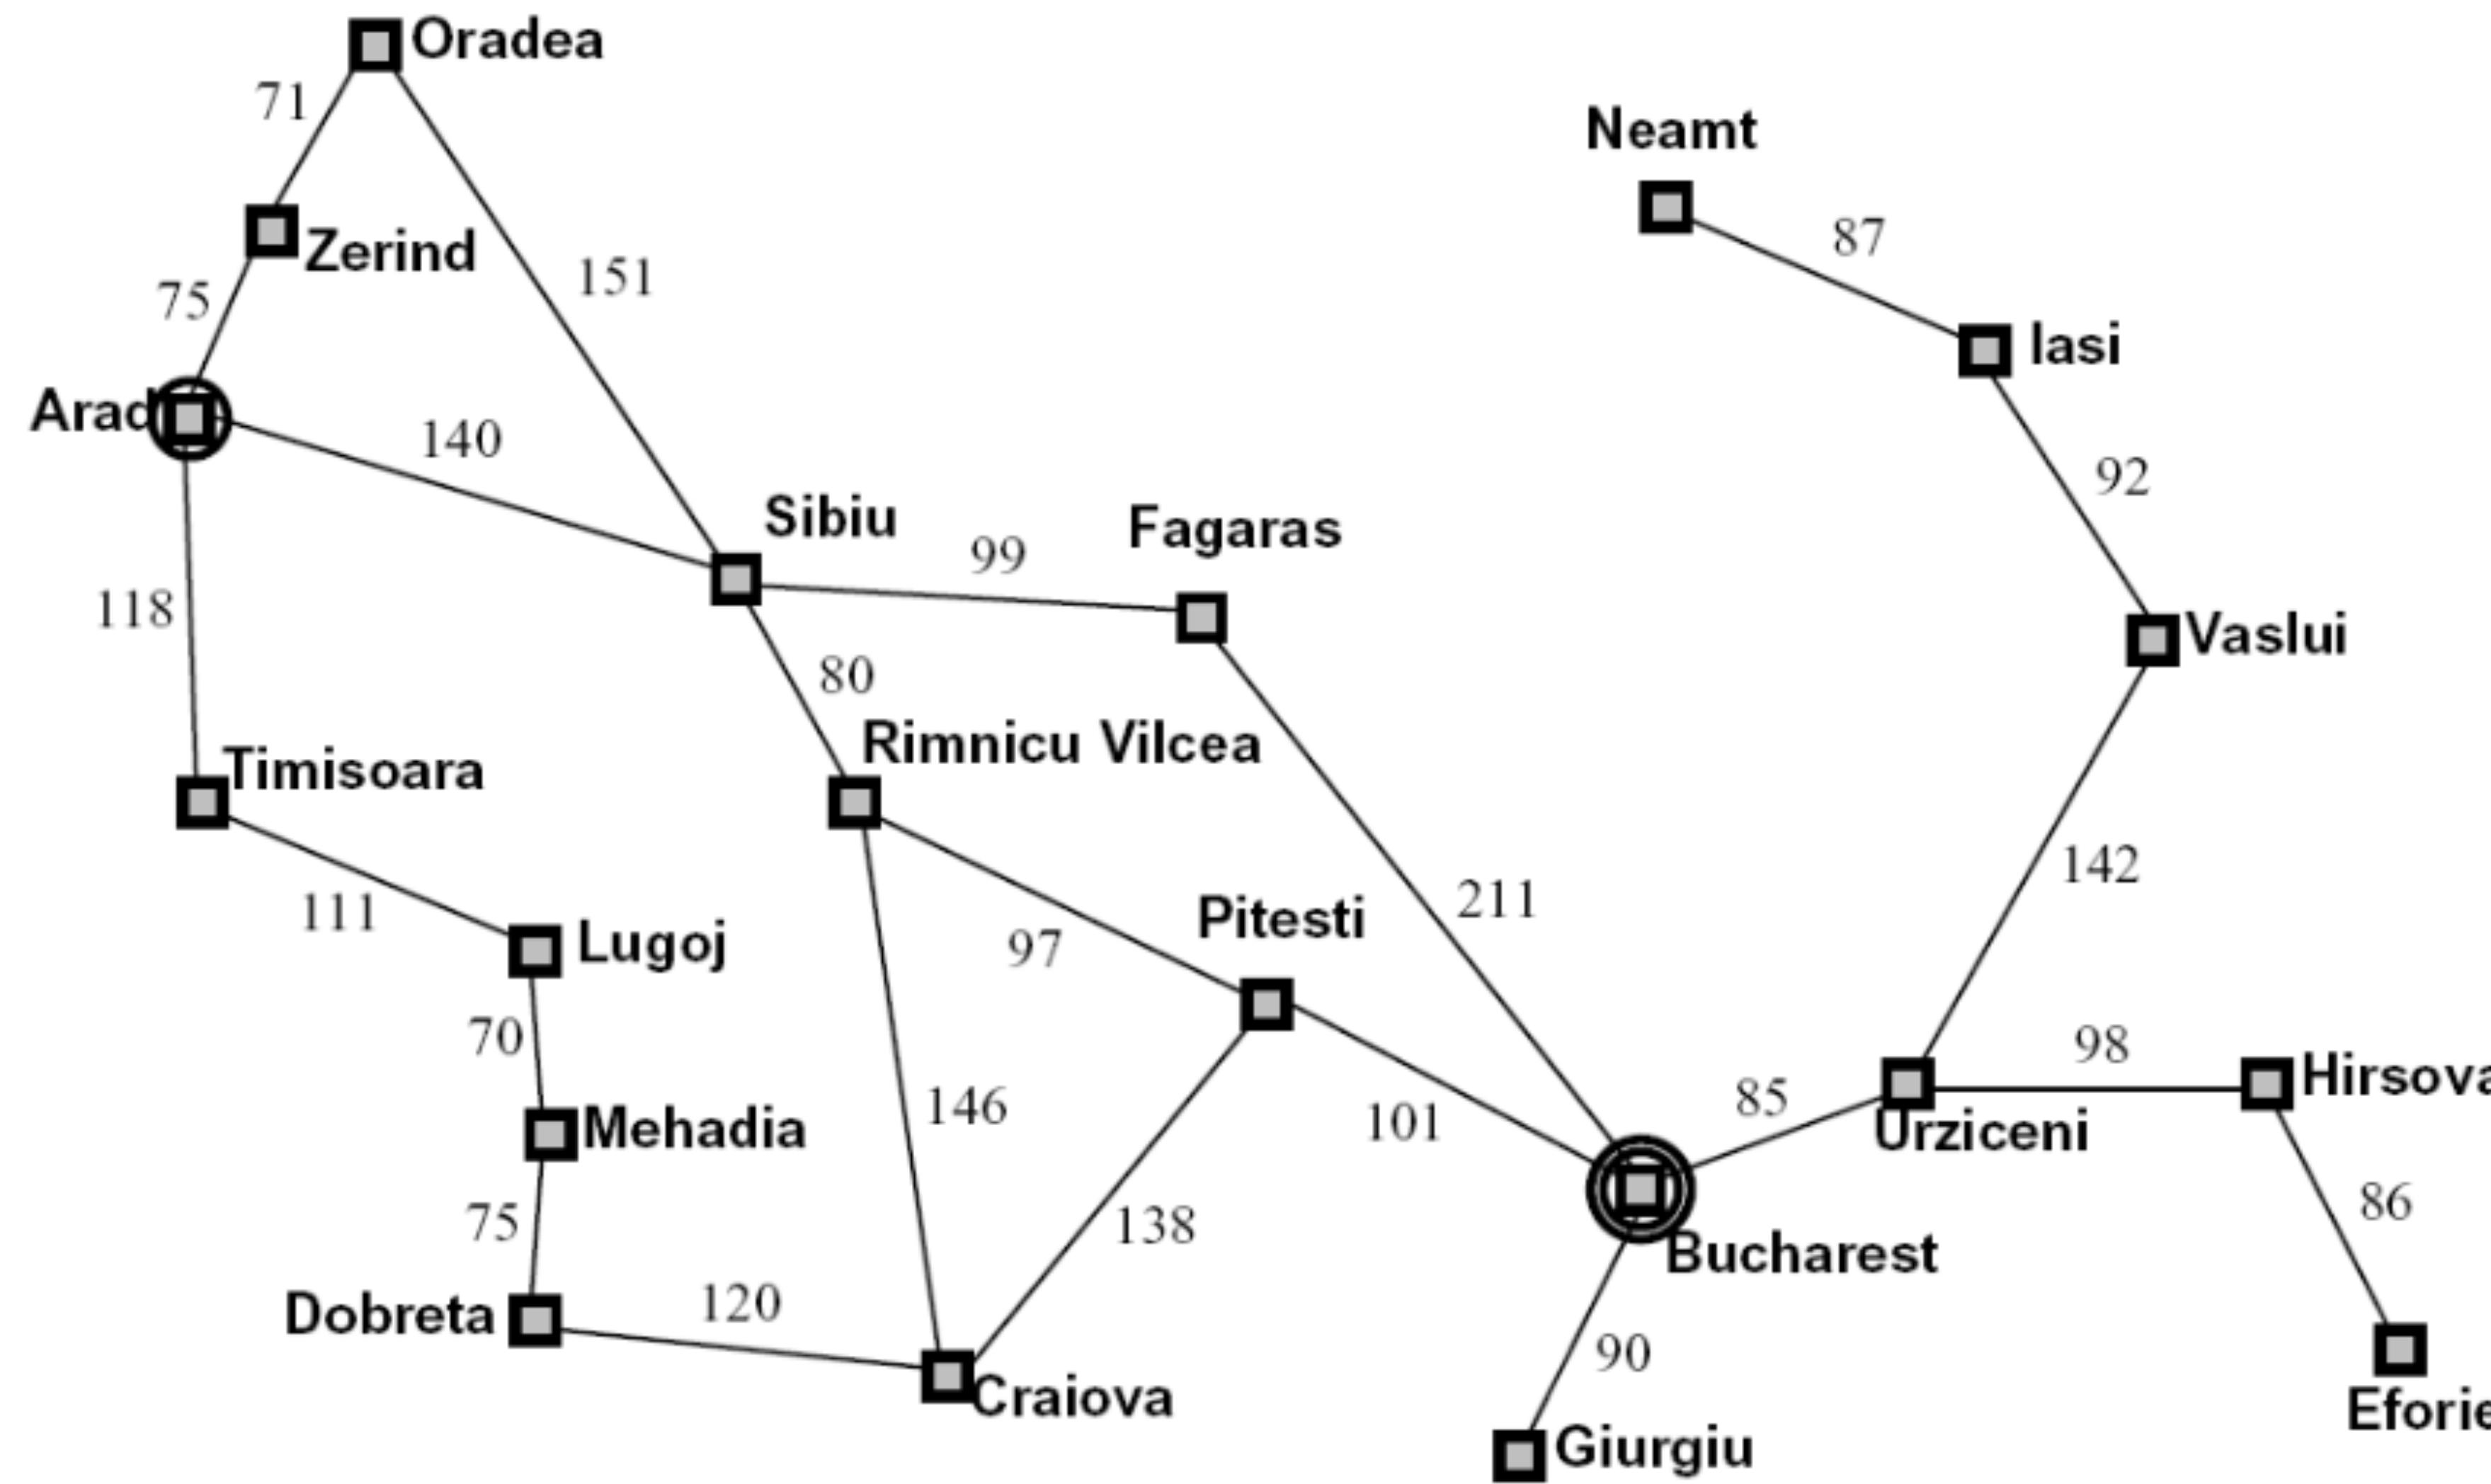


Important: Lots of repeated structure in the search tree!

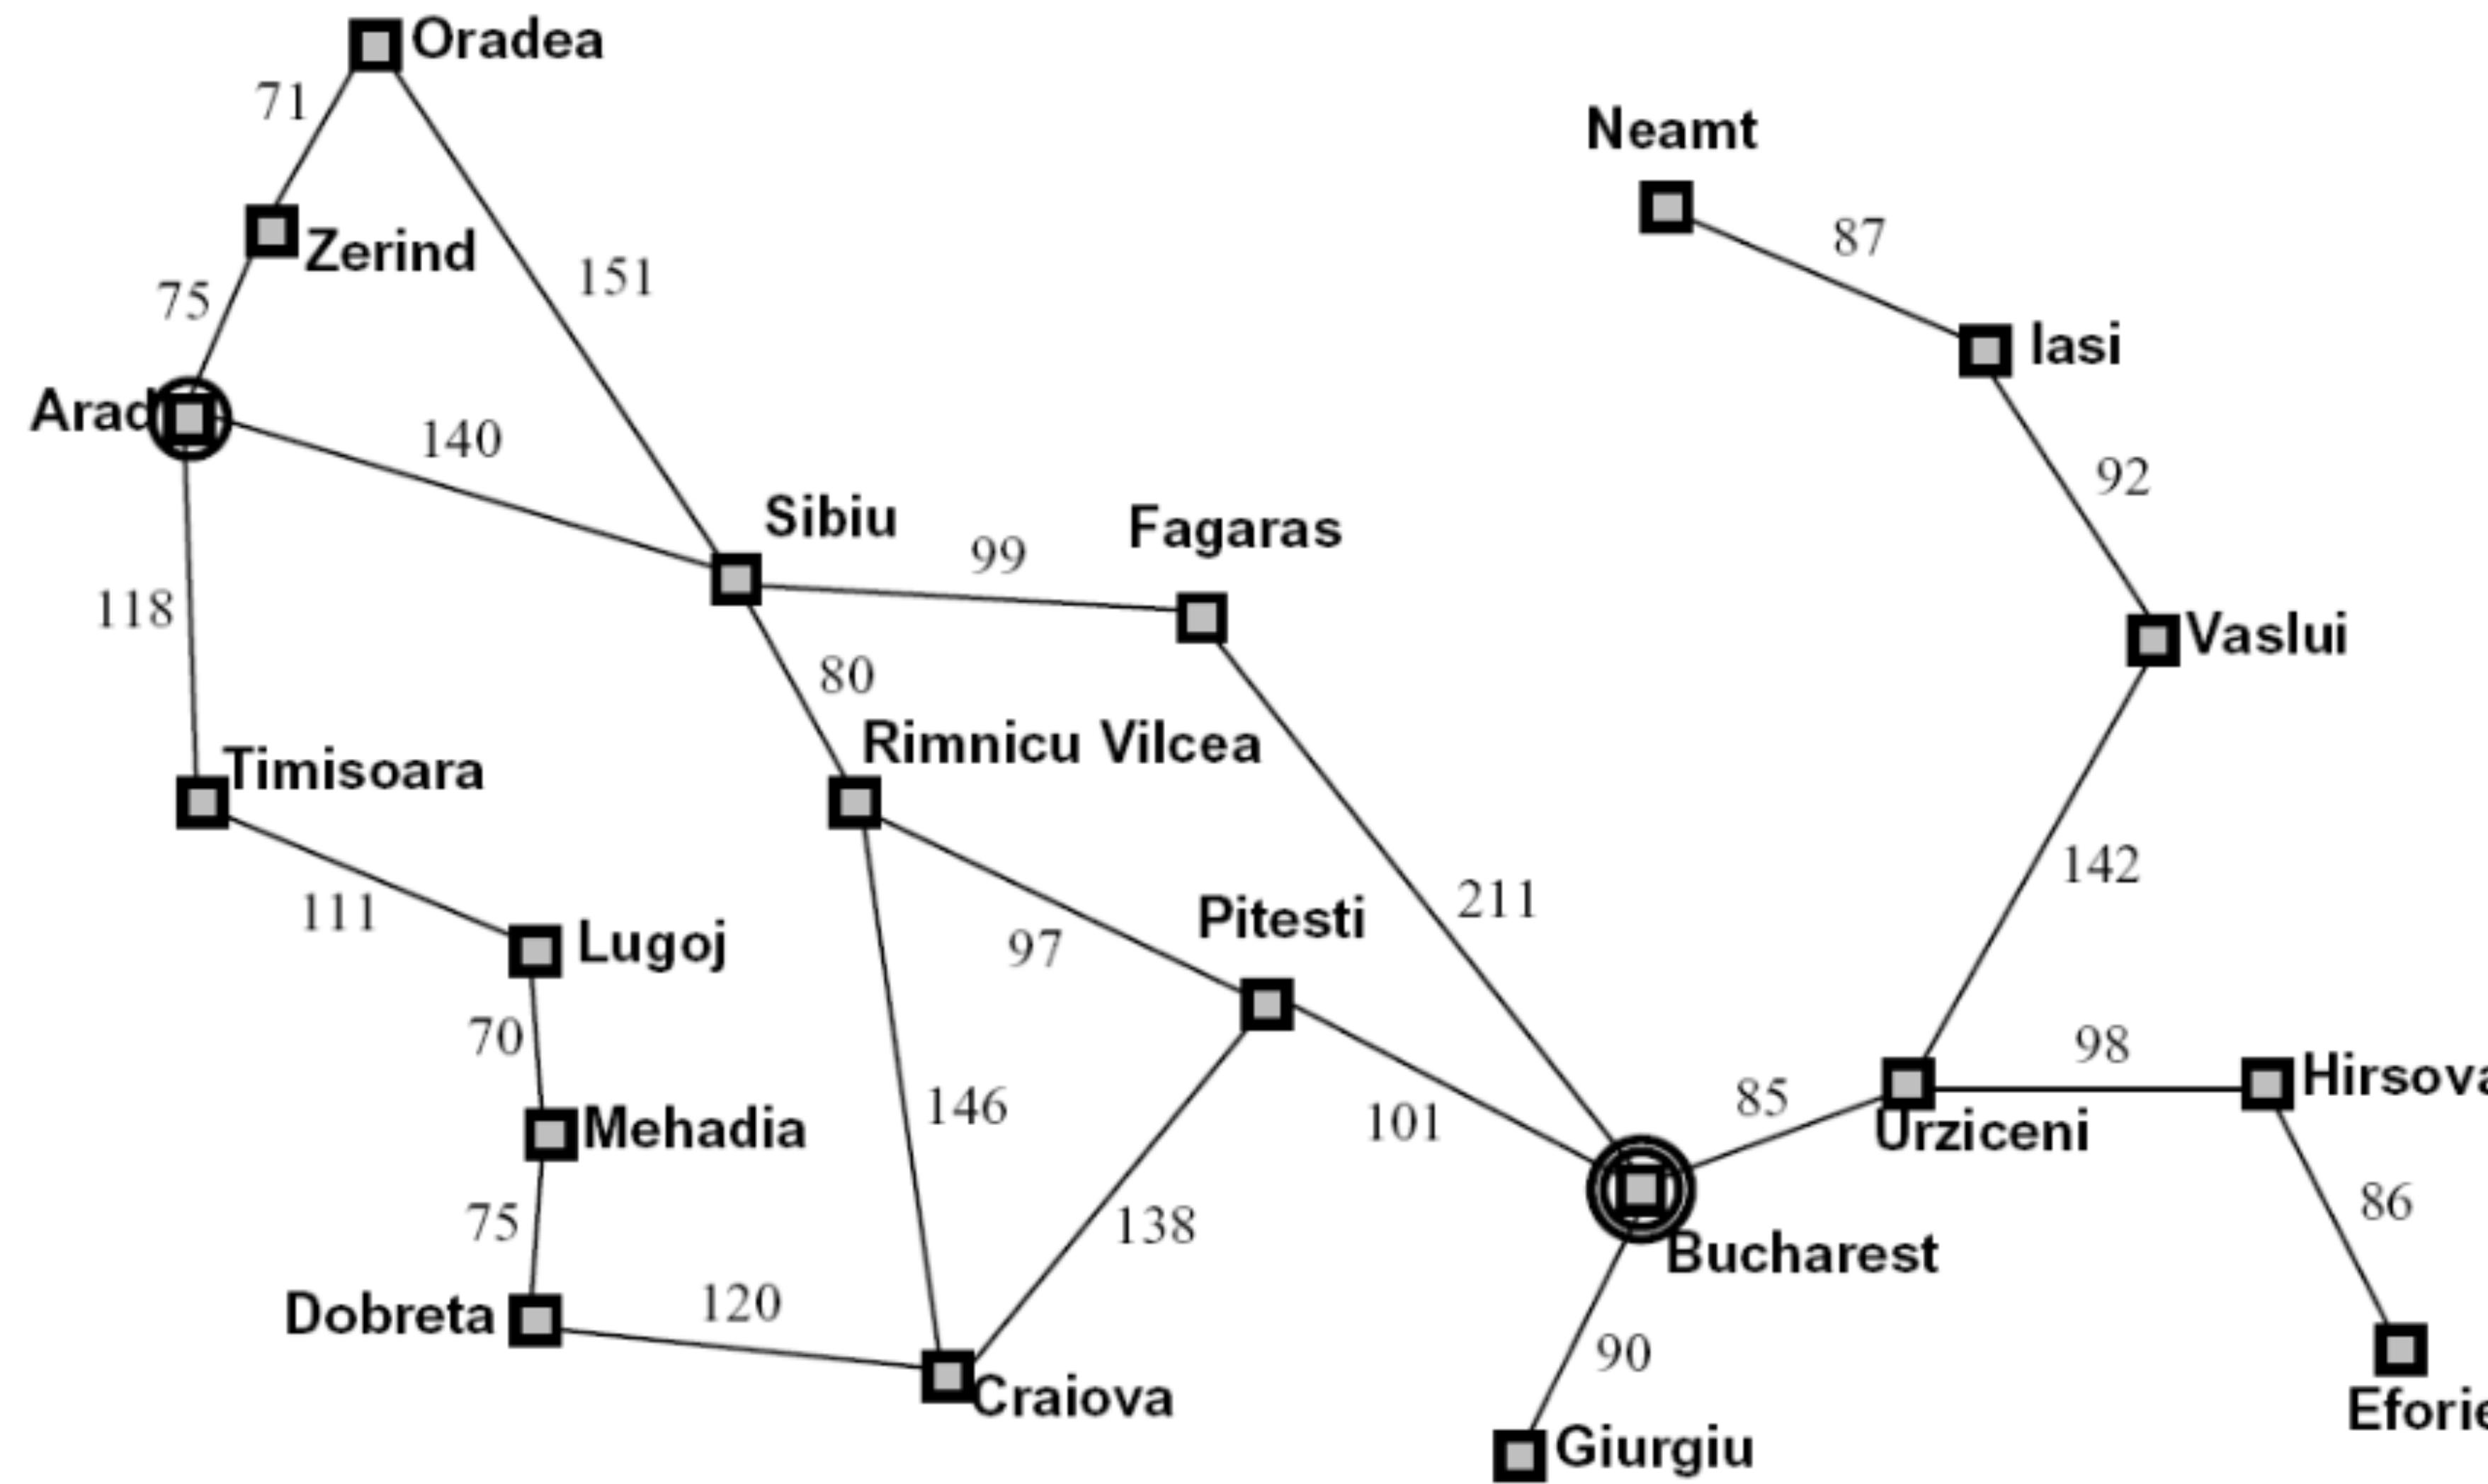
Tree Search



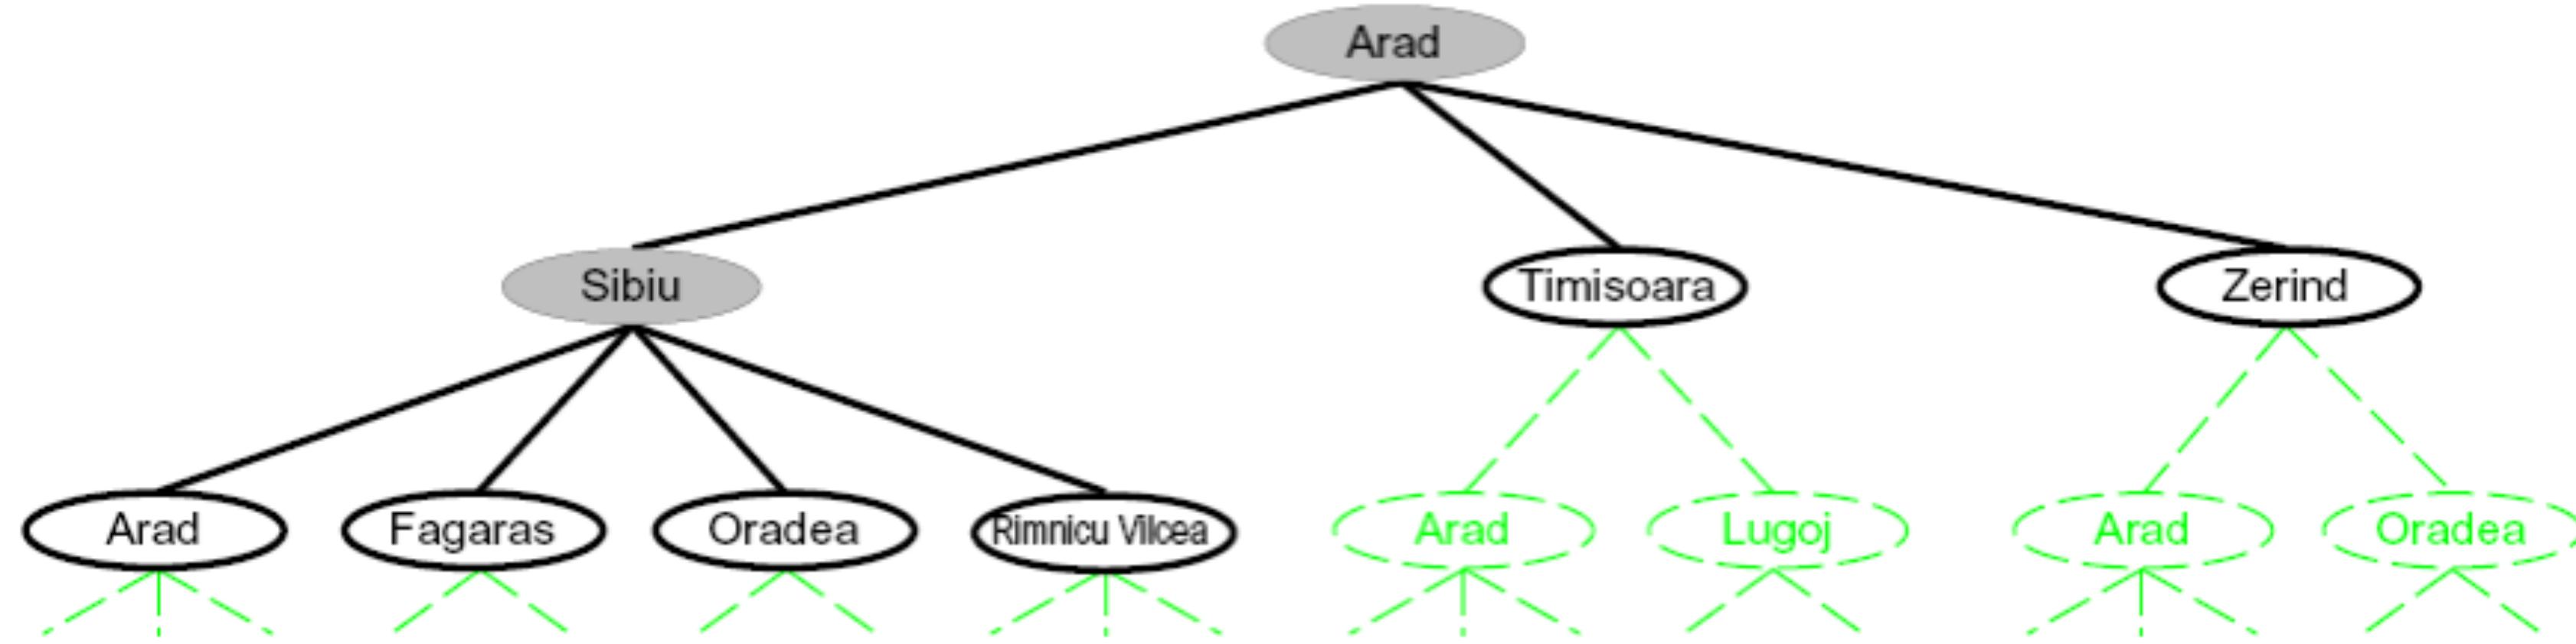
Search Example: Romania



CE 2: Draw/write out the search tree



Another Search Tree



- Search:
 - Expand out possible plans
 - Maintain a **fringe** of unexpanded plans
 - Try to expand as few tree nodes as possible

General Tree Search

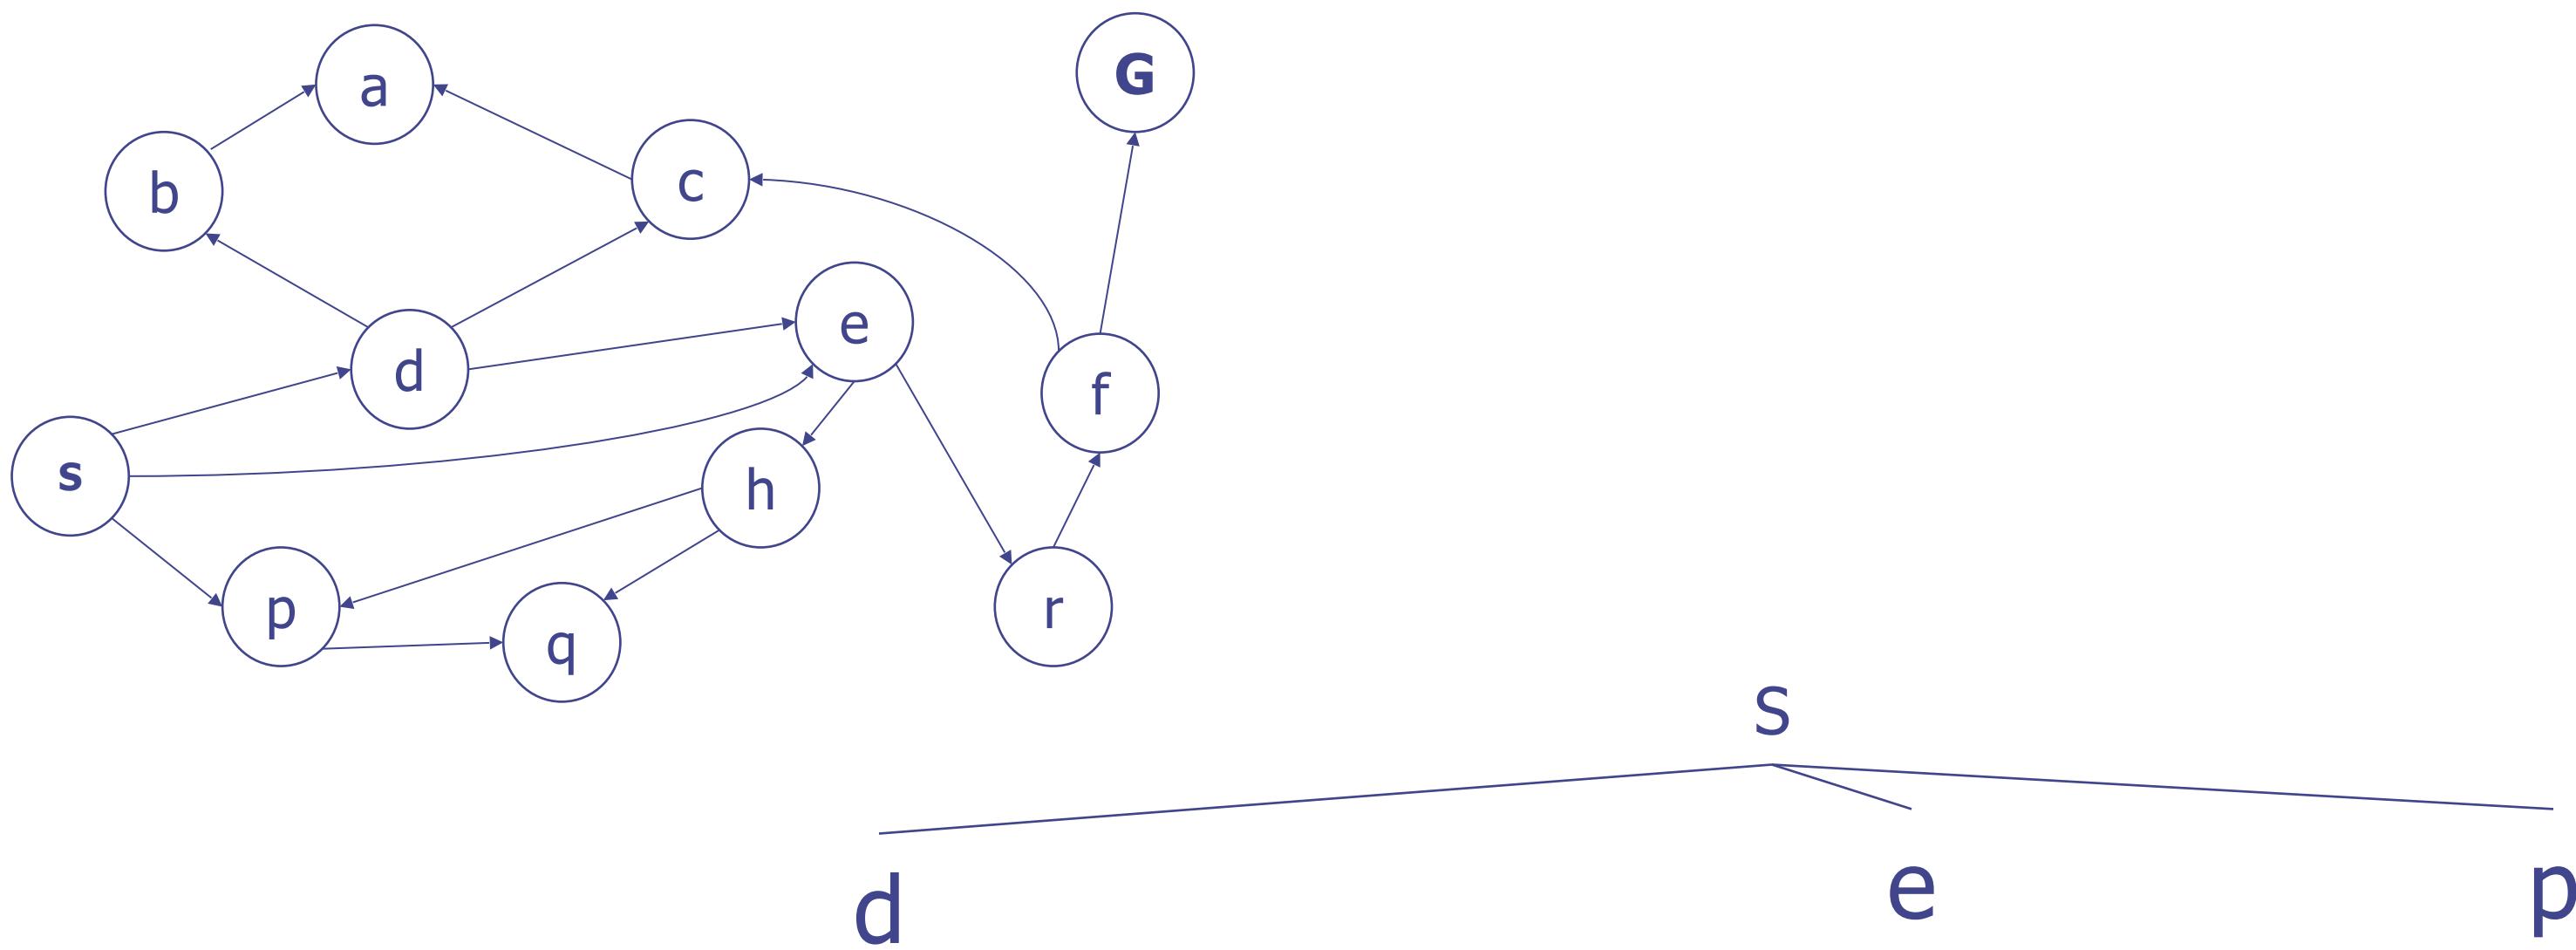
```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
    initialize the search tree using the initial state of problem
    loop do
        if there are no candidates for expansion then return failure
        choose a leaf node for expansion according to strategy
        if the node contains a goal state then return the corresponding solution
        else expand the node and add the resulting nodes to the search tree
    end
```

Important ideas:

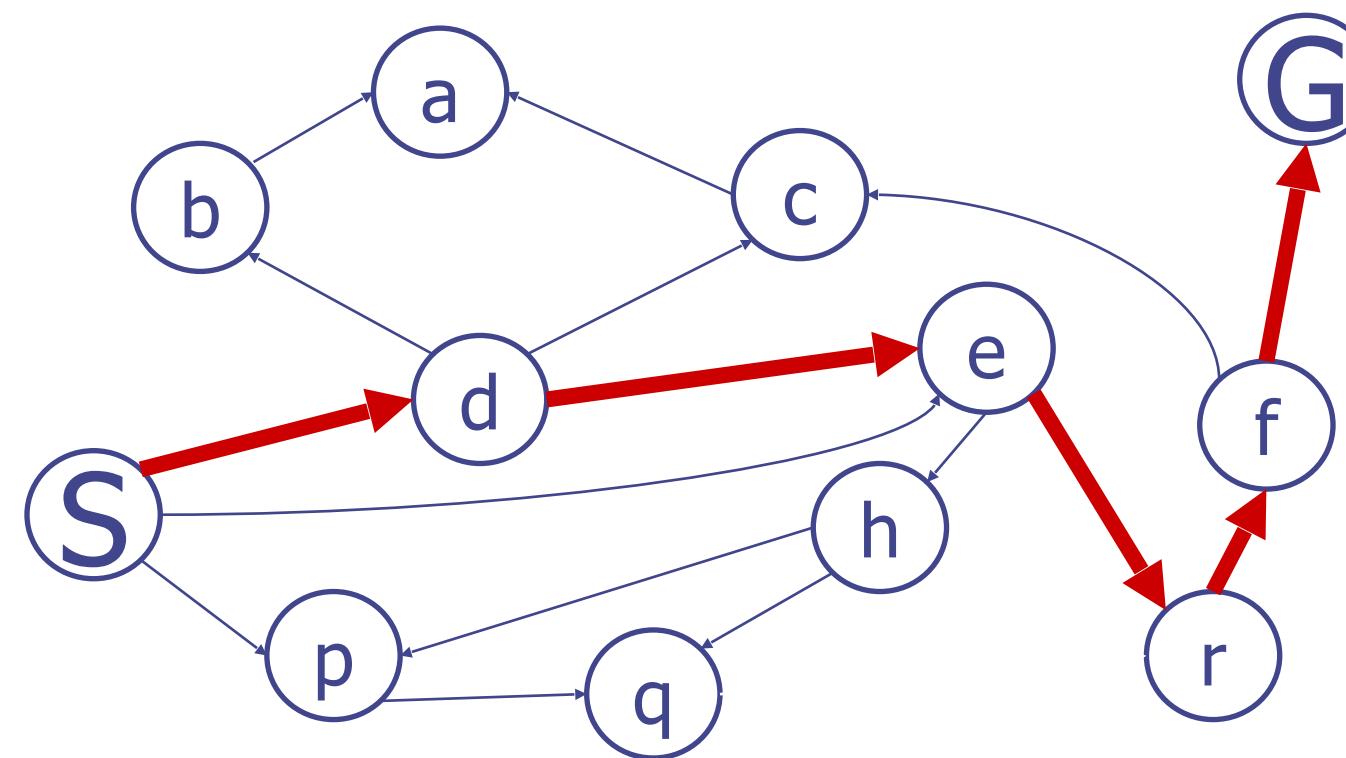
- Fringe or frontier set
- Expansion
- Exploration strategy
- Main question: which fringe nodes to explore?

**Detailed
pseudocode is in
the book!**

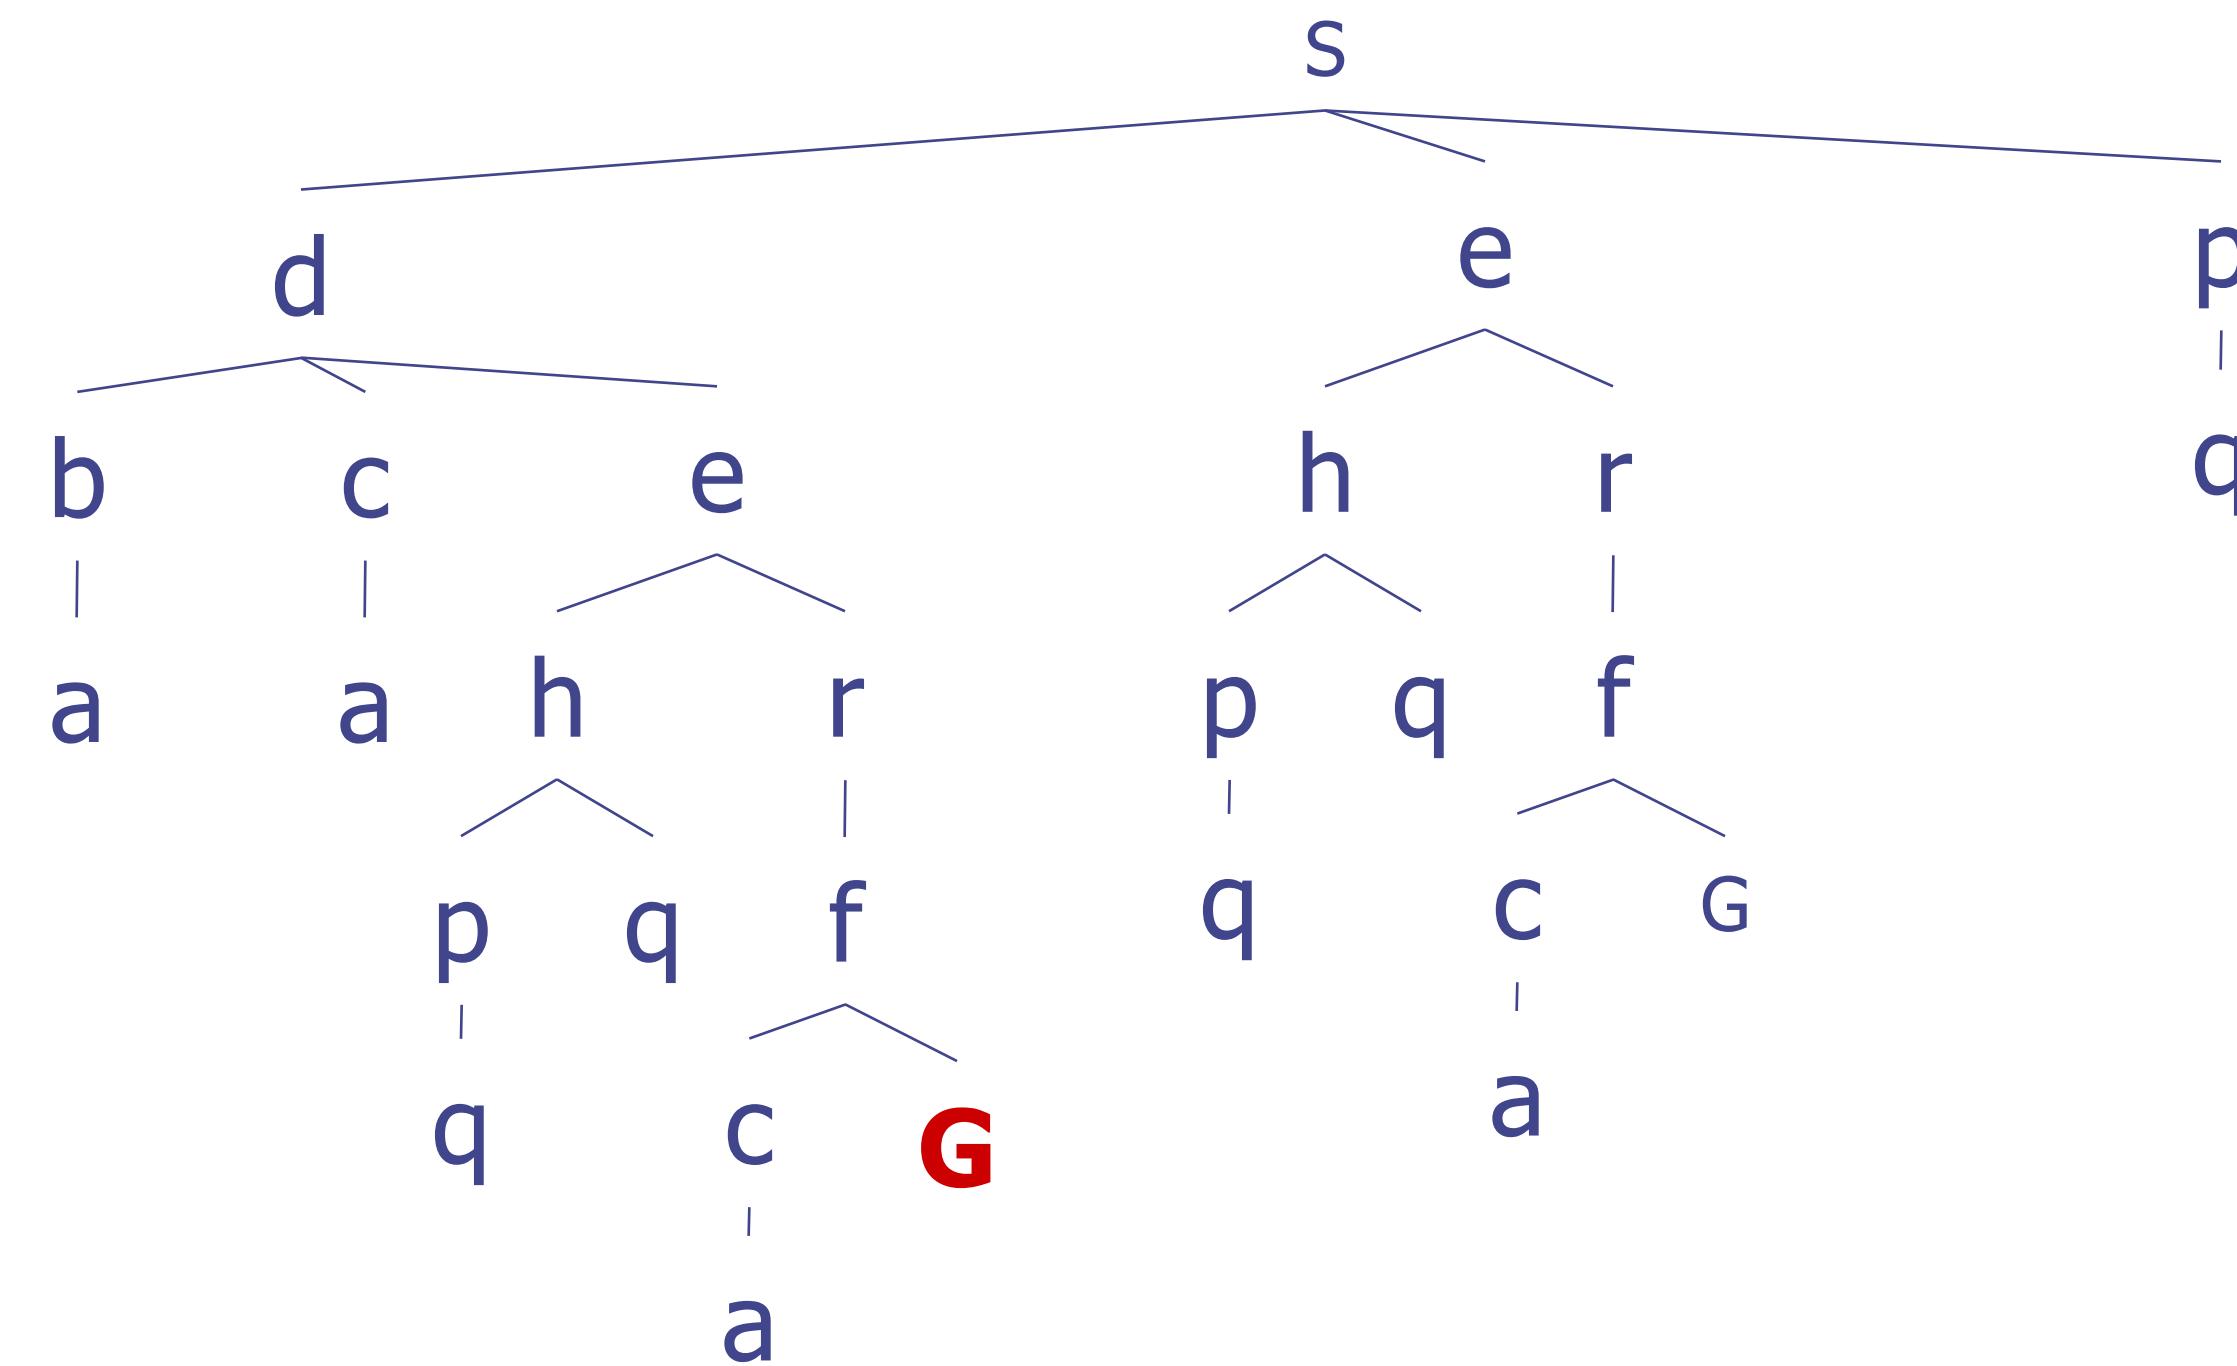
Example: Tree Search



Example: Tree Search



We construct both on demand – and we construct as little as possible.

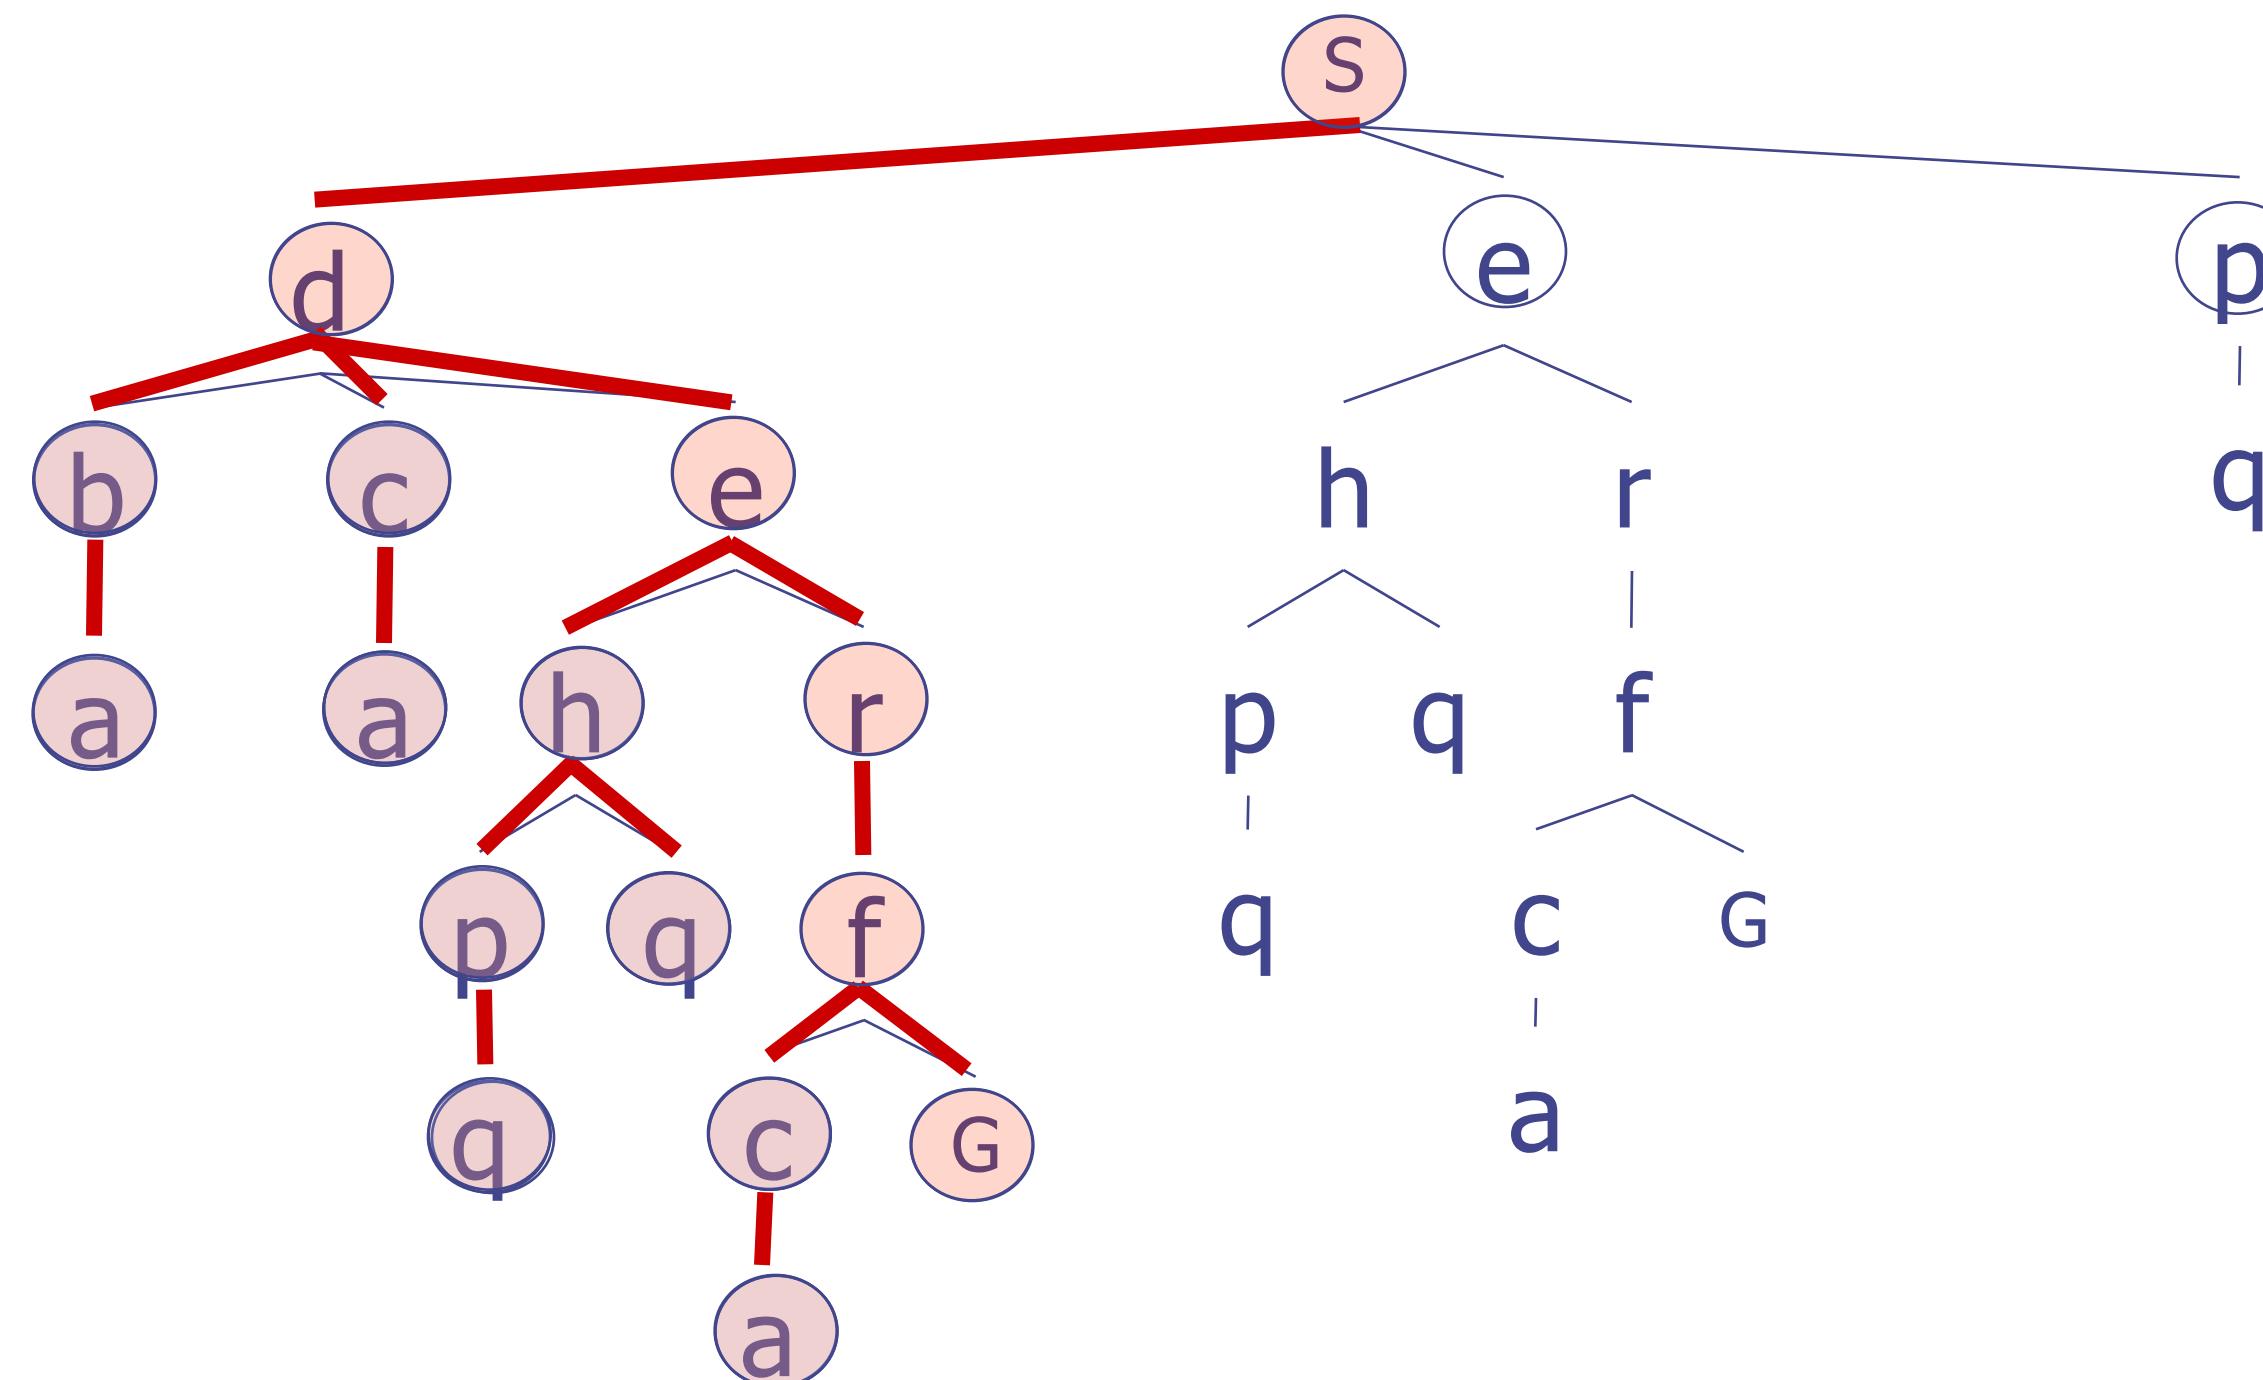
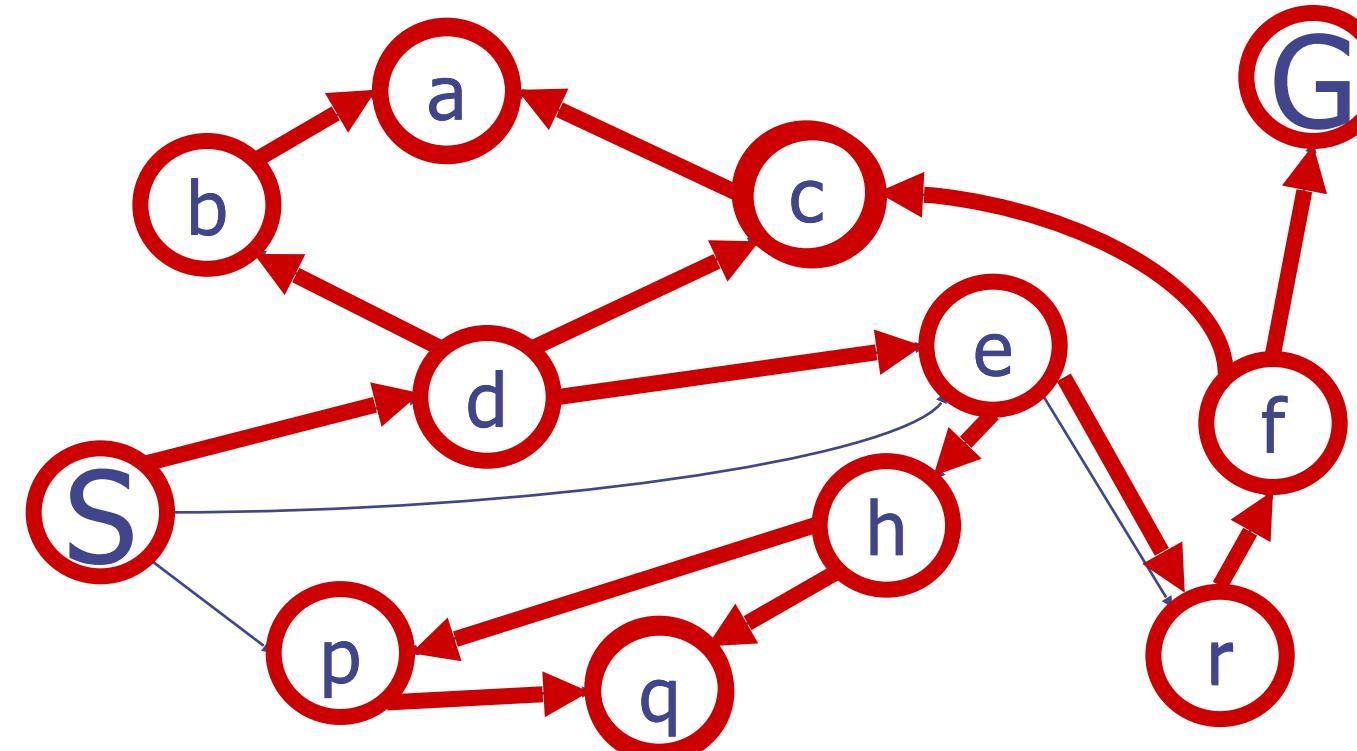


Each NODE in the search tree is an entire PATH in the state space graph.

Depth First Search

Review: Depth First Search

Strategy: expand deepest node first
Implementation:
Fringe is a LIFO stack



Search Algorithm Properties

Search Algorithm Properties

- Complete? **Guaranteed to find a solution if one exists?**
- Optimal? **Guaranteed to find the least cost path?**
- Time complexity?
- Space complexity?

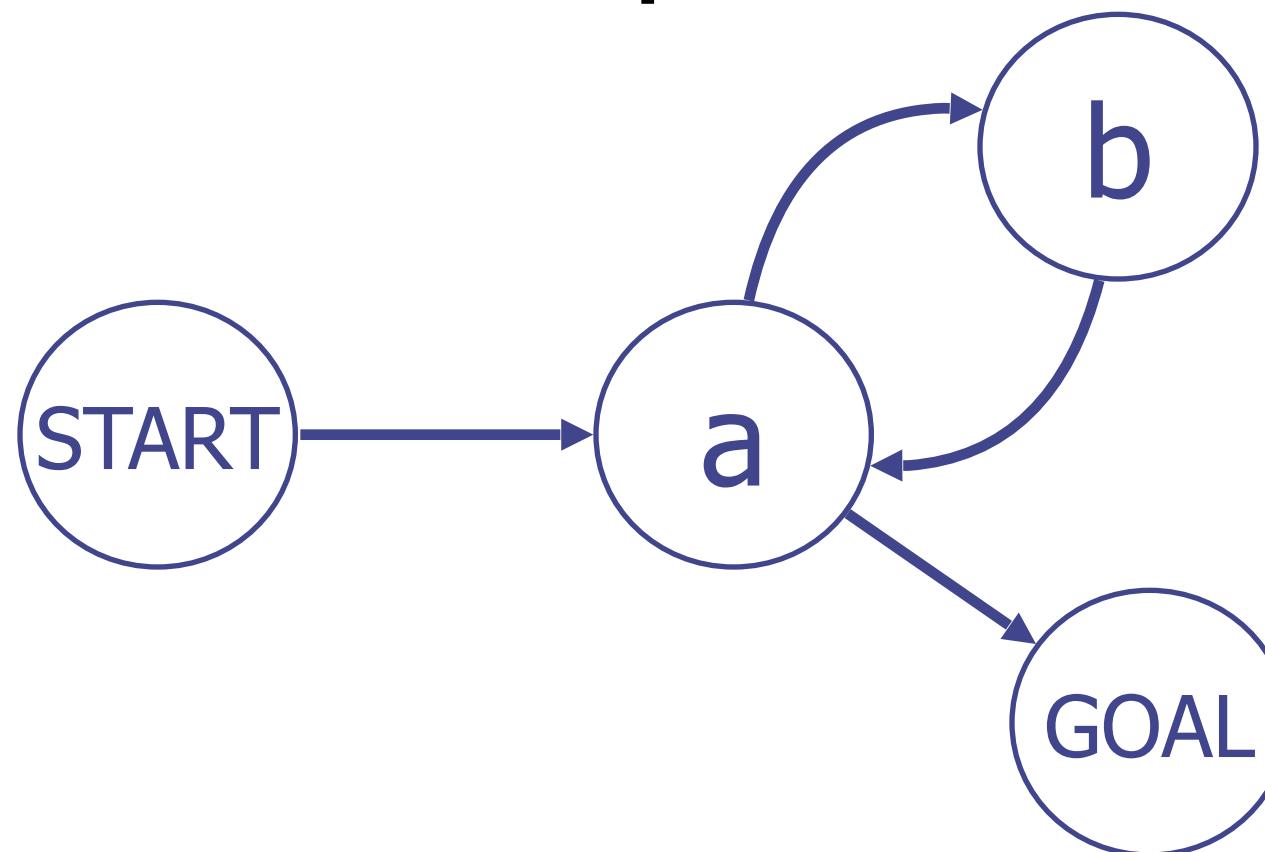
Variables:

n	Number of states in the problem
b	The average branching factor B (the average number of successors)
C^*	Cost of least cost solution
s	Depth of the shallowest solution
m	Max depth of the search tree

DFS

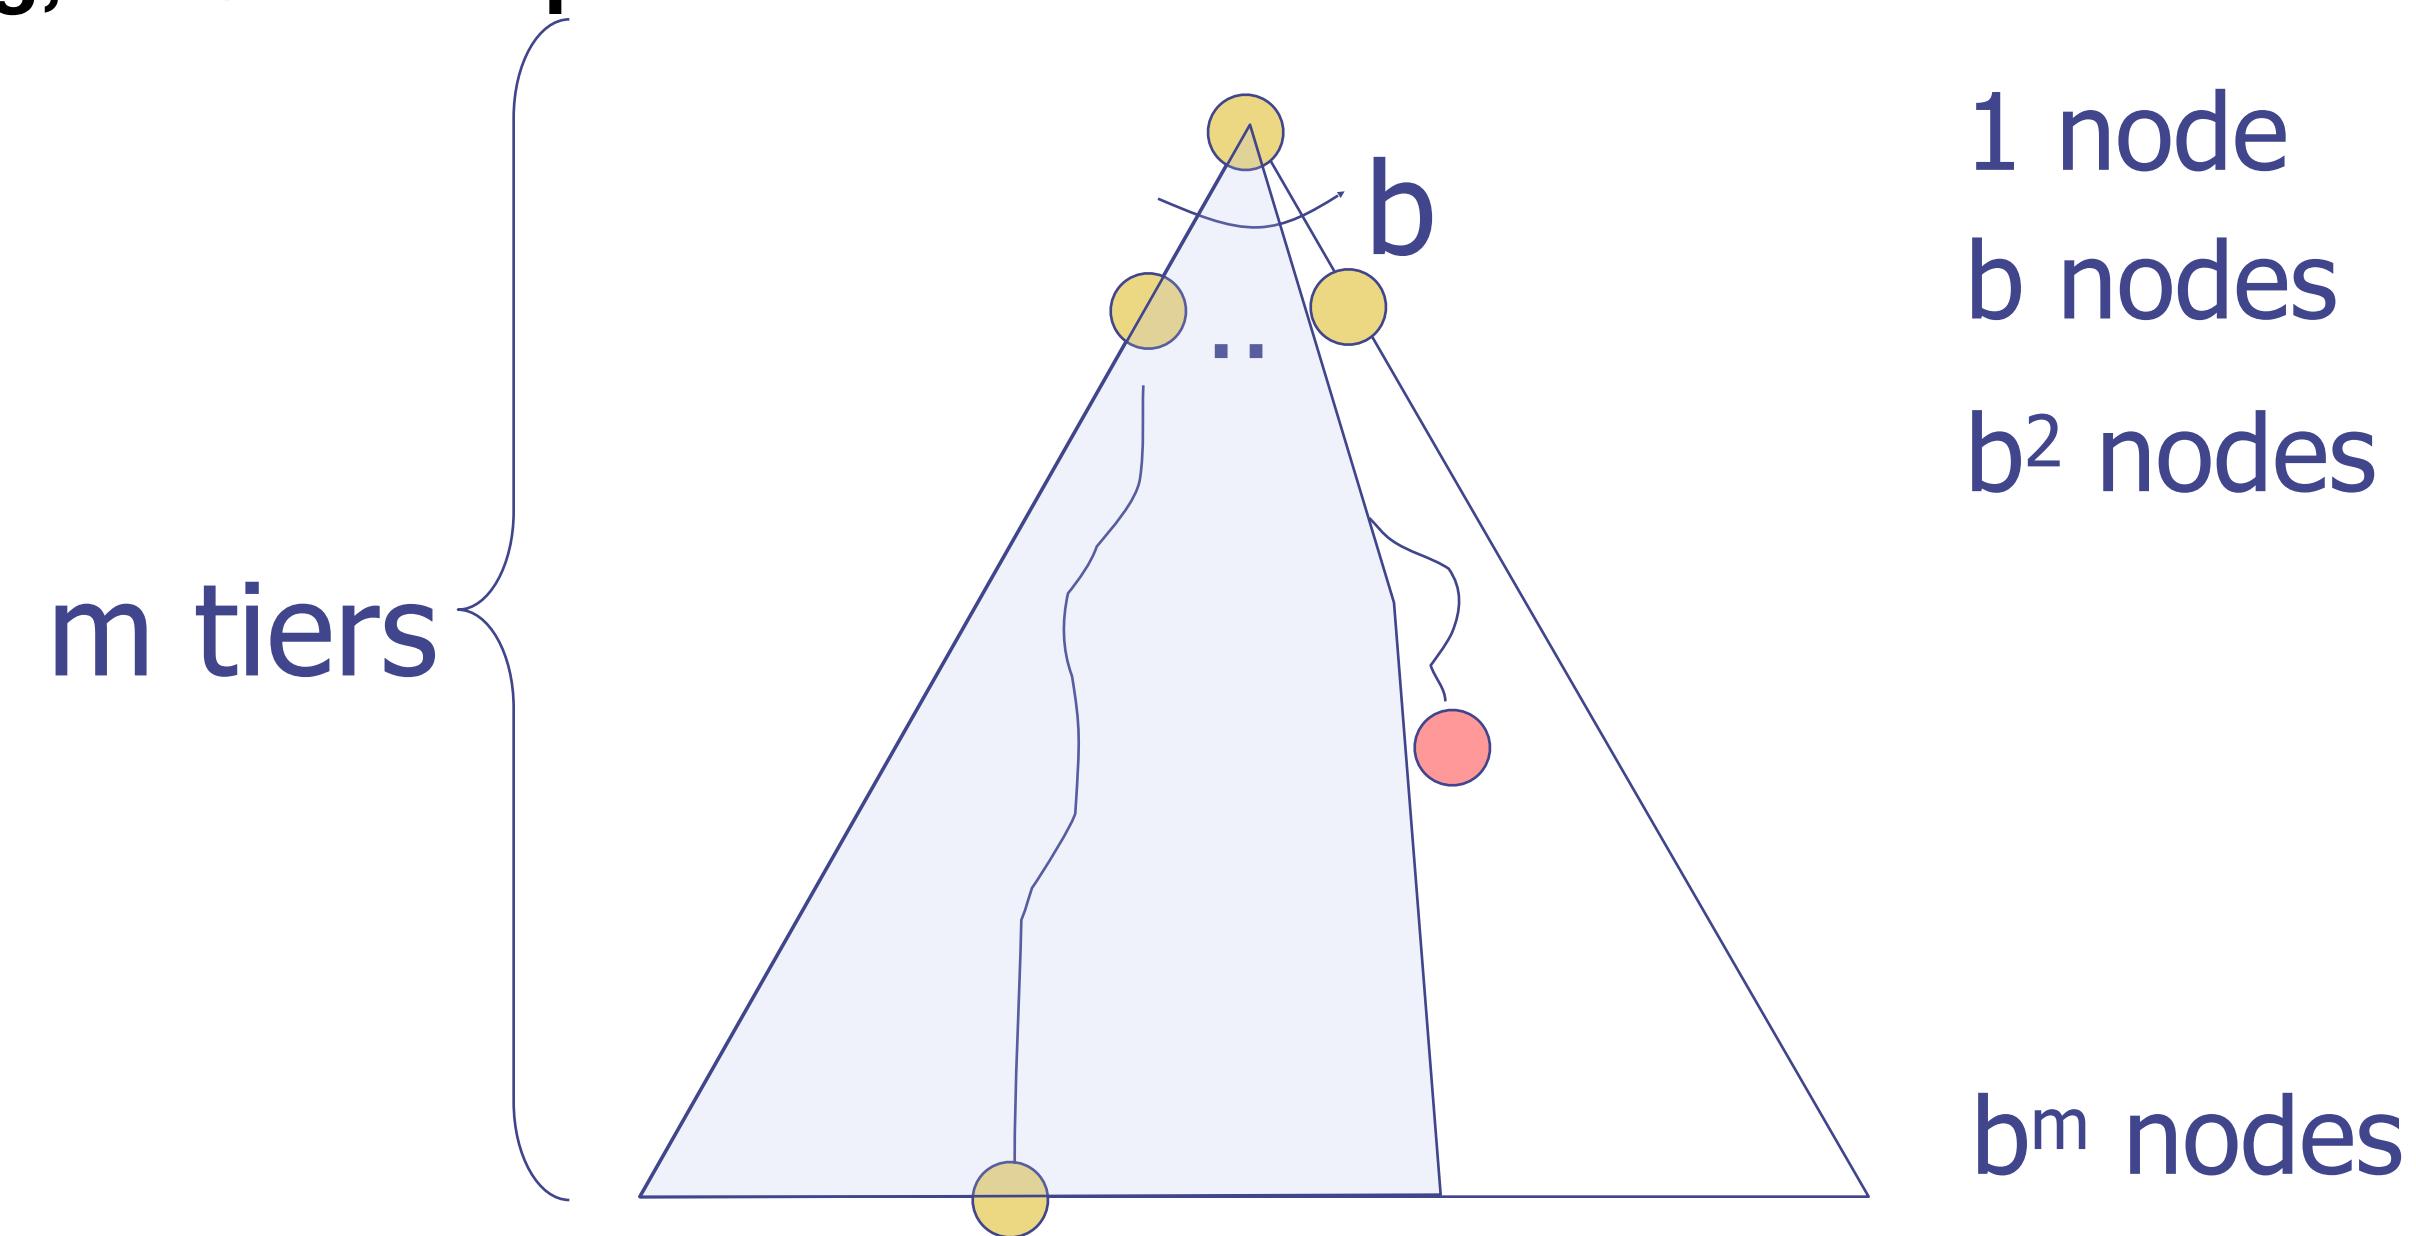
Algorithm	Complete	Optimal	Time	Space
DFS Depth First Search	N	N	Infinite	Infinite

- Infinite paths make DFS incomplete...
- How can we fix this?



DFS

- With cycle checking, DFS is complete.



Algorithm	Complete	Optimal	Time	Space
DFS w/ Path Checking	Y	N	$O(b^{m+1})$	$O(bm)$

- When is DFS optimal?

DFS Assumption

What assumption about the search problem should we have to be able to use DFS to optimally find the goal state?

Backtracking Search



Algorithm: backtracking search

```
def backtrackingSearch( $s$ , path):
    If IsEnd( $s$ ): update minimum cost path
    For each action  $a \in \text{Actions}(s)$ :
        Extend path with  $\text{Succ}(s, a)$  and  $\text{Cost}(s, a)$ 
        Call backtrackingSearch( $\text{Succ}(s, a)$ , path)
    Return minimum cost path
```

Backtracking Search Performance

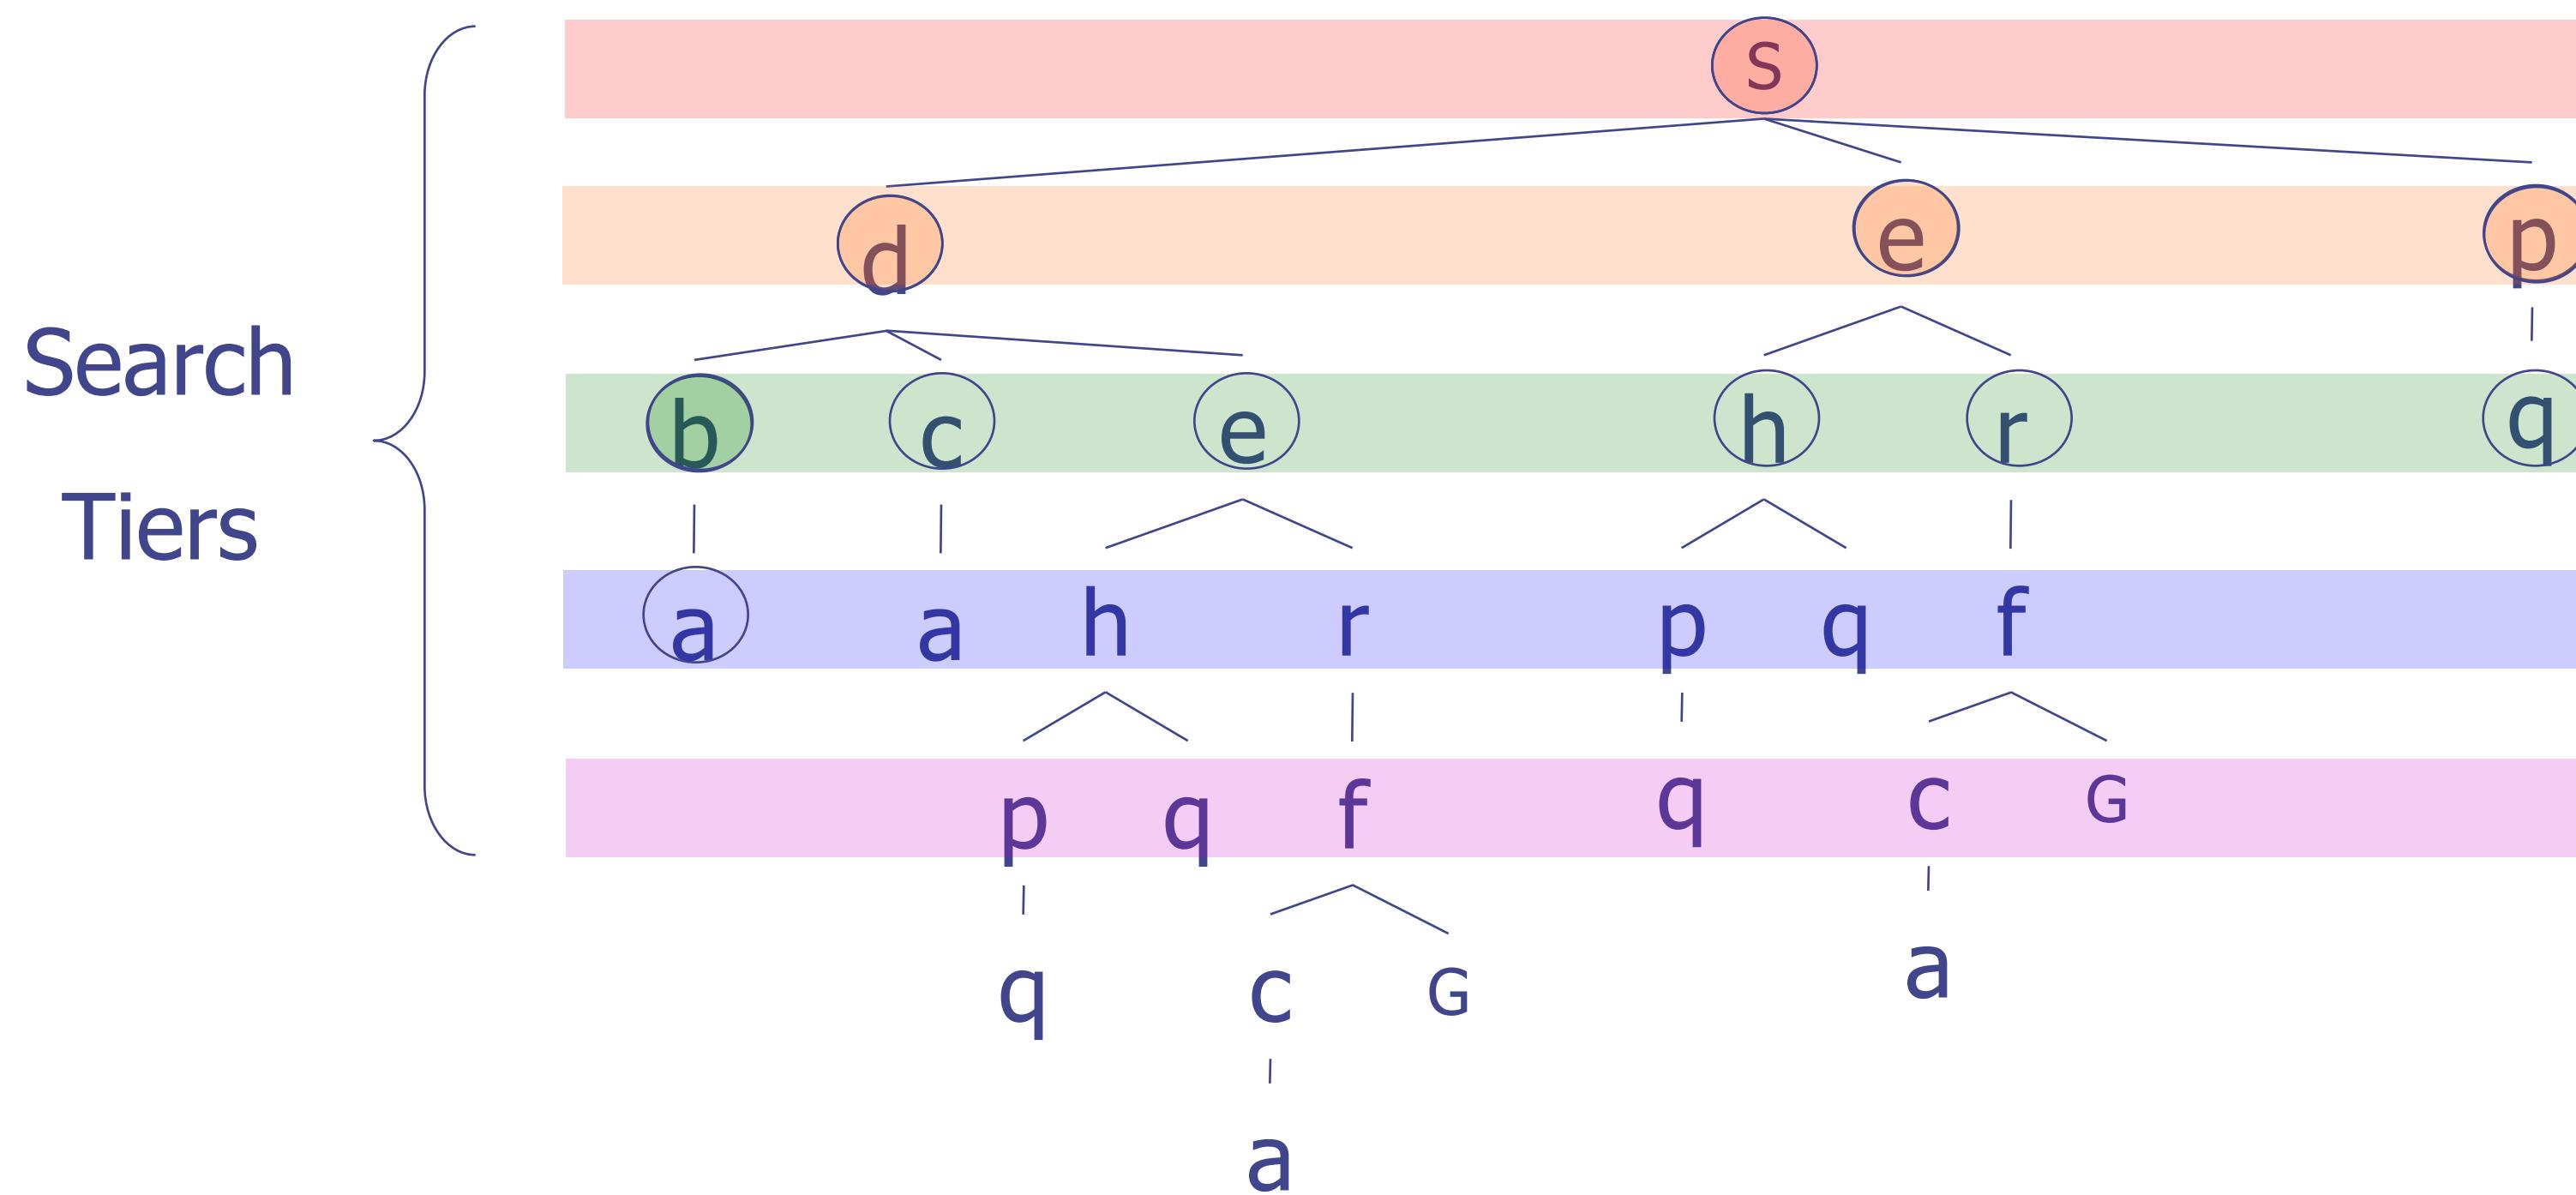
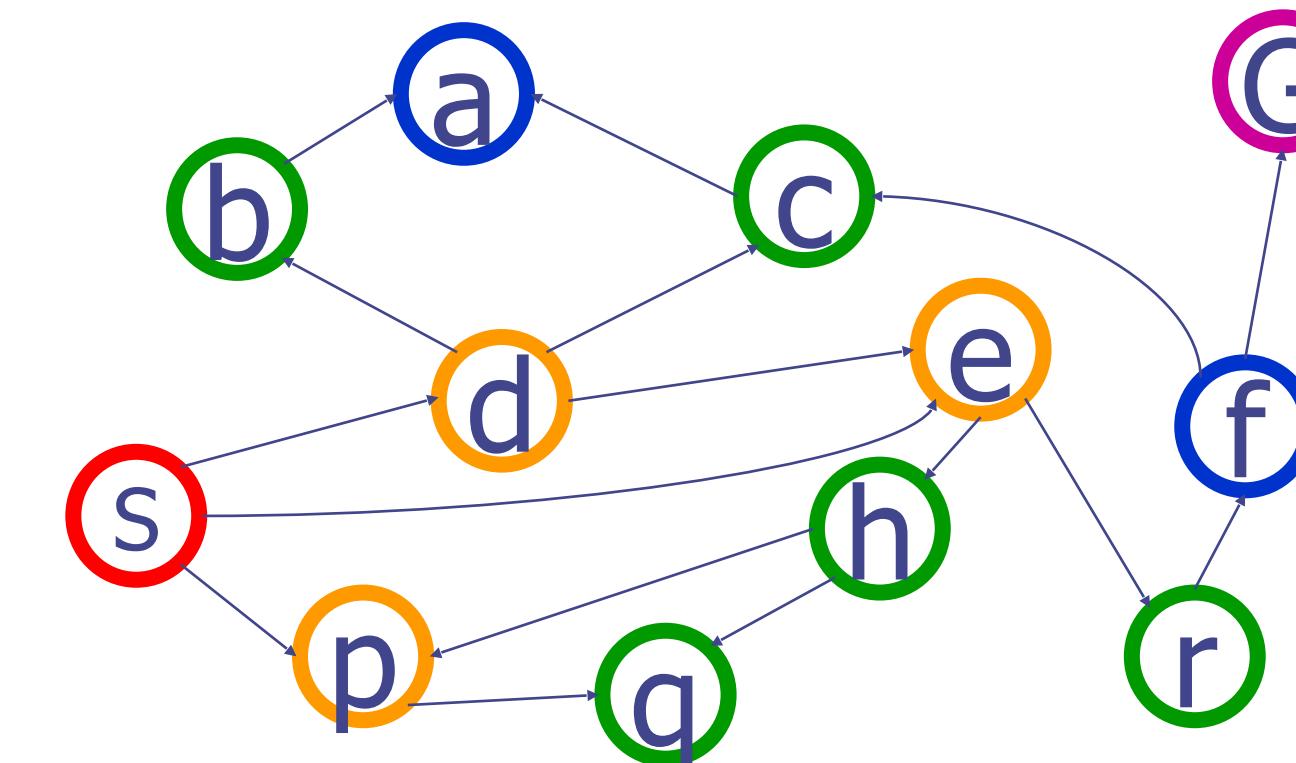
- If ‘b’ actions per state (branching factor) and maximum depth of the tree is ‘m’:
 - Memory complexity: $O(bm)$
 - In each level at most ‘b’ nodes will be added to the frontier set.
 - Time complexity: $O(b^m)$
 - In worst-case scenario, all nodes in the tree will be evaluated once.

Breadth-First Search

Breadth First Search

Strategy: expand shallowest node first

Implementation:
Fringe is a FIFO
queue



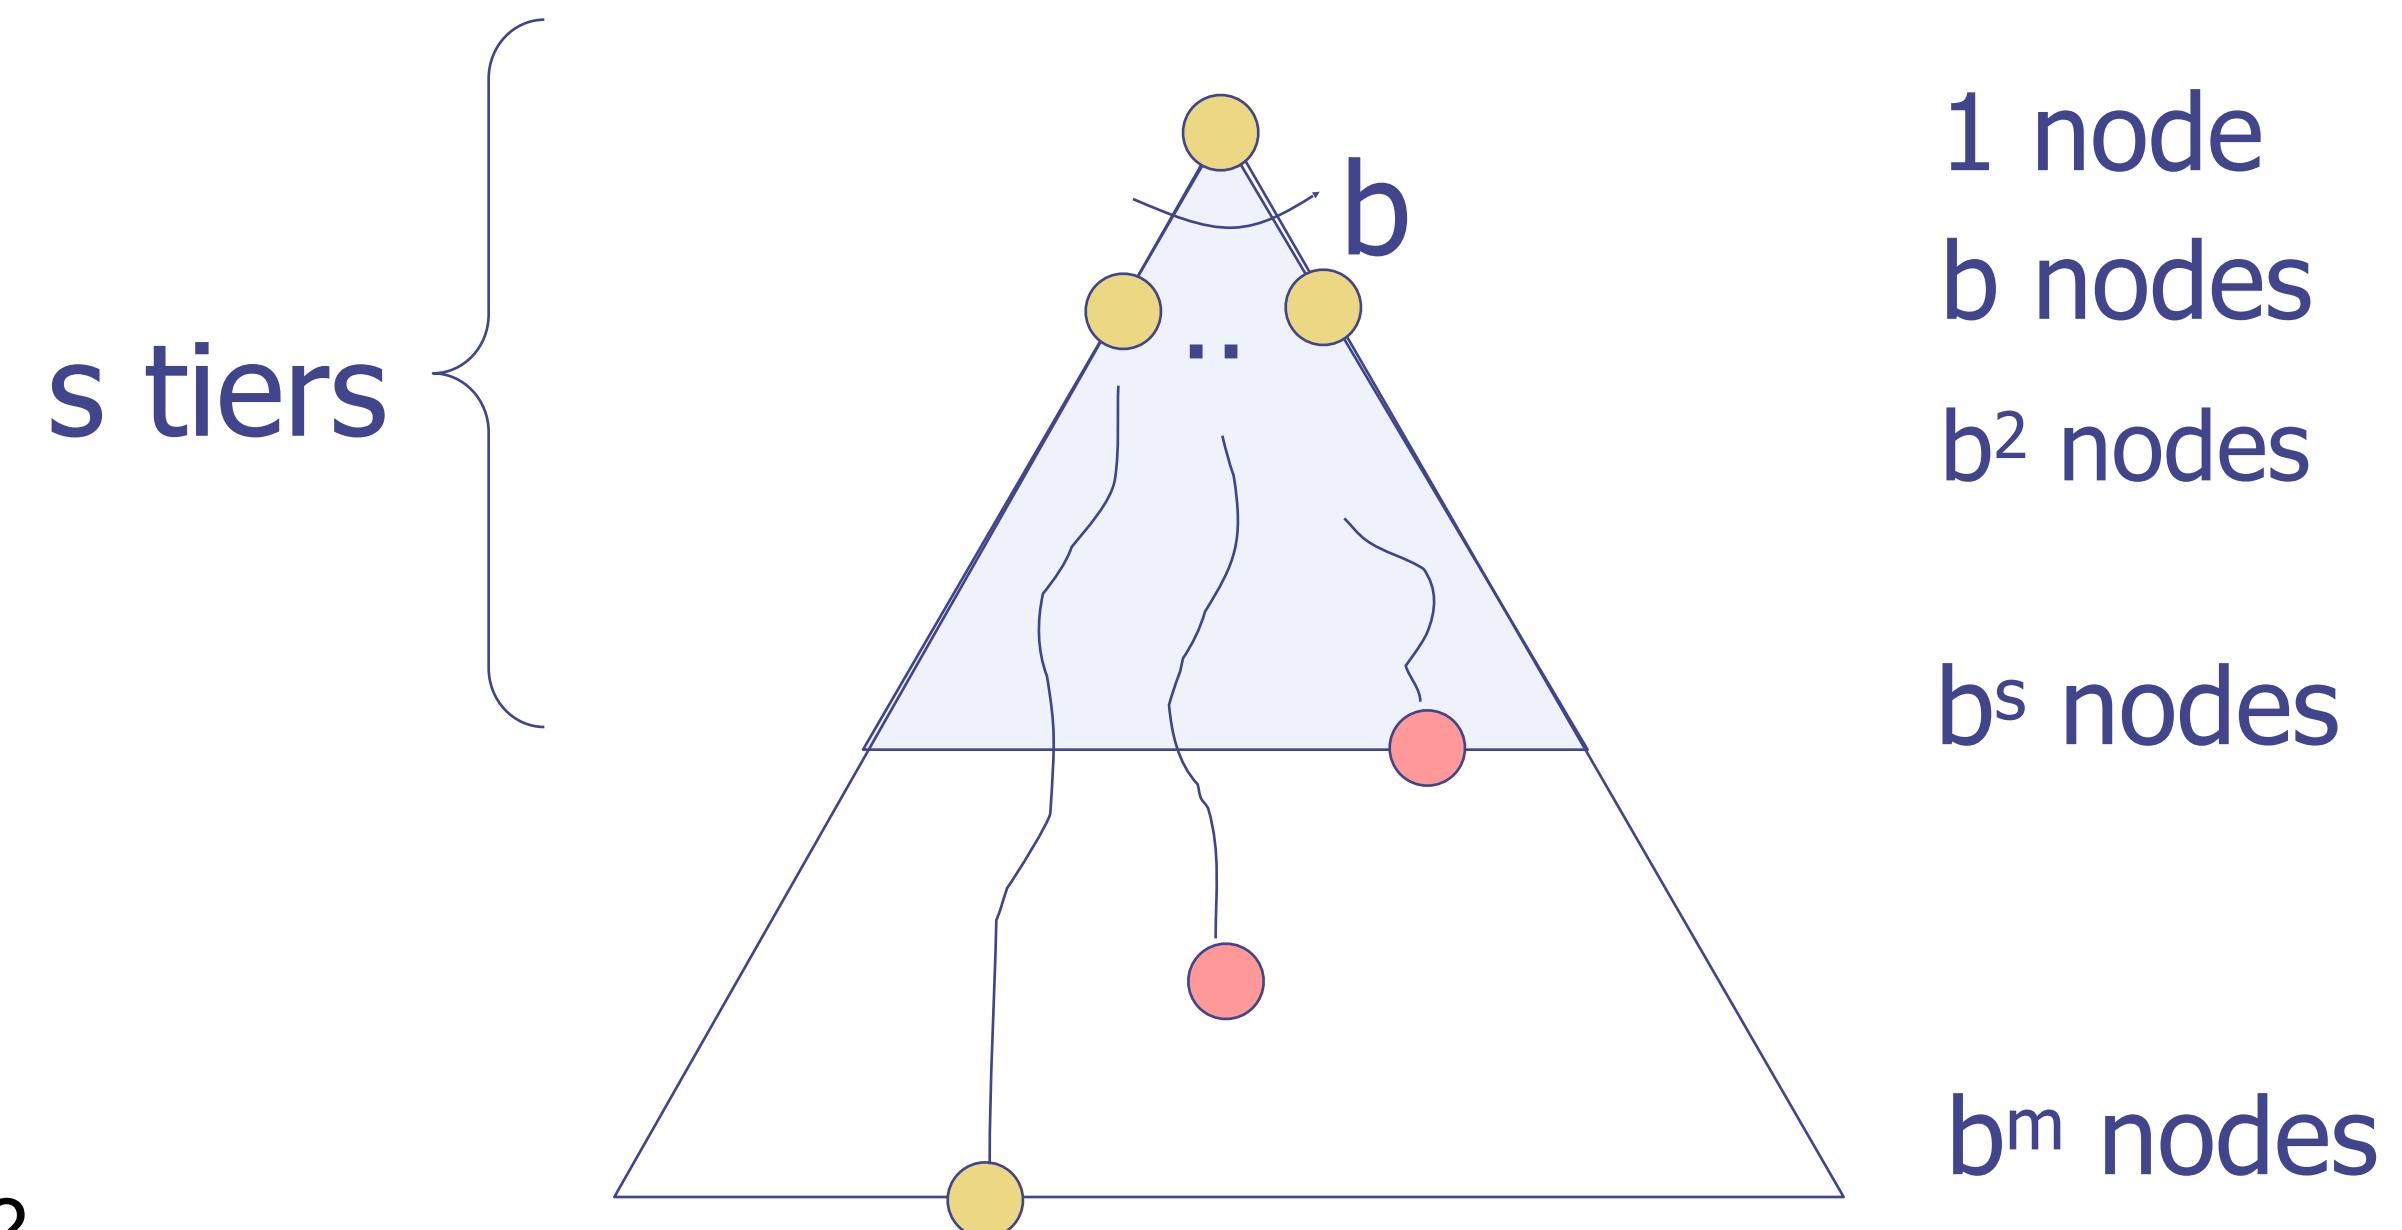
BFS Algorithm

```
function BREADTH-FIRST-SEARCH(problem) returns a solution, or failure
  node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
  frontier  $\leftarrow$  a FIFO queue with node as the only element
  explored  $\leftarrow$  an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node  $\leftarrow$  POP(frontier) /* chooses the shallowest node in frontier */
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        if problem.GOAL-TEST(child.STATE) then return SOLUTION(child)
        frontier  $\leftarrow$  INSERT(child, frontier)
```

BFS

Algorithm		Complete	Optimal	Time	Space
DFS	w/ Path Checking	Y	N	$O(b^{m+1})$	$O(bm)$
BFS		Y	N*	$O(b^{s+1})$	$O(b^s)$

- When is BFS optimal?



BFS Assumptions?

Recap and Upcoming

Today

- Solving problems by searching
 - Uninformed search strategies
 - BFS, DFS, ID-DFS, UCS

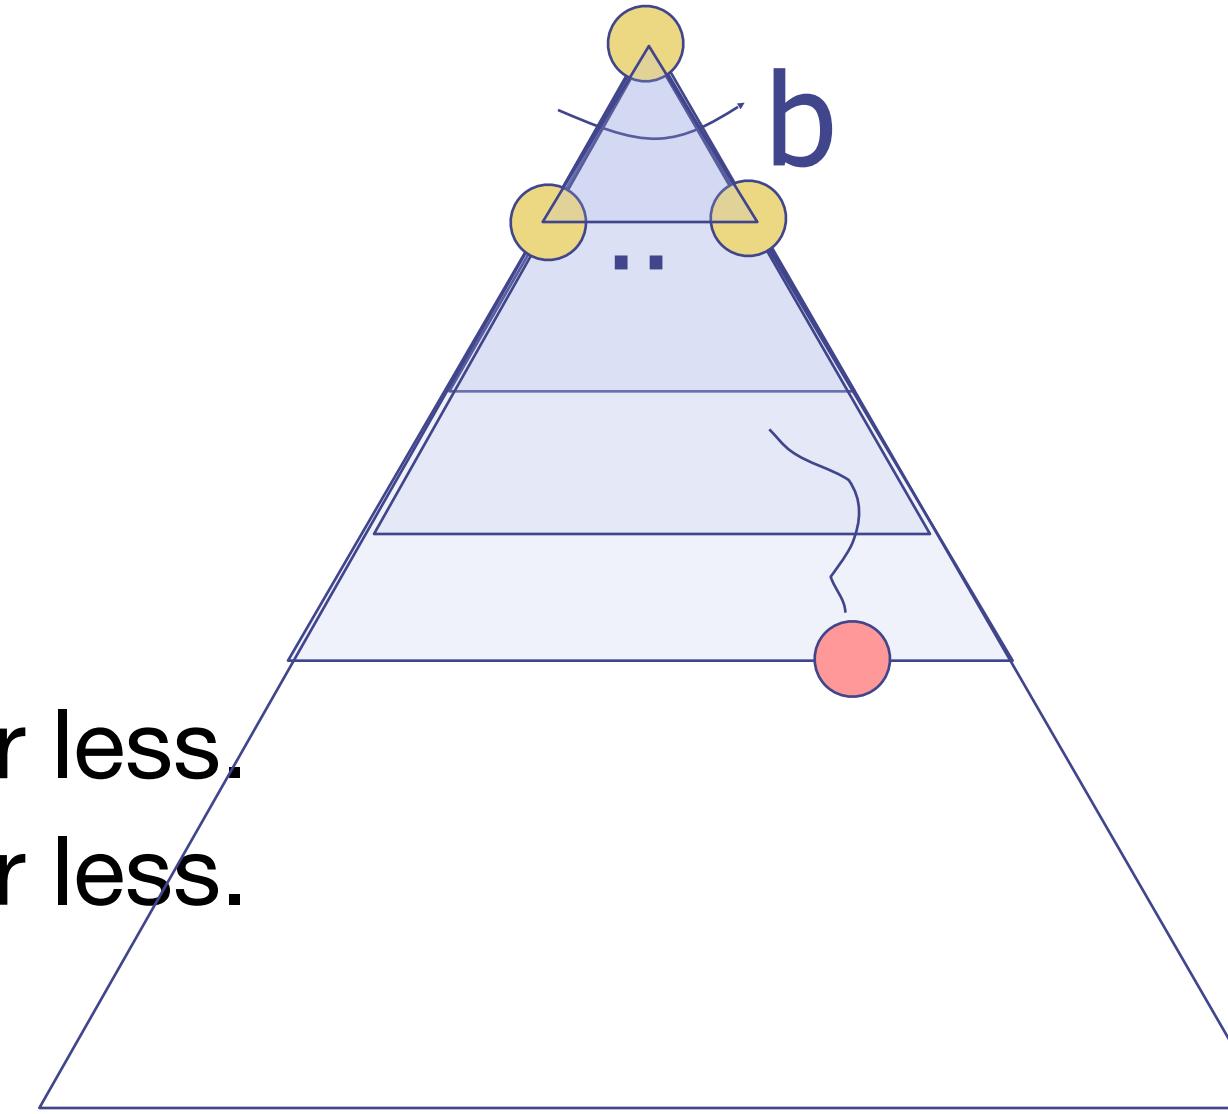
Next Week

- Solving problems by searching (cont.)
 - Informed search strategies
 - Heuristics functions
 - Search in complex environments
 - Hill climbing, simulated annealing, local beam search, evolutionary algorithm.

Iterative Deepening

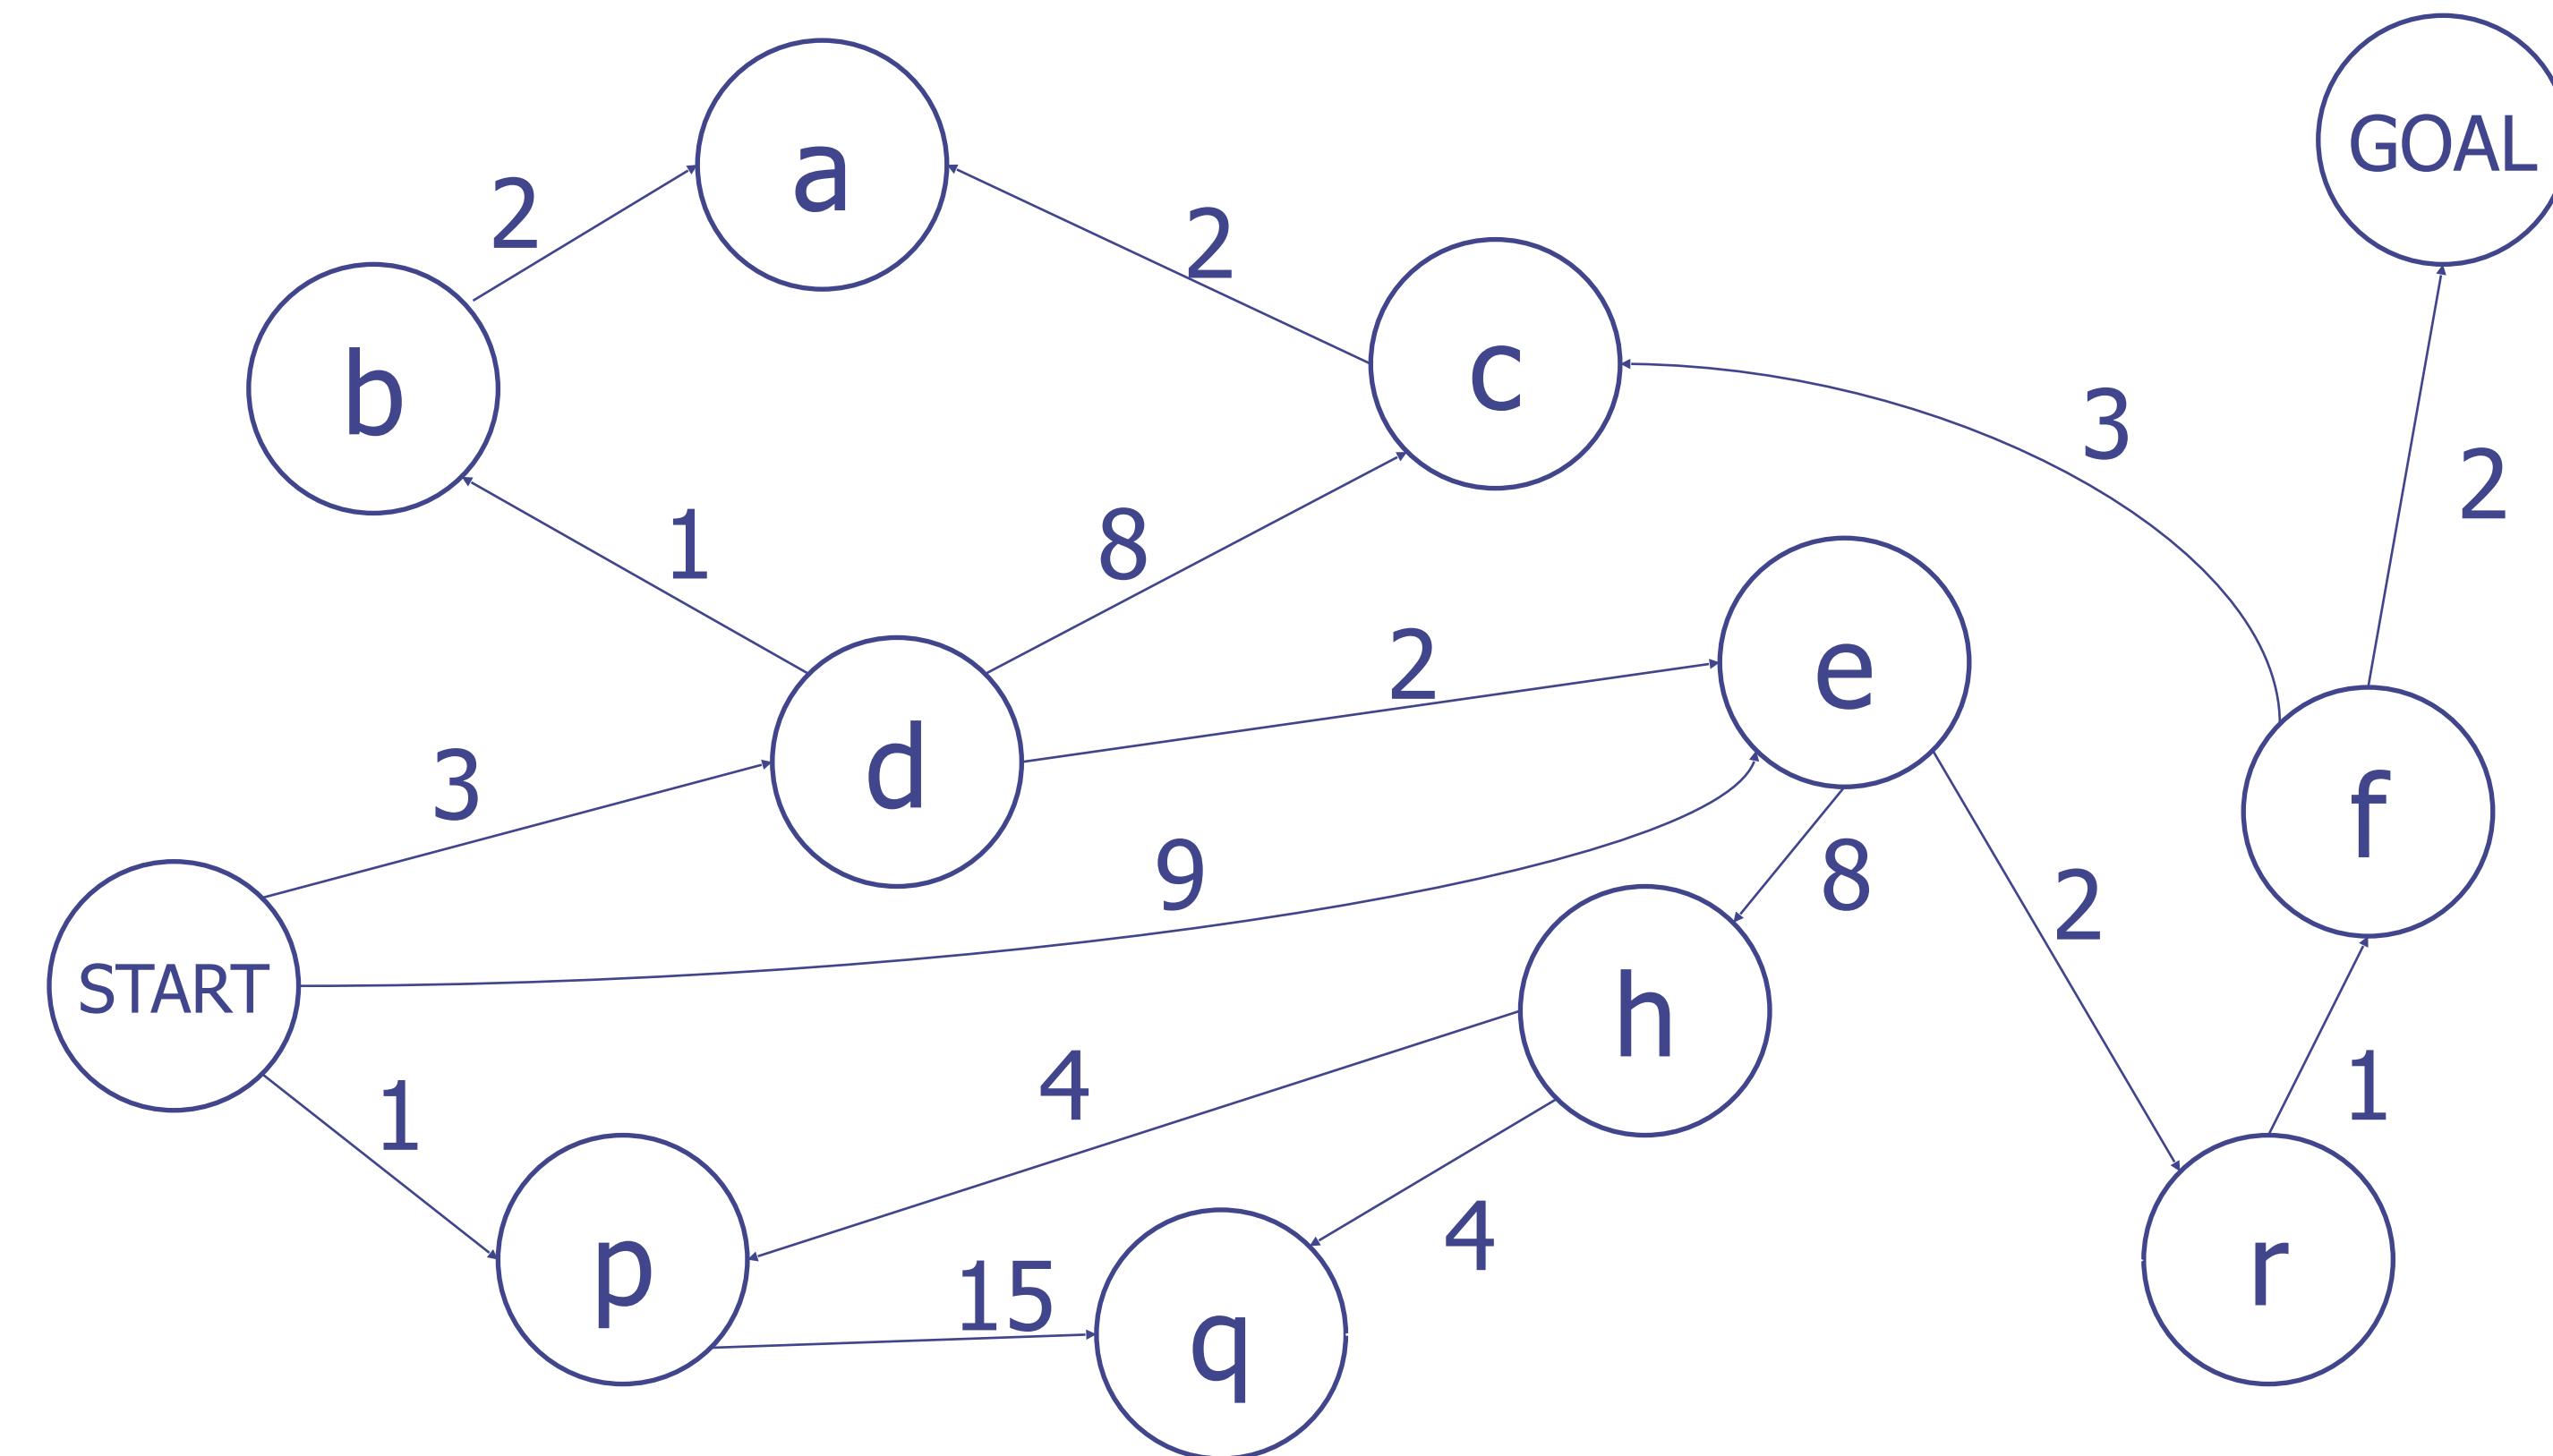
Iterative deepening uses DFS as a subroutine:

1. Do a DFS which only searches for paths of length 1 or less.
2. If “1” failed, do a DFS which only searches paths of length 2 or less.
3. If “2” failed, do a DFS which only searches paths of length 3 or less.
....and so on.



Algorithm		Complete	Optimal	Time	Space
DFS	w/ Path Checking	Y	N	$O(b^{m+1})$	$O(bm)$
BFS		Y	N*	$O(b^{s+1})$	$O(b^s)$
ID		Y	N*	$O(b^{s+1})$	$O(bs)$

Costs on Actions

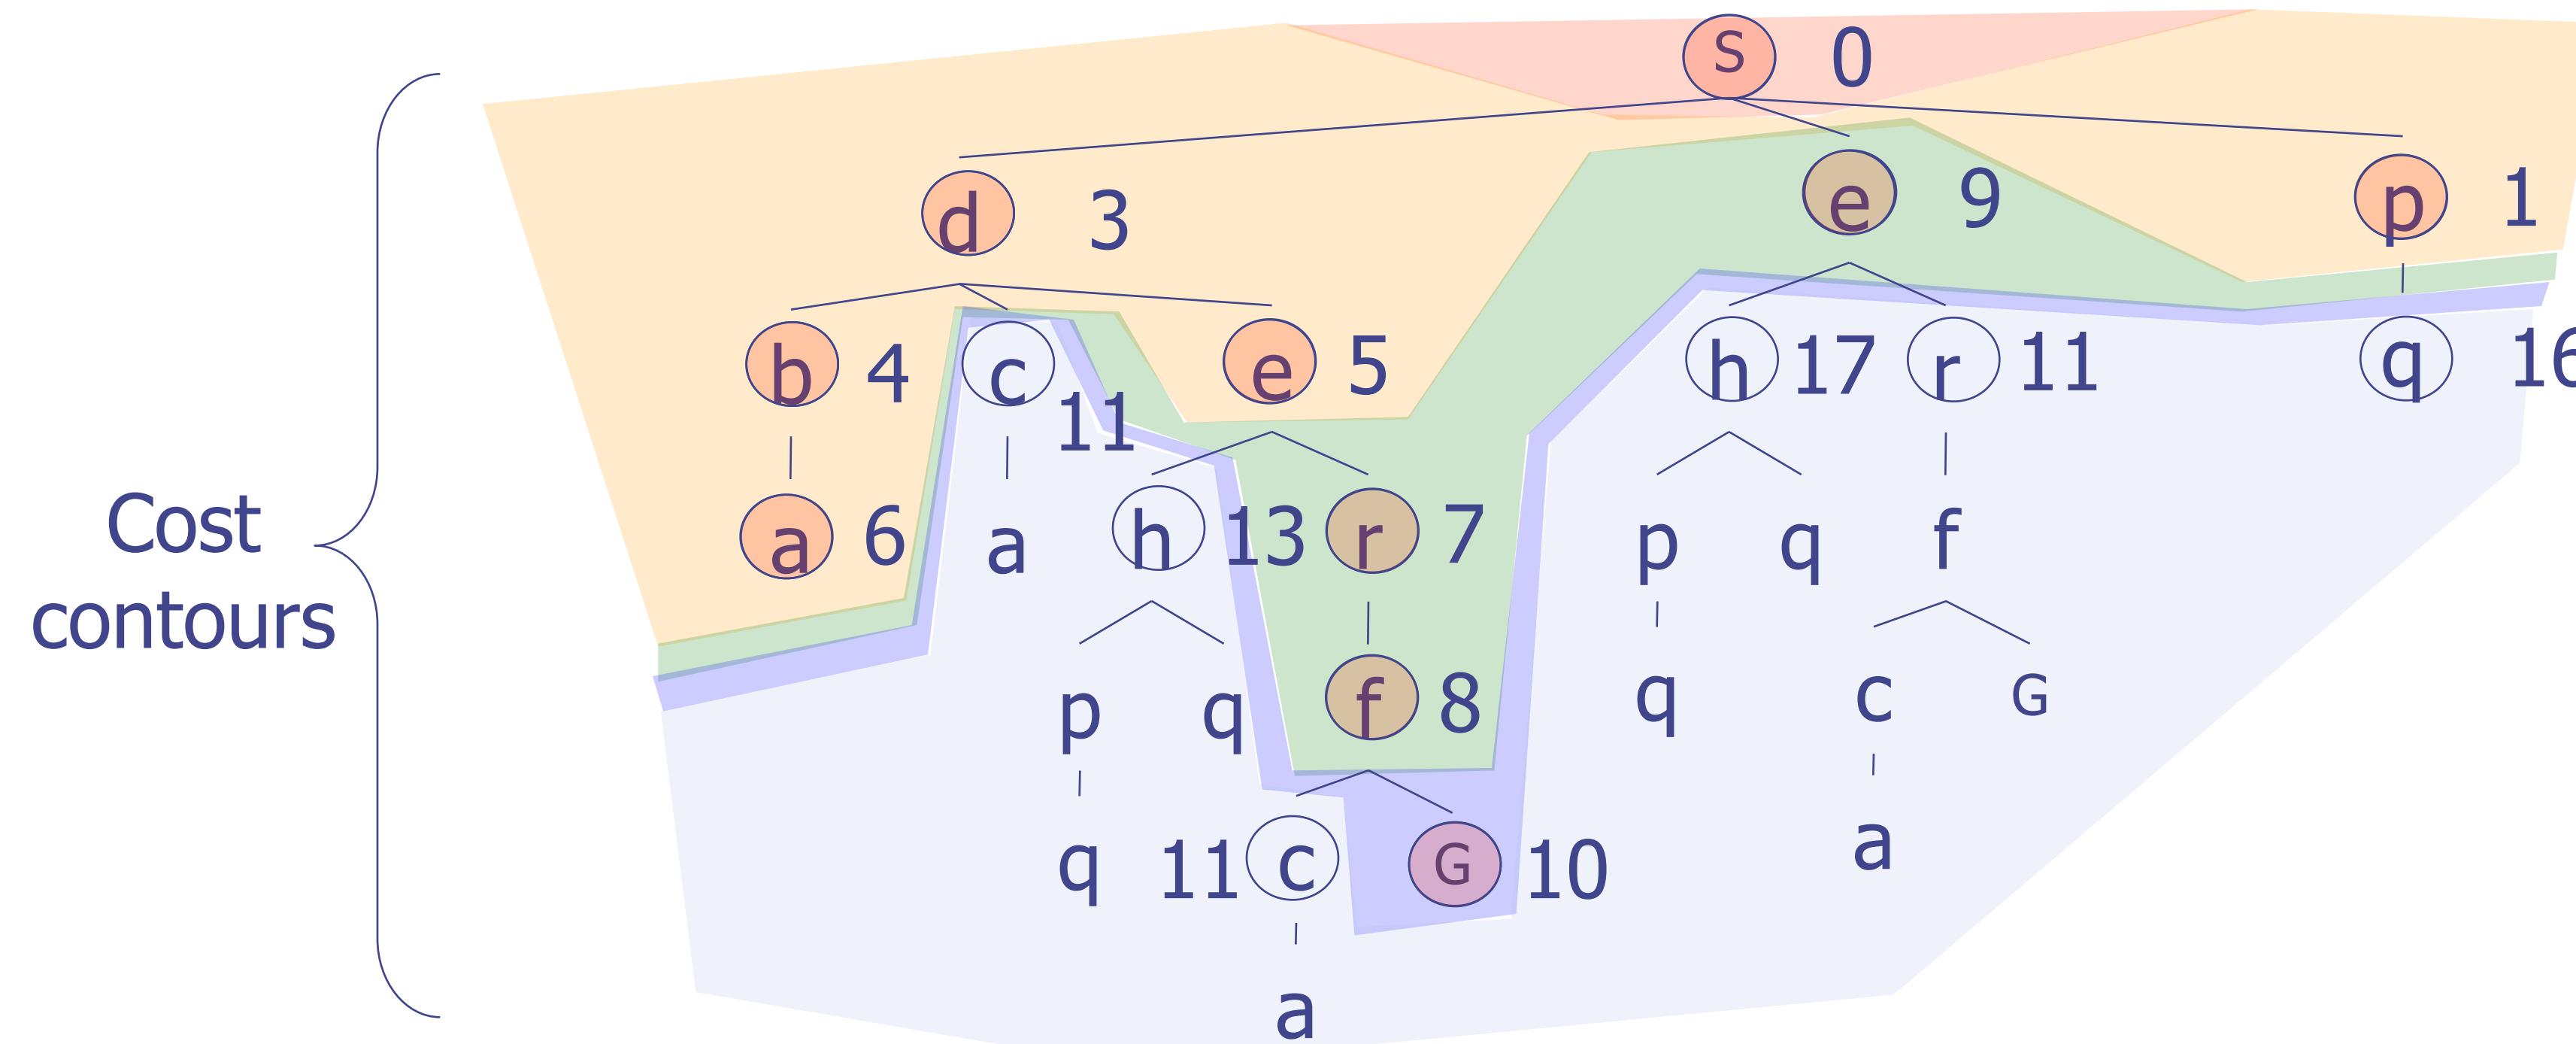
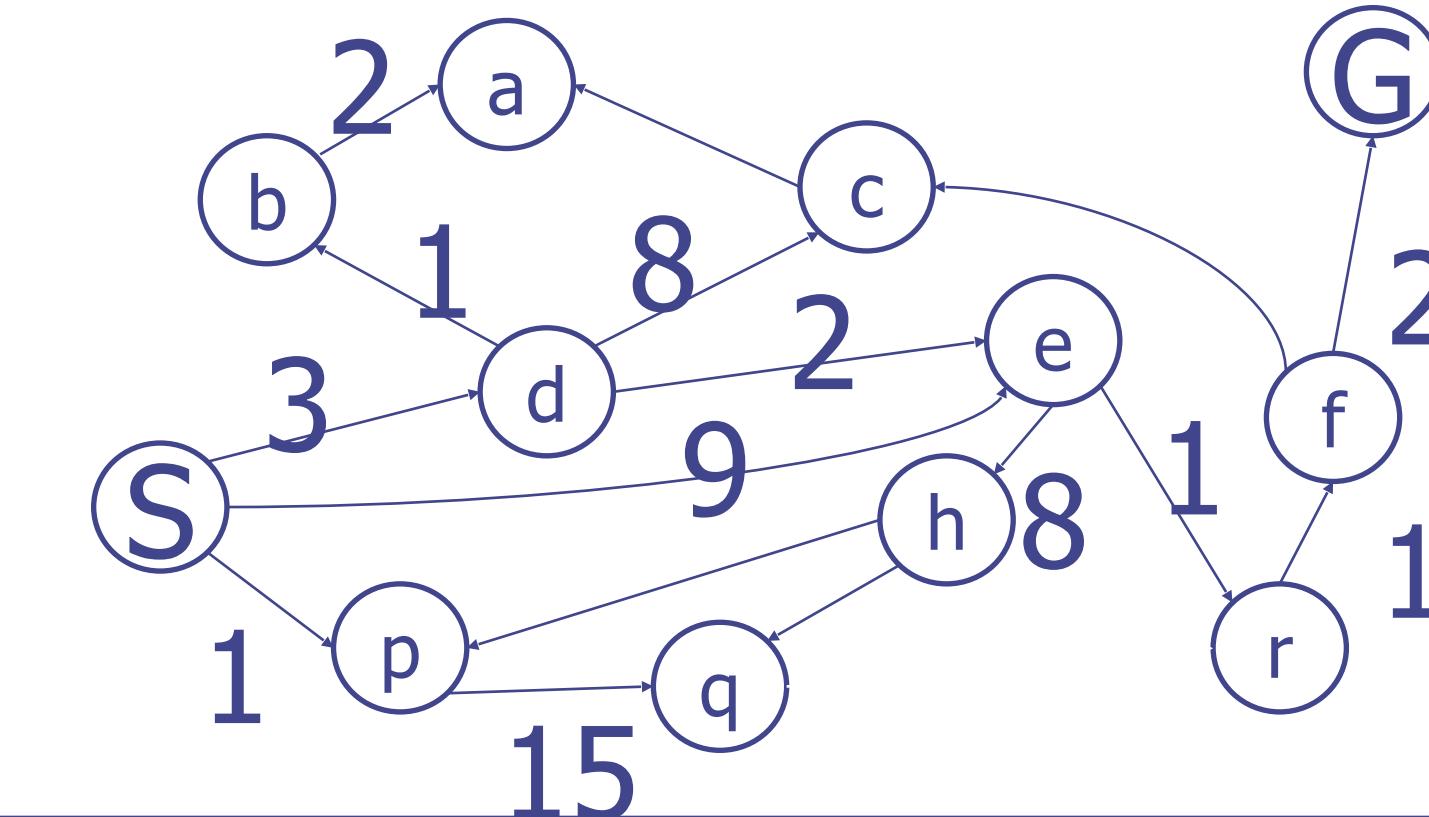


Notice that BFS finds the shortest path in terms of number of transitions. It does not find the least-cost path.

We will quickly cover an algorithm which does find the least-cost path.

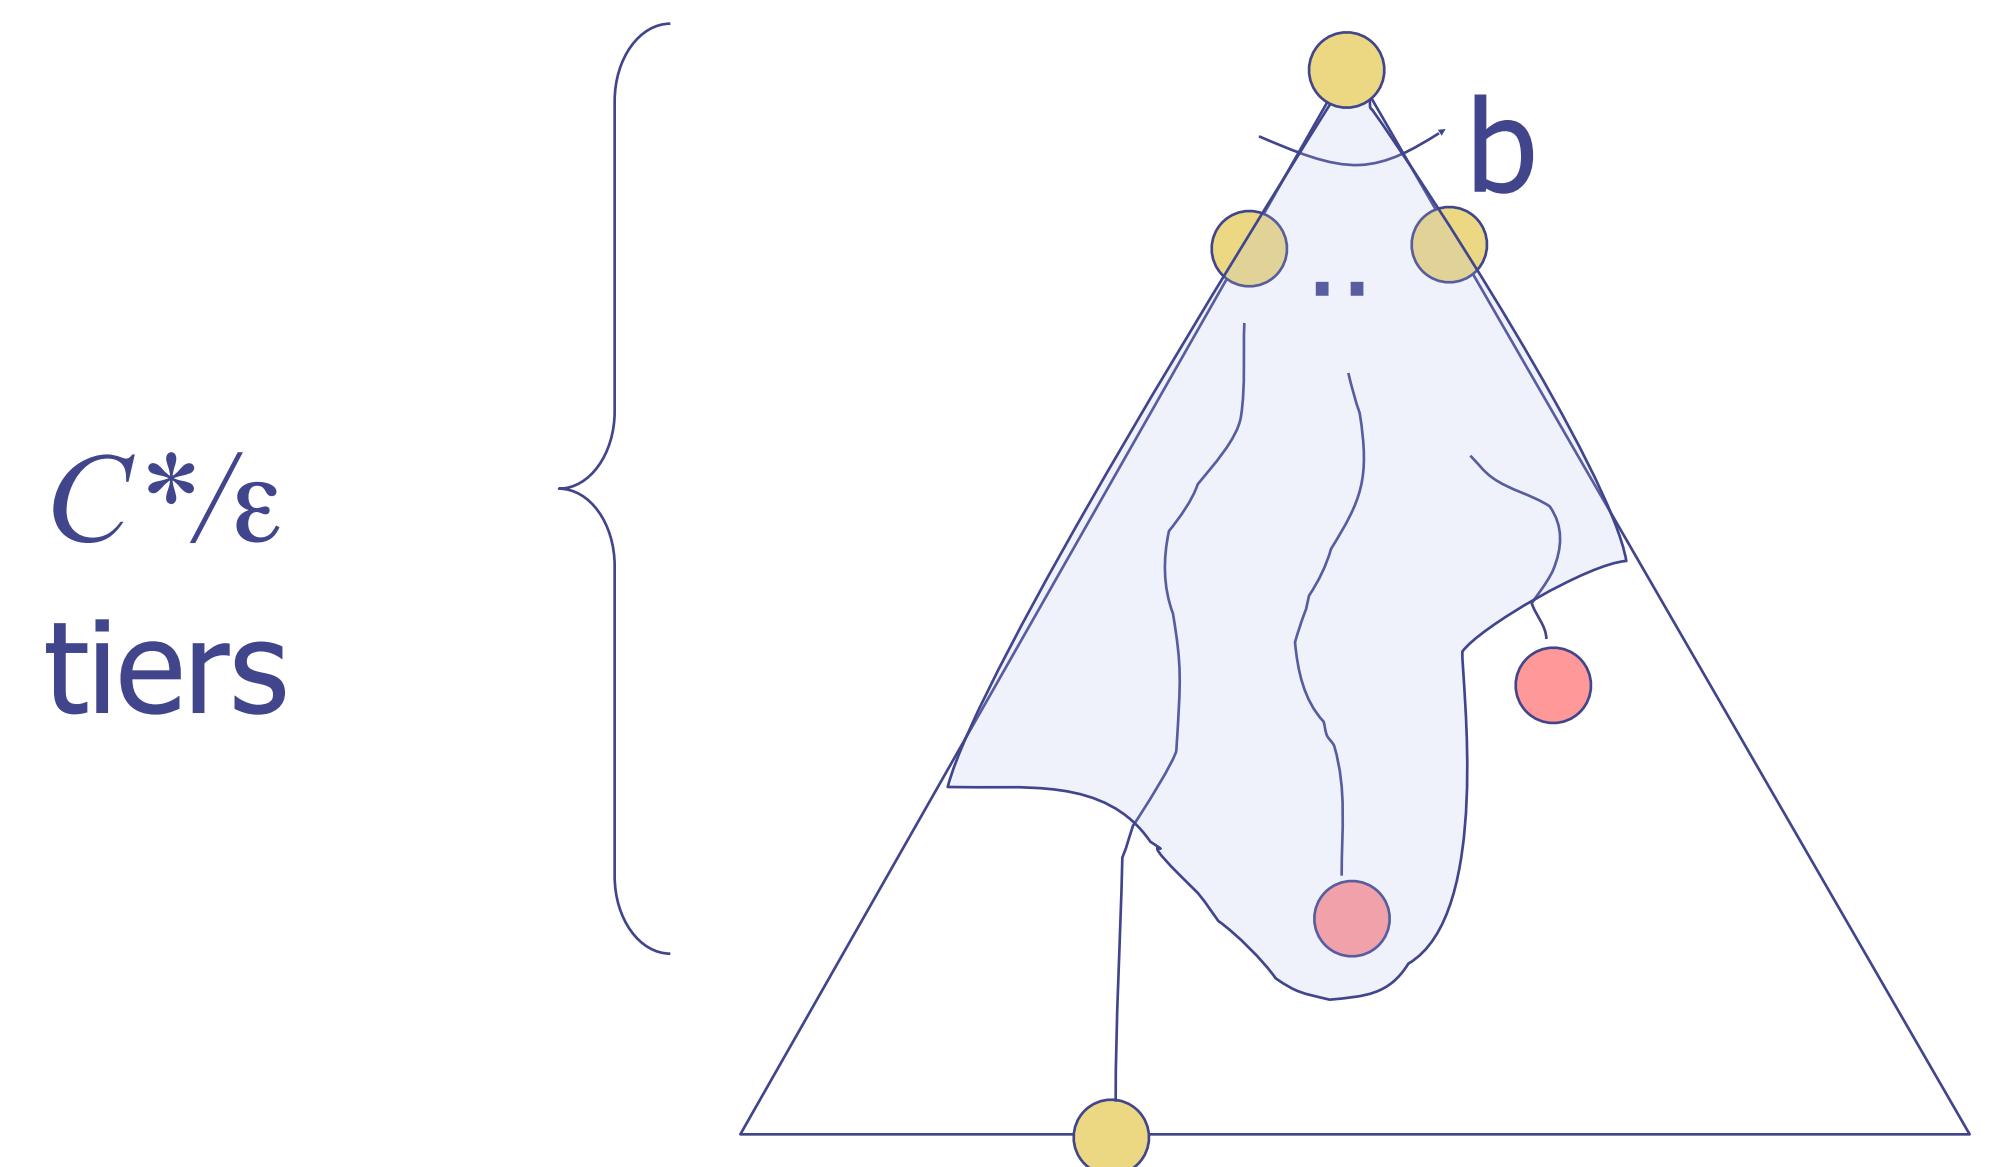
Uniform Cost Search

Expand cheapest node first:
Fringe is a priority queue



Uniform Cost Search

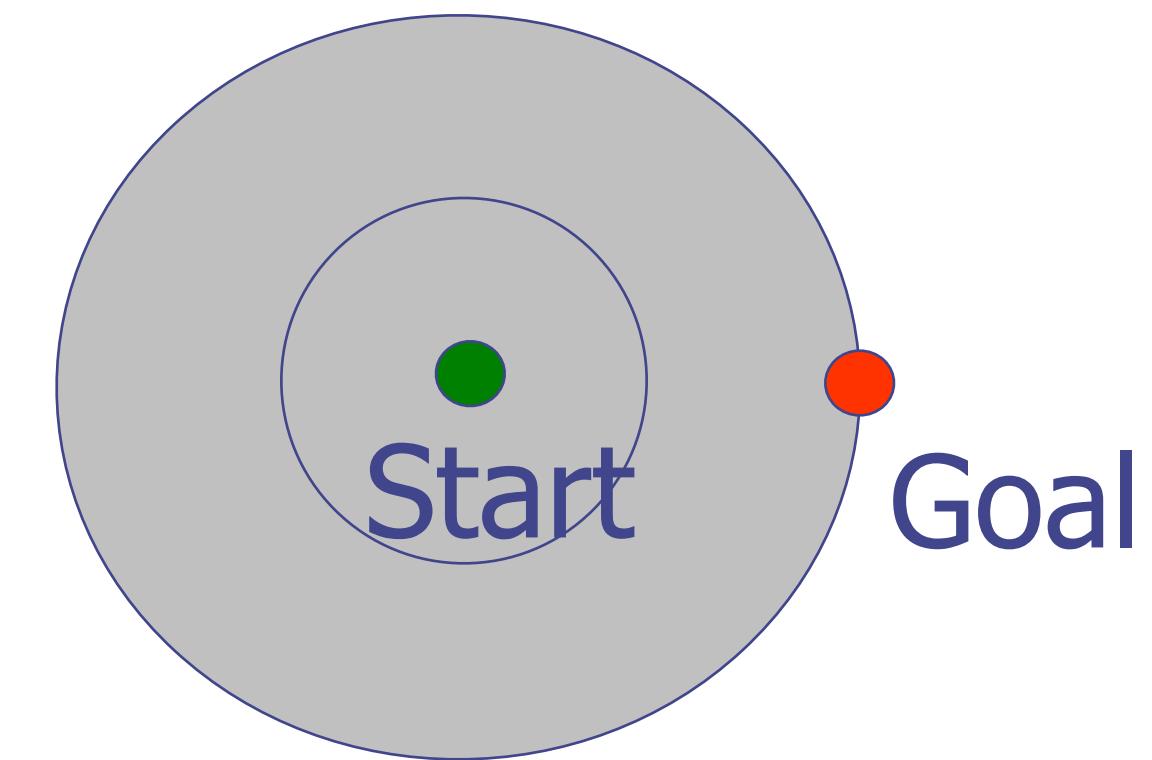
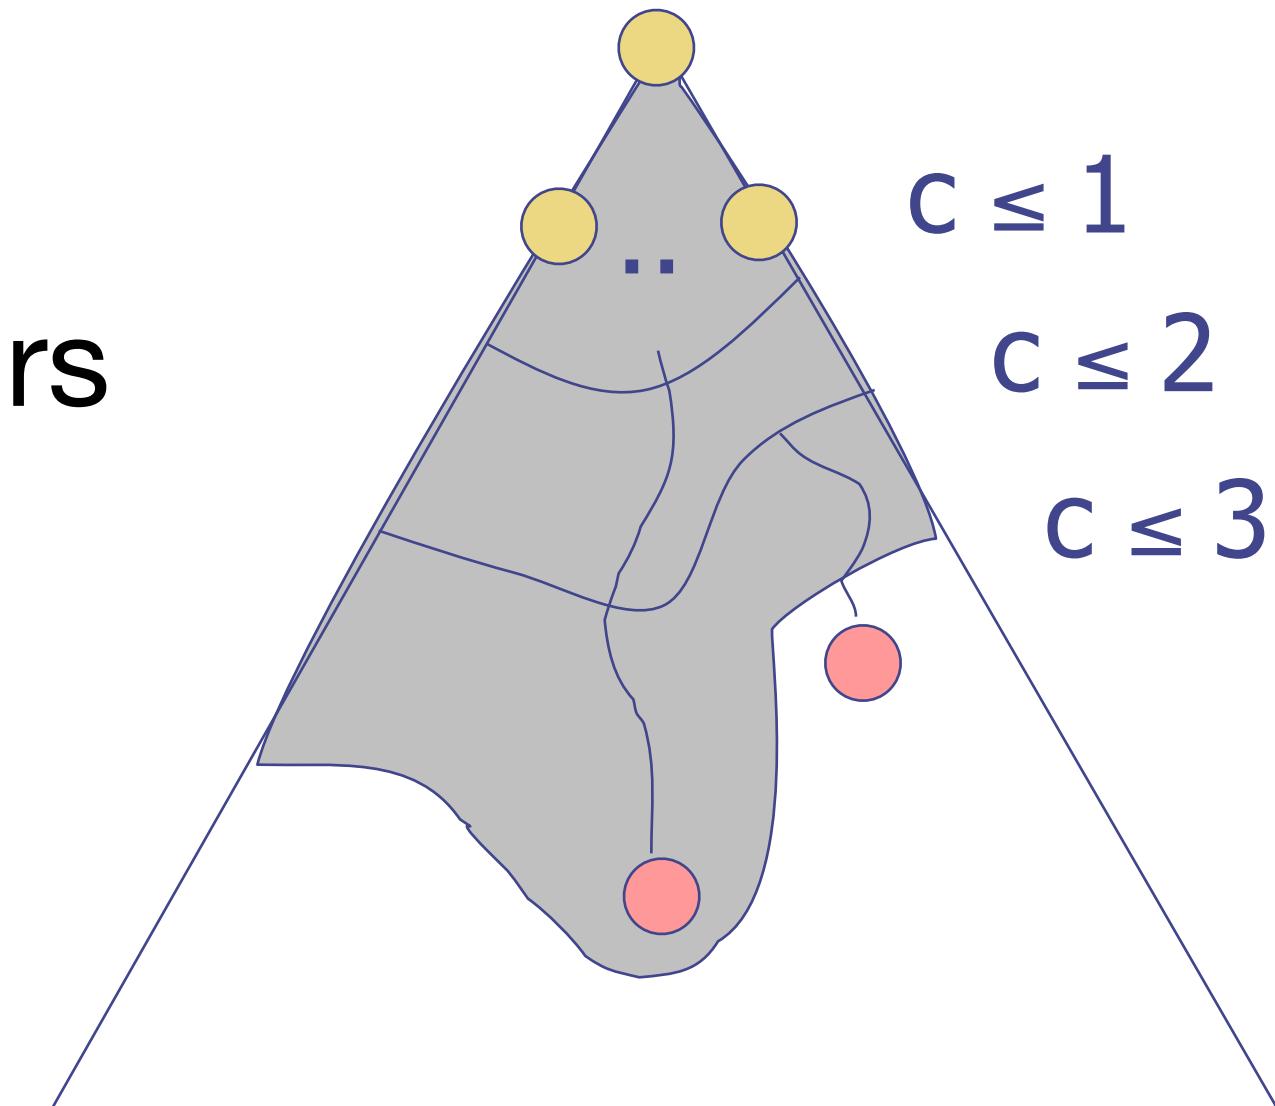
Algorithm		Complete	Optimal	Time	Space
DFS	w/ Path Checking	Y	N	$O(b^{m+1})$	$O(bm)$
BFS		Y	N	$O(b^{s+1})$	$O(b^s)$
UCS		Y*	Y	$O(b^{C^*/\varepsilon})$	$O(b^{C^*/\varepsilon})$



* UCS can fail if actions can get arbitrarily cheap

Uniform Cost Issues

- Remember: explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
 - Explores options in every “direction”
 - No information about goal location



Recap and Upcoming

Today

- Solving problems by searching
 - Uninformed search strategies
 - BFS, DFS, ID-DFS, UCS

Net Week

- Solving problems by searching (cont.)
 - Informed search strategies
 - Heuristics functions
 - Search in complex environments
 - Hill climbing, simulated annealing, local beam search, evolutionary algorithm.