

# **Agents that Plan Ahead: Search**

**Russell and Norvig: Chapter 3.1-3.4, 3.5-3.6**

**CSE 240: Winter 2023**

**Lecture 2**

# Announcements

## Announcements

- Assignment 1 is up
- Prof. Marinescu lecturing on Thursday.
  - I am planning to hold office hours on Thursday (might be subject to change).
- Use slack (#questions)



# Recap and Upcoming

## Today

- Uninformed search strategies
  - ID-DFS
  - Uniform Cost Search UCS
- Informed search strategies
  - Heuristics functions
  - Greedy Search algorithms
  - A\* search algorithm



# Review: Depth First Search

**Strategy: expand deepest node first**  
**Implementation: Fringe is a LIFO stack**




# Breadth First Search

Strategy: expand shallowest node first  
Implementation:  
Fringe is a FIFO queue



# Iterative Deepening DFS


- Idea: get DFS's space advantage with BFSs time / shallow-solution advantages
  - Run a DFS with depth limit 1. If no solution...
  - Run a DFS with depth limit 2. If no solution....
  - Run a DFS with depth limit 3 .....
- Isn't that wastefully redundant?
  - Generally most work happens in the deepest level searched, so not so bad!



# Iterative Deepening


Iterative deepening uses DFS as a subroutine:

1. Do a DFS which only searches for paths of length 1 or less.
2. If “1” failed, do a DFS which only searches paths of length 2 or less.
3. If “2” failed, do a DFS which only searches paths of length 3 or less.  
....and so on.



Algorithm		Complete	Optimal	Time	Space
DFS	w/ Path Checking	Y	N	$O(b^{m+1})$	$O(bm)$
BFS		Y	N*	$O(b^{s+1})$	$O(b^s)$
ID		Y	N*	$O(b^{s+1})$	$O(bs)$

# Cost-Sensitive Search



Notice that BFS finds the shortest path in terms of number of transitions. It does not find the least-cost path.



We will quickly cover an algorithm which does find the least-cost path.



# Uniform Cost Search


# Uniform Cost Search

Expand cheapest node first:  
Fringe is a priority queue



# High level Strategy

- Explored: states we have found the optimal path to
- Frontier: states we have seen, still figuring out how to get there cheaply
- Unexplored: states we have not seen




# Algorithm

```
function UNIFORM-COST-SEARCH(problem) returns a solution, or failure  
node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0  
frontier  $\leftarrow$  a priority queue ordered by PATH-COST, with node as the only element  
explored  $\leftarrow$  an empty set  
loop do  
  if EMPTY?(frontier) then return failure  
  node  $\leftarrow$  POP(frontier) /* chooses the lowest-cost node in frontier */  
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)  
  add node.STATE to explored  
  for each action in problem.ACTIONS(node.STATE) do  
    child  $\leftarrow$  CHILD-NODE(problem, node, action)  
    if child.STATE is not in explored or frontier then  
      frontier  $\leftarrow$  INSERT(child, frontier)  
    else if child.STATE is in frontier with higher PATH-COST then  
      replace that frontier node with child
```

# UCS Example




Example: UCS example



Start state: A, end state: D


# Analysis of UCS



## Theorem: correctness

When a state  $s$  is popped from the frontier and moved to explored, its priority is  $\text{PastCost}(s)$ , the minimum cost to  $s$ .


Proof:



# Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
  - Processes all nodes with cost less than cheapest solution!
  - If that solution costs  $C^*$  and arcs cost at least  $\epsilon$ , then the “effective depth” is roughly  $C^* / \epsilon$
  - Takes time  $O(b^{C^*/\epsilon})$  (exponential in effective depth)
- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^{C^*/\epsilon})$
- Is it complete?
  - Assuming best solution has a finite cost and minimum arc cost is positive, yes!
- Is it optimal?
  - Yes! (Proof in textbook)


$C^*/\epsilon$   
tiers





# Uniform Cost Search



Algorithm		Complete	Optimal	Time	Space
DFS	w/ Path Checking	Y	N	$O(b^{m+1})$	$O(bm)$
BFS		Y	N	$O(b^{s+1})$	$O(b^s)$
UCS		Y*	Y	$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$






# Uniform Cost Issues

- Remember: explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
  - Explores options in every “direction”
  - No information about goal location



# Search Heuristics

- Any estimate of how close a state is to a goal
- Designed for a particular search problem
- Examples: Manhattan distance, Euclidean distance

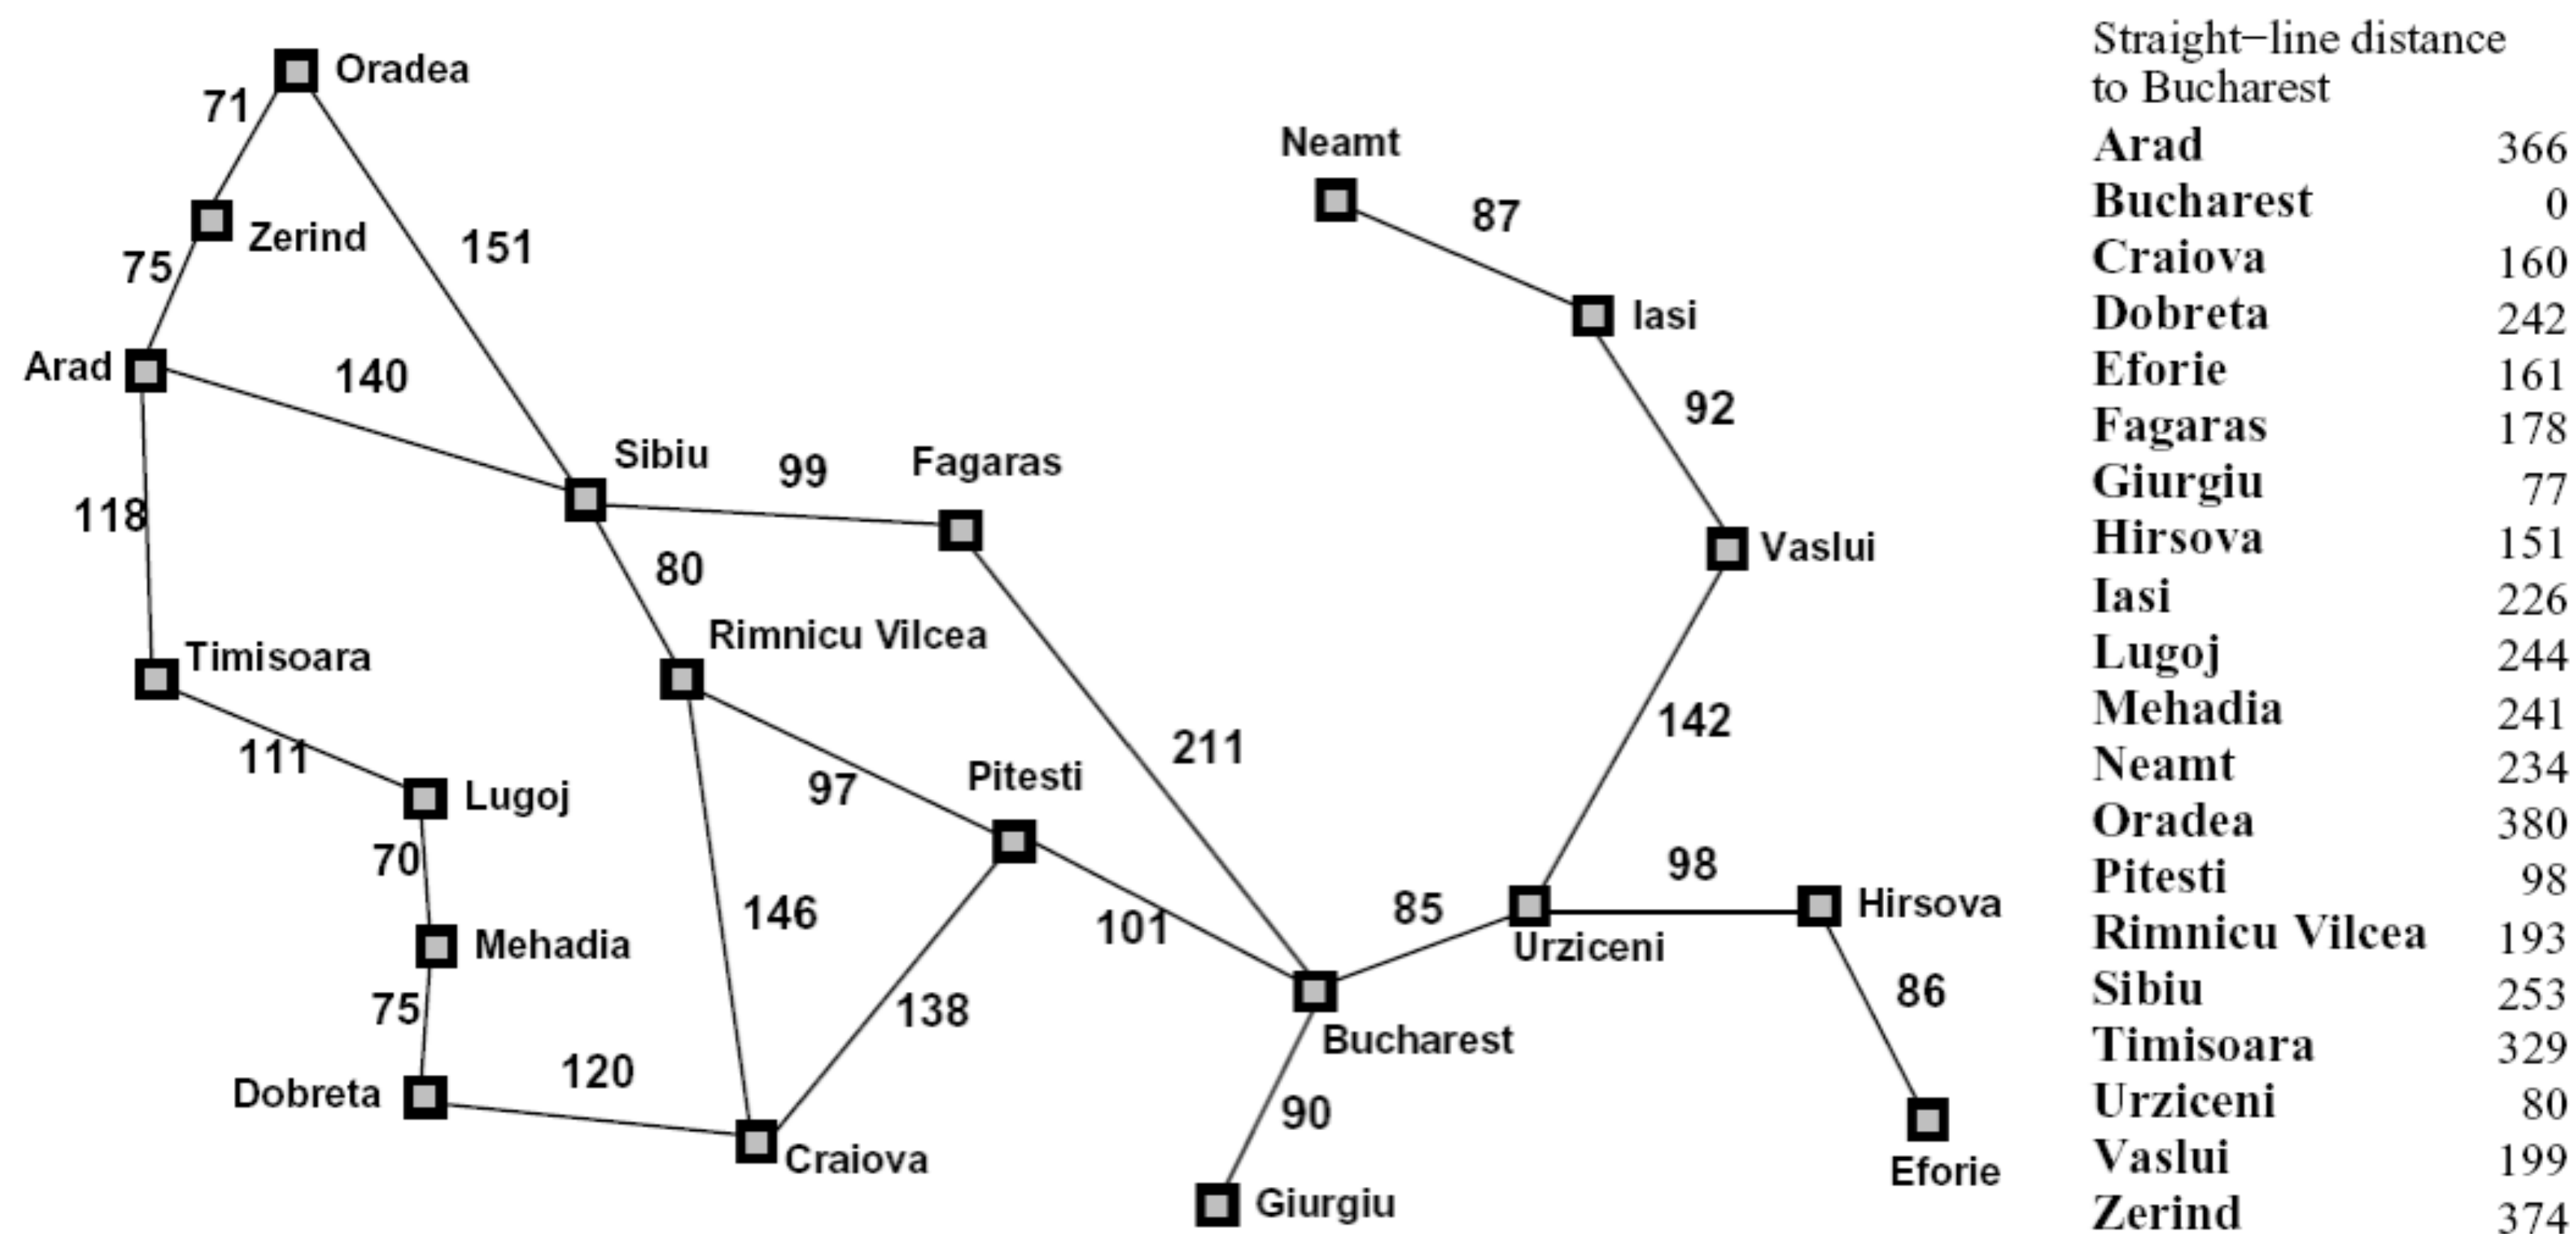


# CE 3: Heuristic brainstorm

- Heuristic for a specific search problem.
- Doesn't have to be optimal (just start thinking about it)
- Example: get a fruit and cereal from the grocery store.
- Heuristic:  $\text{dist}(\text{fruit}) + \text{dist}(\text{cereal})$

**Take ~5 minutes**

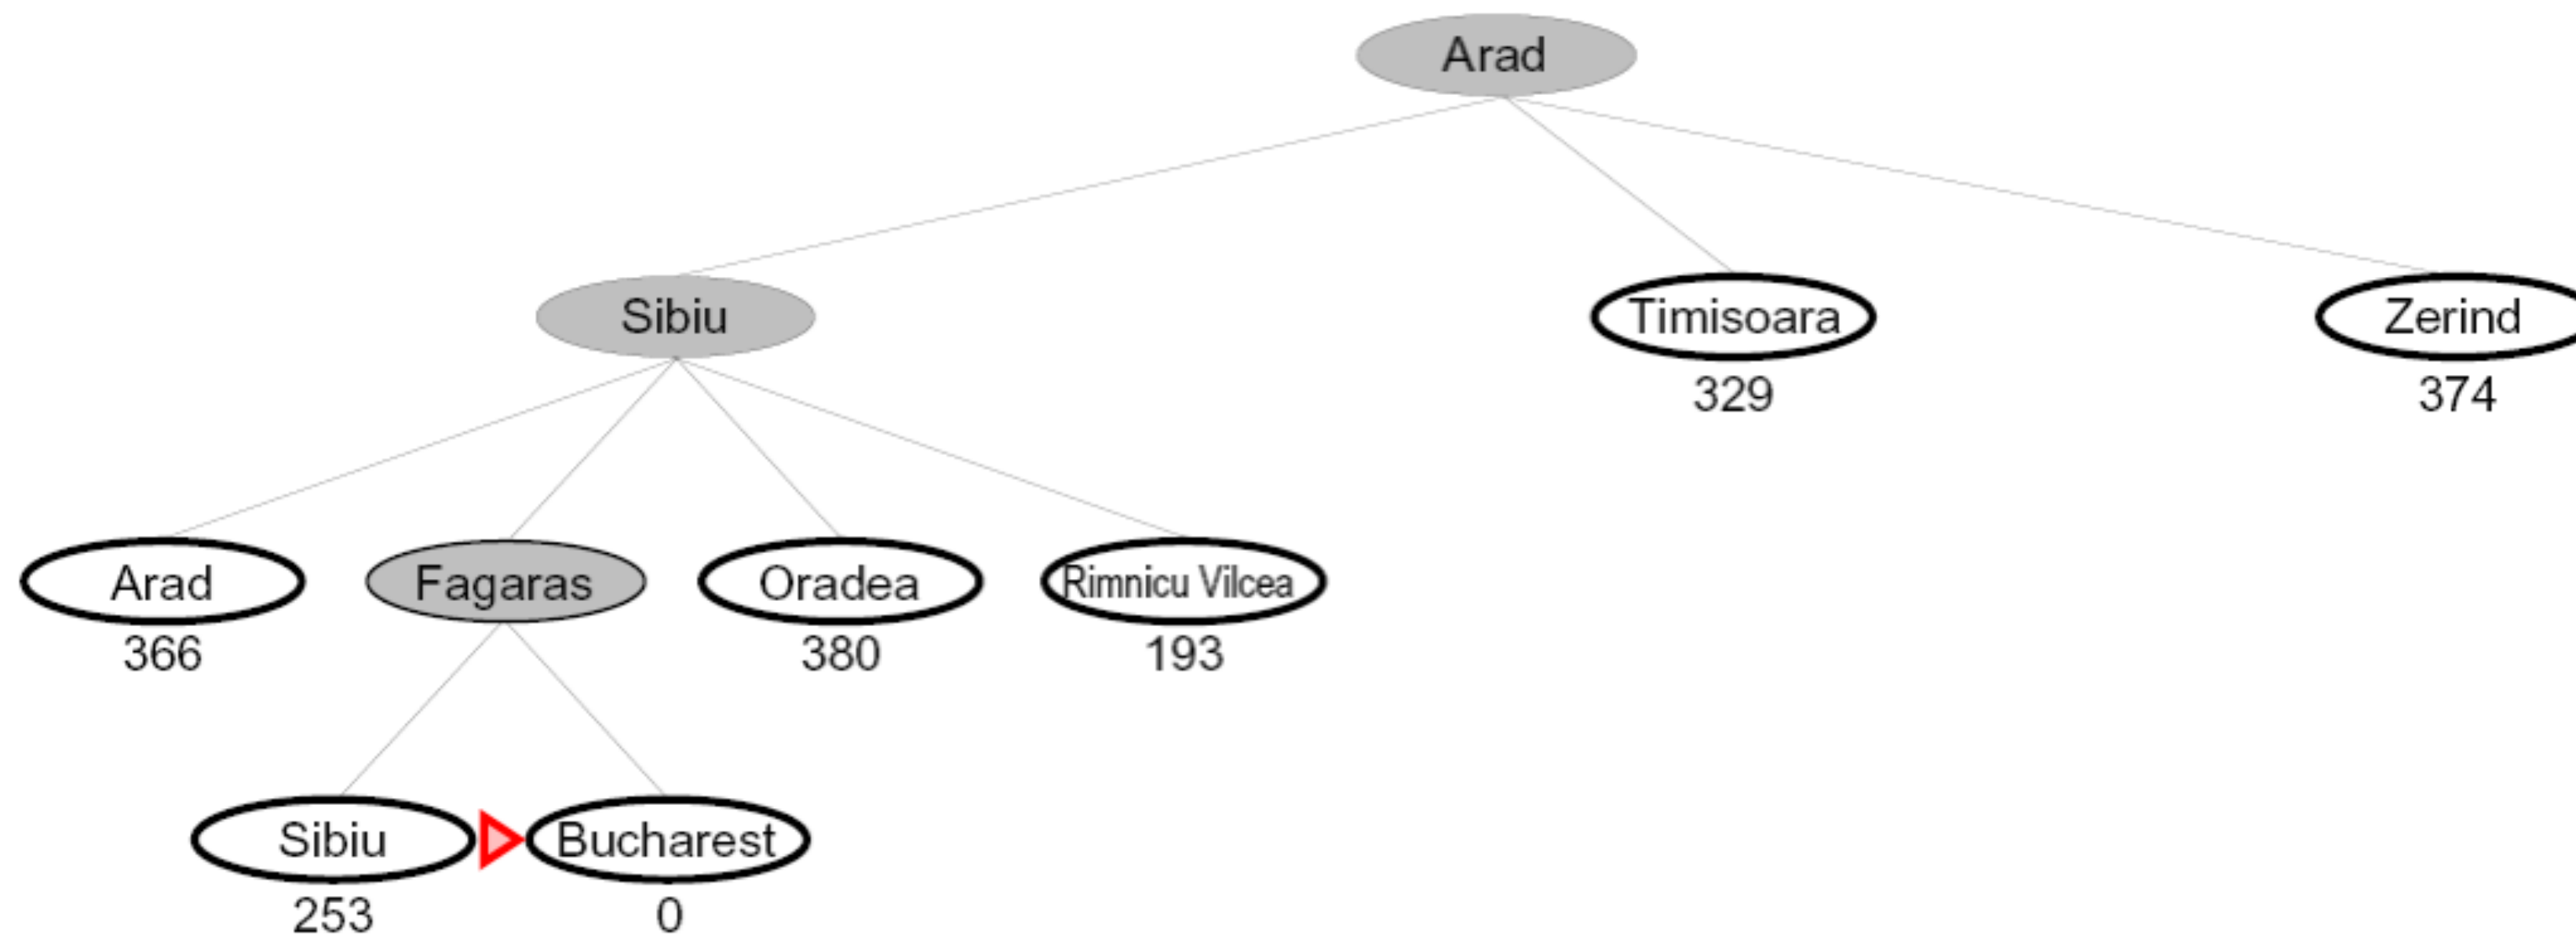
# Example of a Heuristic Function



# Greedy Search

# Best First / Greedy Search

- Expand the node that seems closest...




- What can go wrong?

# A\* Search

# Combining UCS and Greedy

$f(n) = g(n)$

- Uniform-cost orders by path cost, or backward cost  $g(n)$




Node	Fringe	f(n)
s	s->a	1
s->a	s->a->b	2
s->a	s->a->d	4
s->a	s->a->e	6



# Combining UCS and Greedy

- Uniform-cost orders by path cost, or backward cost  $g(n)$
- Greedy orders by goal proximity, or forward cost  $h(n)$




Node	Fringe	$f(n)$
s	s->a	5
s->a	s->a->b	6
s->a	s->a->d	2
s->a	s->a->e	1

# Combining UCS and Greedy

$f(n) = g(n) + h(n)$

- Uniform-cost orders by path cost, or backward cost  $g(n)$
- Greedy orders by goal proximity, or forward cost  $h(n)$



Node	Fringe	f(n)
s	s->a	6
s->a	s->a->b	8
s->a	s->a->d	6
s->a	s->a->e	7


- A\* Search orders by the sum:  $f(n) = g(n) + h(n)$

# Another Way to Implement A\*

Run UCS with modified edge costs in order to account for closeness to the goal state

$$\text{Cost}'(s, a) = \text{Cost}(s, a) + h(\text{succ}(s, a)) - h(s)$$


- Intuition: add a penalty for how much action 'a' takes us away from the end state



$$\text{Cost}'(C, B) = \text{Cost}(C, B) + h(B) - h(C) = 1 + (3 - 2) = 2$$


# When should A\* terminate?

- Should we stop when we enqueue a goal?



- No: only stop when we dequeue a goal

# Is A\* Optimal?



- What went wrong?
- Actual bad goal cost < estimated good goal cost
- We need estimates to be less than actual costs!


# Admissible Heuristics

- A heuristic  $h$  is **admissible** (optimistic) if:

$$h(n) \leq h^*(n)$$

- where  $h^*(n)$  is the true cost to a nearest goal

- Example:




- Coming up with admissible heuristics is most of what's involved in using  $A^*$  in practice.

# Consistent Heuristic

A heuristic  $h$  is “consistent” if

- $Cost'(s, a) = Cost(s, a) + h(succ(s, a)) - h(s) \geq 0$
- $h(s_{end}) = 0$






# Correctness of A\*

- If  $h$  is consistent, A\* returns the minimum cost path
- Consider any path

- Key identity:



$$\underbrace{\sum_{i=1}^L \text{Cost}'(s_{i-1}, a_i)}_{\text{modified path cost}} = \underbrace{\sum_{i=1}^L \text{Cost}(s_{i-1}, a_i)}_{\text{original path cost}} + \underbrace{h(s_L) - h(s_0)}_{\text{constant}}$$

- Therefore, A\* solves the original problem using UCS and therefore the algorithm has complete



# Efficiency of A\*

A\* explores all states satisfying

$$f(s) \leq f(s_{end}) - h(s)$$

- Interpretation: the larger  $h(s)$ , the better
- Proof: A\* explores all nodes 's' such that

$$f(s) + h(s) \leq f(s_{end}) + h(s_{end})$$

$$f(s) + h(s) \leq f(s_{end})$$

$$f(s) \leq f(s_{end}) - h(s)$$

# Recap

## Week 2

- **Solving problems by searching (cont.)**
  - **Informed search strategies**
  - Heuristics functions
- Search in complex environments
  - Hill climbing, simulated annealing, local beam search, evolutionary algorithm.