

Agents that Plan Ahead: A* Search

Russell and Norvig: Chapter 3.1-3.4, 3.5-3.6

CSE 240: Winter 2023

Lecture 4

Guest Lecture: Prof. Marinescu

Announcements

- Assignment 1 is up
- Prof. Marinescu lecturing today.
- Quizzes will be all remote on Canvas.
- Prof. Gilpin will update the class on Tuesday.



Agenda

Today

- Informed search strategies
 - A* search algorithm
 - Heuristics

Recap: Uniform Cost Issues

- Remember: explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
 - Explores options in every “direction”
 - No information about goal location




A* Search

Combining UCS and Greedy

$f(n) = g(n) + h(n)$

- Uniform-cost orders by path cost, or backward cost $g(n)$
- Greedy orders by goal proximity, or forward cost $h(n)$



Node	Fringe	f(n)
s	s->a	6
s->a	s->a->b	8
s->a	s->a->d	6
s->a	s->a->e	7


- A* Search orders by the sum: $f(n) = g(n) + h(n)$

Another Way to Implement A*

Run UCS with modified edge costs in order to account for closeness to the goal state

$$\mathbf{Cost'}(s, a) = \mathbf{Cost}(s, a) + \mathbf{h}(\mathbf{succ}(s, a)) - \mathbf{h}(s)$$


- Intuition: add a penalty for how much action 'a' takes us away from the end state



$$\mathbf{Cost'}(C, B) = \mathbf{Cost}(C, B) + \mathbf{h}(B) - \mathbf{h}(C) = 1 + (3 - 2) = 2$$


When should A* terminate?

- Should we stop when we enqueue a goal?



- No: only stop when we dequeue a goal


Is A* Optimal?



- What went wrong?
- Actual bad goal cost < estimated good goal cost
- We need estimates to be less than actual costs!

An Example Heuristic

- Would any heuristic work?
- Doesn't work because of the negative modified edge costs (or being pessimistic about the correct path)




Admissible Heuristics

- A heuristic h is **admissible** (optimistic) if:

$$h(n) \leq h^*(n)$$

- where $h^*(n)$ is the true cost to a nearest goal

- Example:




- Coming up with admissible heuristics is most of what's involved in using A^* in practice.

Consistent Heuristic

A heuristic h is “consistent” if

- $Cost'(s, a) = Cost(s, a) + h(succ(s, a)) - h(s) \geq 0$
- $h(s_{end}) = 0$




Correctness of A*

- If h is consistent, A* returns the minimum cost path.

- Consider any path

- Key identity:



$$\underbrace{\sum_{i=1}^L \text{Cost}'(s_{i-1}, a_i)}_{\text{modified path cost}} = \underbrace{\sum_{i=1}^L \text{Cost}(s_{i-1}, a_i)}_{\text{original path cost}} + \underbrace{h(s_L) - h(s_0)}_{\text{constant}}$$

- Therefore, A* solves the original problem using UCS, and therefore the algorithm is complete.

Efficiency of A*


A* explores all states satisfying
$$f(s) \leq f(s_{end}) - h(s)$$

- Interpretation: the larger $h(s)$, the better
- Proof: A* explores all nodes 's' such that:

$$\begin{aligned} f(s) + h(s) &\leq f(s_{end}) + h(s_{end}) \\ f(s) + h(s) &\leq f(s_{end}) \\ f(s) &\leq f(s_{end}) - h(s) \end{aligned}$$



Amount Explored

- If $h(s)=0$, then A^* is the same as UCS.
- If $h(s) = \text{FutureCost}(s)$, then A^* only explores nodes on a minimum cost path.
- Usually $h(s)$ is somewhere in between.



UCS versus A* Contours

- Uniform-cost expands equally in all “directions”
- A* expands mainly toward the goal, but does hedge its bets to ensure optimality



How Do we Get Good Heuristics?



Just Relax!

Relaxation

- Ideally, we use $h(s) = \text{FutureCost}(s)$, but that's as hard as solving the original problem.




Key idea: relaxation

Constraints make life hard. Get rid of them.
But this is just for the heuristic!



Relaxation Overview




Closed Form Solution




Example: knock down walls

Goal: move from triangle to circle



Hard



Easy

Heuristic:

$$h(s) = \text{ManhattanDistance}(s, (2, 5))$$


$$\text{e.g., } h((1, 1)) = 5$$

CE 4: What is a Relaxation of this Problem?




Example: knock down walls

Goal: move from triangle to circle



Hard



Easy

Heuristic:

$$h(s) = \text{ManhattanDistance}(s, (2, 5))$$

$$\text{e.g., } h((1, 1)) = 5$$

Easier Search



Example: original problem

Start state: 1

Walk action: from s to $s + 1$ (cost: 1)

Tram action: from s to $2s$ (cost: 2)

End state: n

Constraint: can't have more tram actions than walk actions.

State: (location, **#walk - #tram**)

Number of states goes from $O(n)$ to $O(n^2)$!

Easier Search



Example: relaxed problem

Start state: 1

Walk action: from s to $s + 1$ (cost: 1)

Tram action: from s to $2s$ (cost: 2)

End state: n

~~Constraint: can't have more tram actions than walk actions.~~

Original state: (location, #walk - #tram)

Relaxed state: location

Easier Search

- Compute relaxed $\text{FutureCost}_{\text{rel}}(\text{location})$ for each location $(1, \dots, n)$ using dynamic programming or UCS
- Modify UCS to compute all past costs in reversed relaxed problem (equivalent to future costs in relaxed problem!)



Example: reversed relaxed problem

Start state: n

Walk action: from s to $s - 1$ (cost: 1)

Tram action: from s to $s/2$ (cost: 2)

End state: 1

- Define heuristic for original problem: $h(\text{location}, \# \text{walk} - \# \text{tram}) = \text{FutureCost}_{\text{rel}}(\text{location})$

Independent Subproblems

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8


Goal State

- Original problem: tiles cannot overlap (constraint)
- Relaxed problem: tiles can overlap (no constraint)
- Relaxed solution: 8 indep. problems, each in closed form

General Framework

- Removing constraints (knock down walls, walk/tram freely, overlap pieces)
- Reducing edge costs (from ∞ to some finite cost)

- Example:



- Original: $\text{Cost}((1, 1), \text{East}) = \infty$
- Relaxed: $\text{Cost_rel}((1,1), \text{East}) = 1$

General Framework



Definition: relaxed search problem

A **relaxation** P_{rel} of a search problem P has costs that satisfy:


$$\text{Cost}_{\text{rel}}(s, a) \leq \text{Cost}(s, a).$$



Definition: relaxed heuristic

Given a relaxed search problem P_{rel} , define the **relaxed heuristic** $h(s) = \text{FutureCost}_{\text{rel}}(s)$, the minimum cost from s to an end state using $\text{Cost}_{\text{rel}}(s, a)$.

Consistency



Theorem: consistency of relaxed heuristics

Suppose $h(s) = \text{FutureCost}_{\text{rel}}(s)$ for some relaxed problem P_{rel} .

Then $h(s)$ is a consistent heuristic.

- Proof:

$$\begin{aligned} h(s) &\leq \text{Cost}_{\text{rel}}(s, a) + h(\text{Succ}(s, a)) \text{ [triangle inequality]} \\ &\leq \text{Cost}(s, a) + h(\text{Succ}(s, a)) \text{ [relaxation]} \end{aligned}$$

Trade-off

- Efficiency
 - $h(s) = \text{FutureCost_rel}(s)$ must be easy to compute
 - Closed form, easier search, independent subproblems
- Tightness
 - heuristic $h(s)$ should be close to $\text{FutureCost}(s)$
 - Don't remove too many constraints

Recap

Week 2 Summary

- Solving problems by searching
 - Informed search strategies
 - Heuristics functions

Next Week

- Search in complex environments
- Hill climbing, simulated annealing, local beam search, evolutionary algorithm.